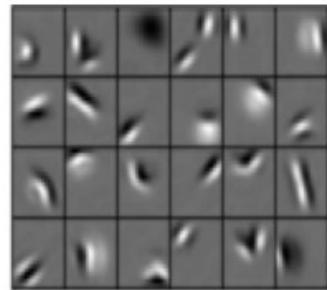


CNN Architectures

Learning Visual Features

**learn hierarchy of features
directly from the data**
(rather than hand-engineering them)

Low level features



Edges, dark spots

Mid level features



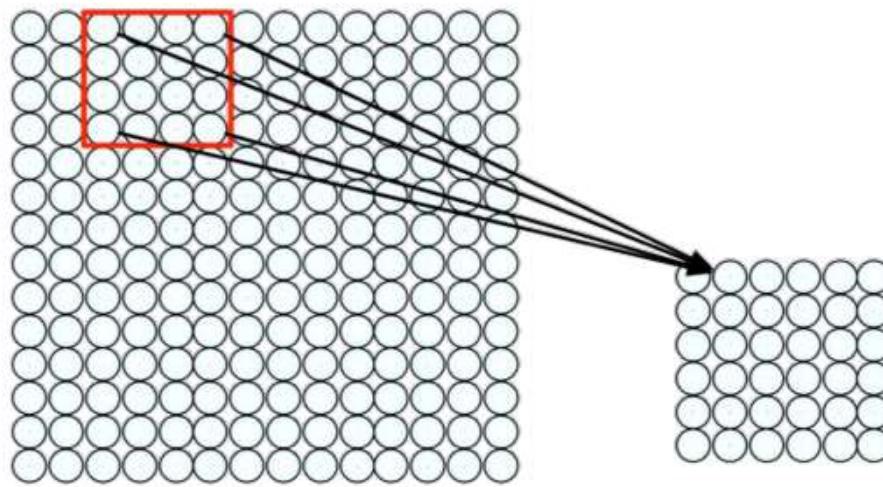
Eyes, ears, nose

High level features



Facial structure

Feature Extraction with Convolution

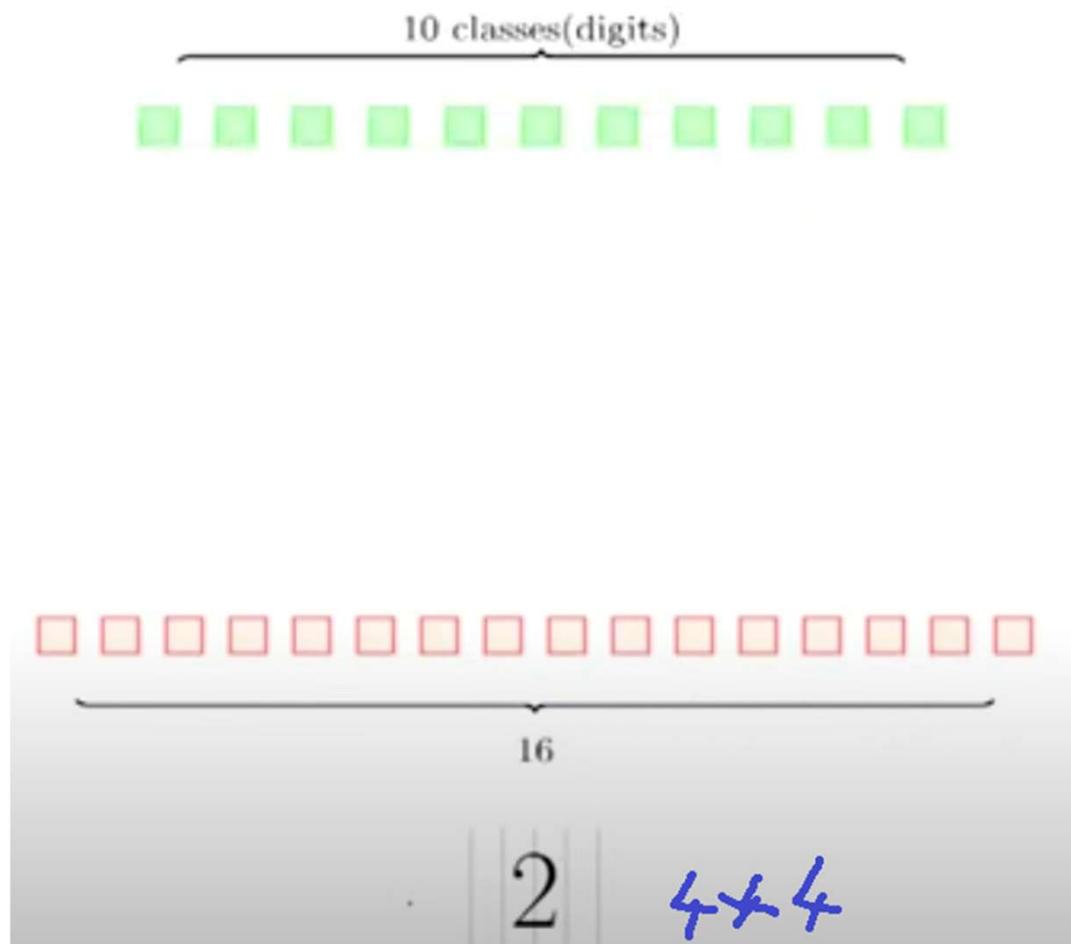


- 1) Apply a set of weights – a filter – to extract **local features**
- 2) Use **multiple filters** to extract different features
- 3) **Spatially share** parameters of each filter

How CNN different from FFN?

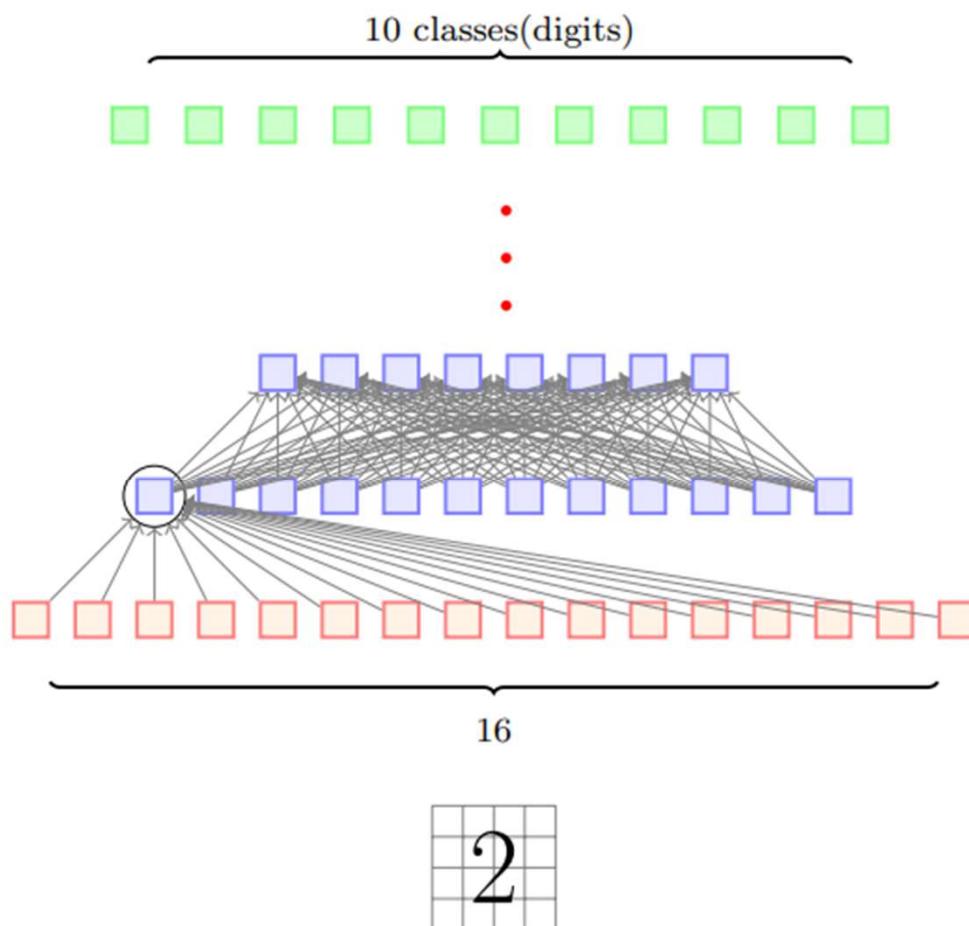
- Sparse Connections
- Weight Sharing

FFN for digit classification

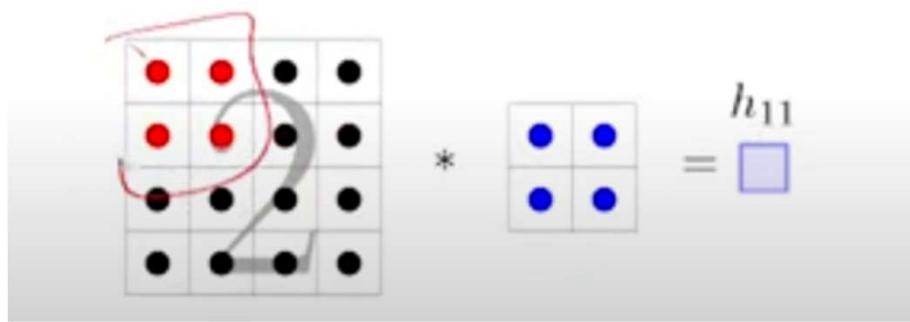
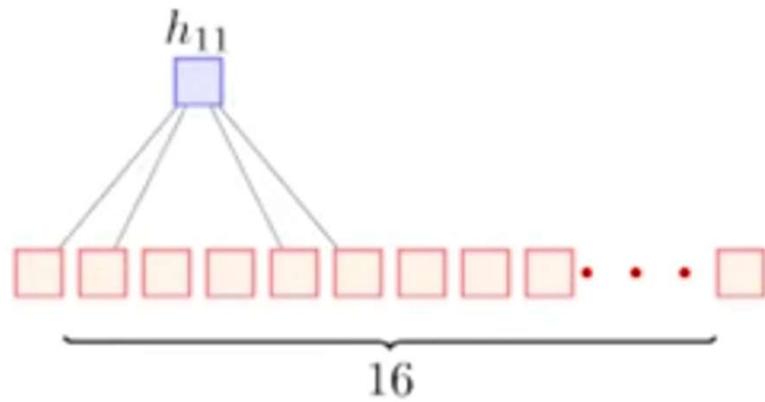


The MNIST dataset consists of **70,000 grayscale images** of handwritten digits,
Training Set: 60,000 images (digits 0–9)
•**Test Set:** 10,000 images (digits 0–9)
•**Classes:** 10 (one for each digit from 0 to 9)

FFN for digit classification

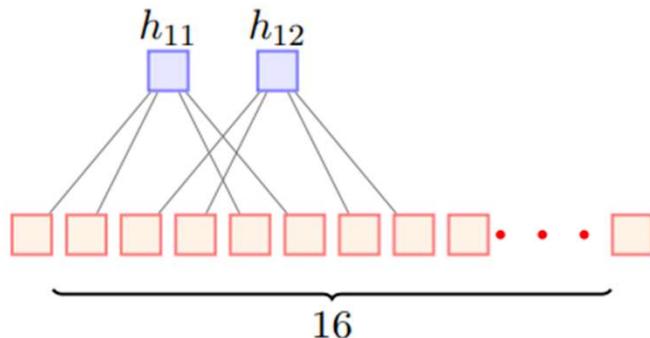


CNN



- Only a few local neurons participate in the computation of h_{11}
- For example, only pixels 1, 2, 5, 6 contribute to h_{11}
- The connections are much sparser

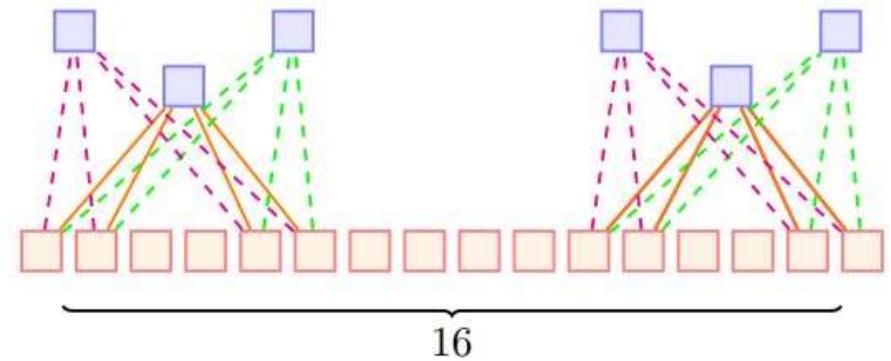
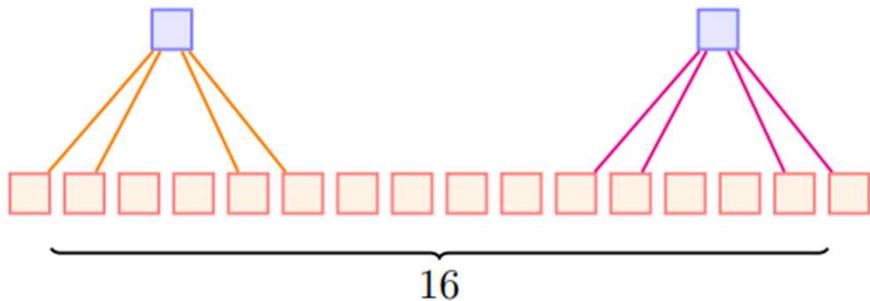
CNN – Sparse Connectivity



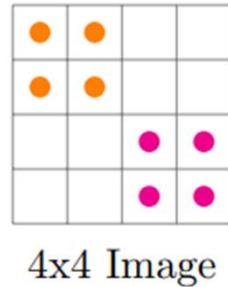
The diagram illustrates the convolutional operation $\text{Input} * \text{Kernel} = \text{Output}$. An input image of size 4x4 with red dots at positions (1,1), (1,2), (2,1), and (2,2) is multiplied by a 3x3 kernel with blue dots at positions (1,1), (1,2), and (2,1). The result is an output neuron h_{13} , represented by a blue square.

- Only a few local neurons participate in the computation of h_{11}
- For example, only pixels 1, 2, 5, 6 contribute to h_{11}
- The connections are much sparser
- We are taking advantage of the structure of the image(interactions between neighboring pixels are more interesting)
- This **sparse connectivity** reduces the number of parameters in the model

CNN – Weight Sharing



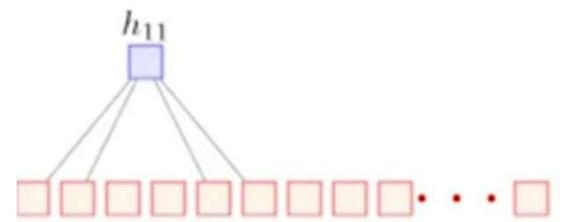
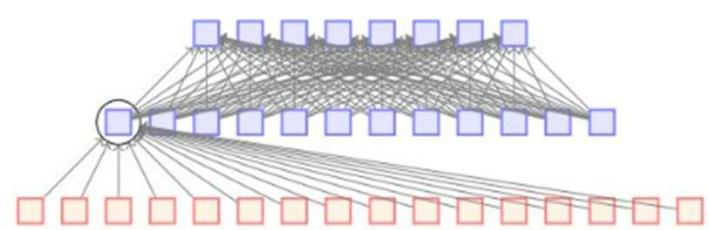
- Kernel 1
- Kernel 2



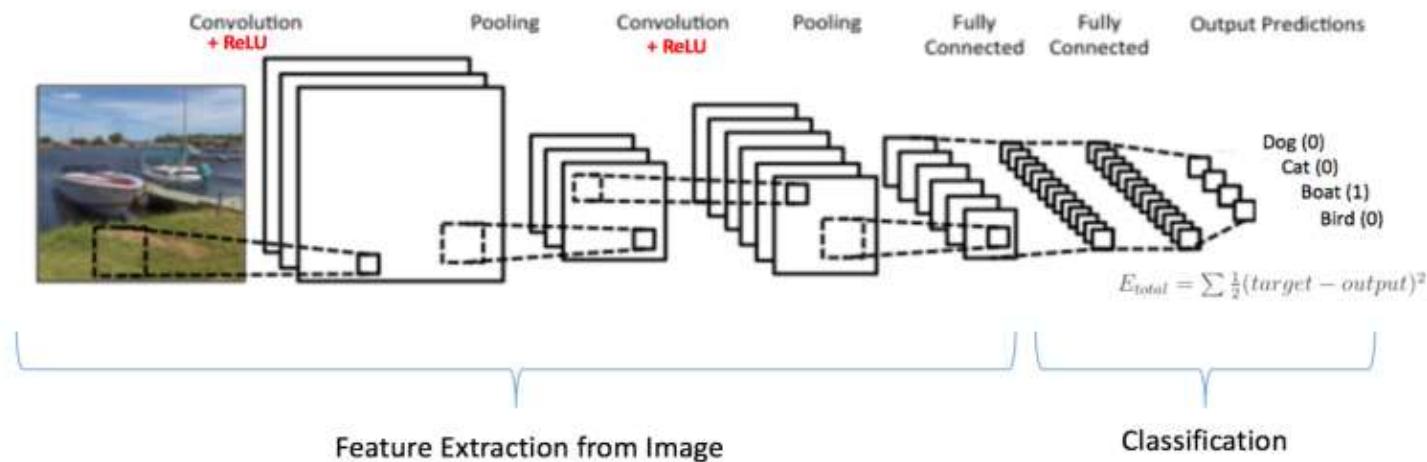
Each convolutional layer applies filters (or kernels) that slide over the input image to detect features, such as edges, textures, or shapes. The same set of weights is used across the entire spatial region of the input for each filter.

Weight Sharing Edge Detection Example

- In a fully connected layer:
 - Every neuron is connected to every input pixel in the image. This means there's a separate weight for each connection between the neuron and the pixel, with no local pattern focus.
- In a convolutional layer:
 - each filter has a small, fixed set of weights that it uses to look for a specific pattern—like an edge—in every part of the image.
- To detect an edge in a 32x32 pixel image, Fully Connected Layer: it would need 1,024 unique weights, one for each pixel.
- CNN with a 3x3 Filter: 3x3 edge-detection filter that is shared across all 1,024 possible 3x3 locations in the image. This requires only 9 weights,



An image classification CNN



Input Layer: Receives the image data

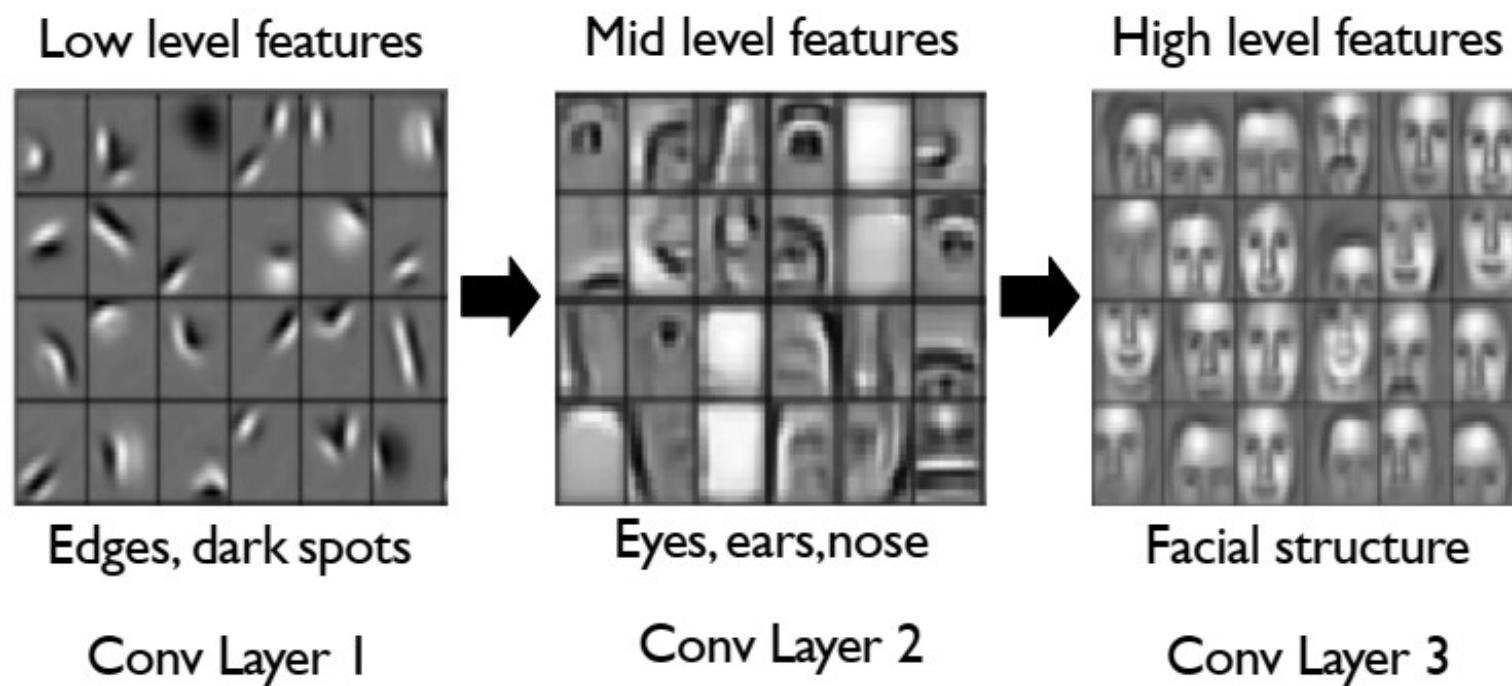
Convolutional Layers: Extracts local features from the input image by applying convolutional filters (kernels) that detect patterns such as edges, textures, and shapes.

Pooling Layers : Reduces the spatial dimensions (height and width) of the feature maps while retaining important features, thus reducing computational load and adding translational invariance.

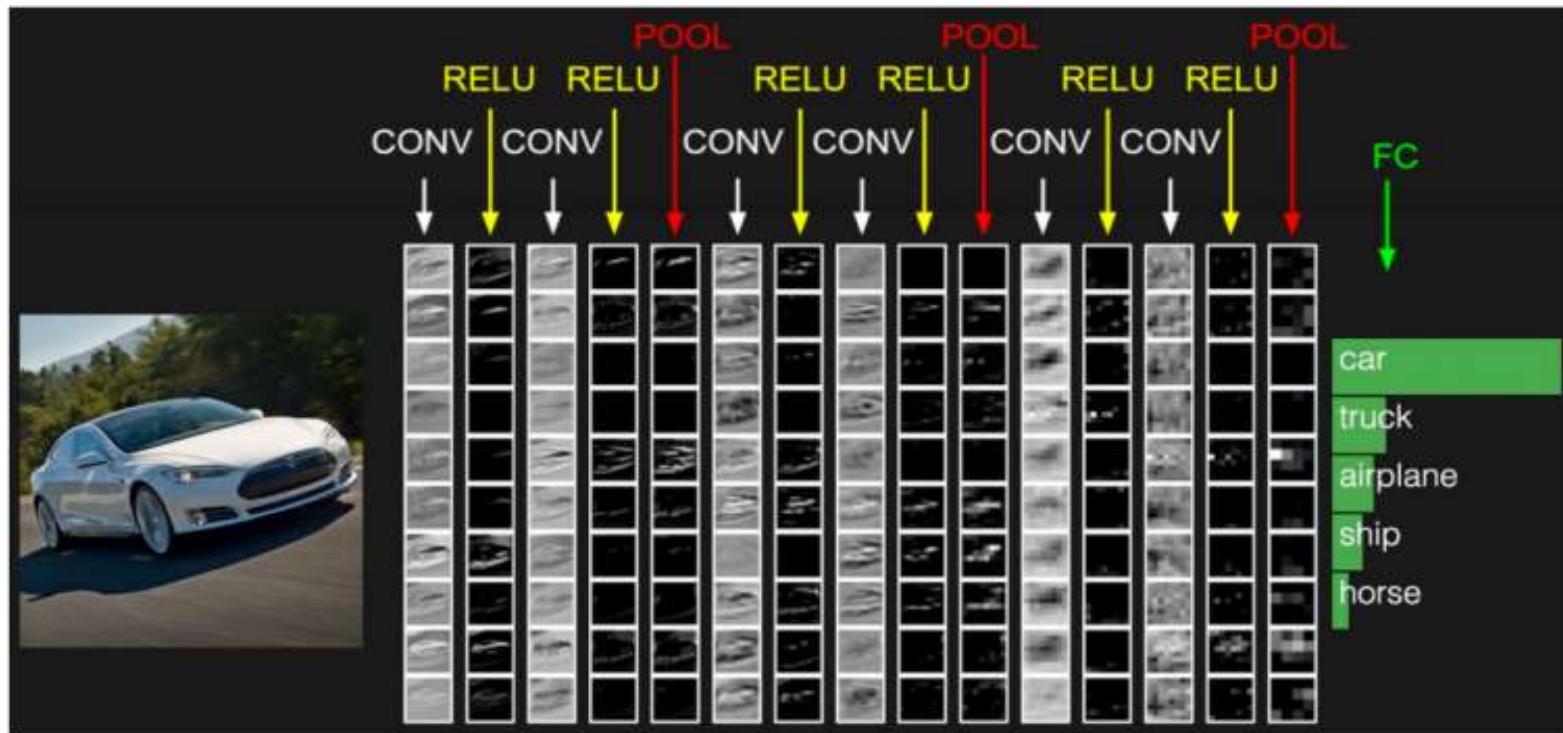
Flattening Layer: Converts the 2D feature maps produced by the final pooling layer into a 1D vector, making it compatible with fully connected layers.

Fully Connected Layers: Combines the features extracted by convolutional layers to perform the final classification.

Representation Learning in Deep CNNs

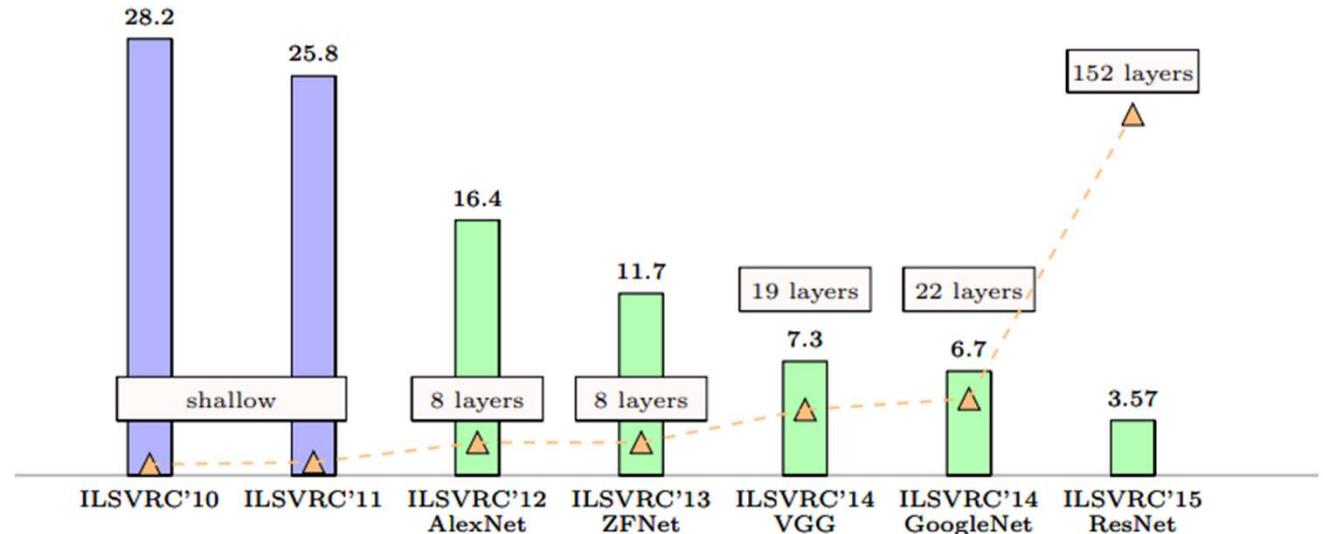


Example – Six convolutional layers



Successful Architectures

- LeNet-5
- AlexNet
- ZFNet
- VGGNet
- GoogLeNet
- ResNet



LeNet-5

- *Gradient Based Learning Applied To Document Recognition - Y. Lecun, L. Bottou, Y. Bengio, P. Haffner; 1998*
- Helped establish how we use CNNs today
- Replaced manual feature extraction

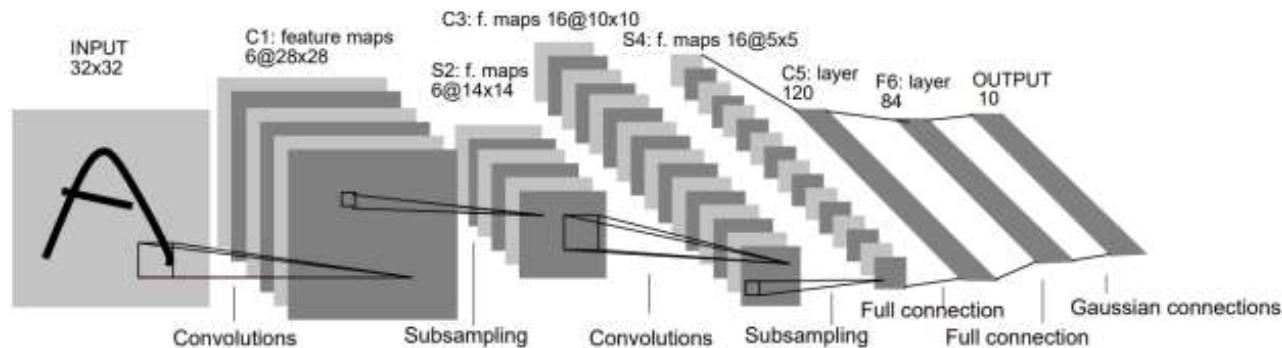
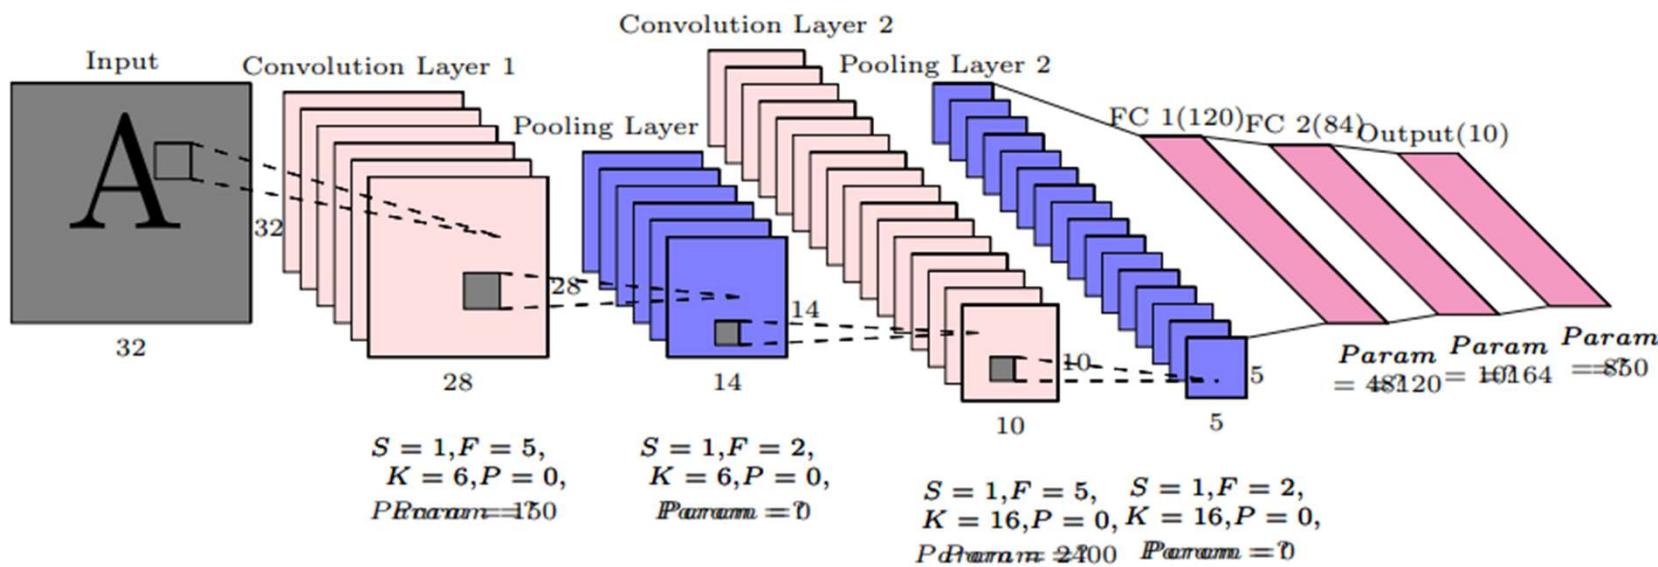


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

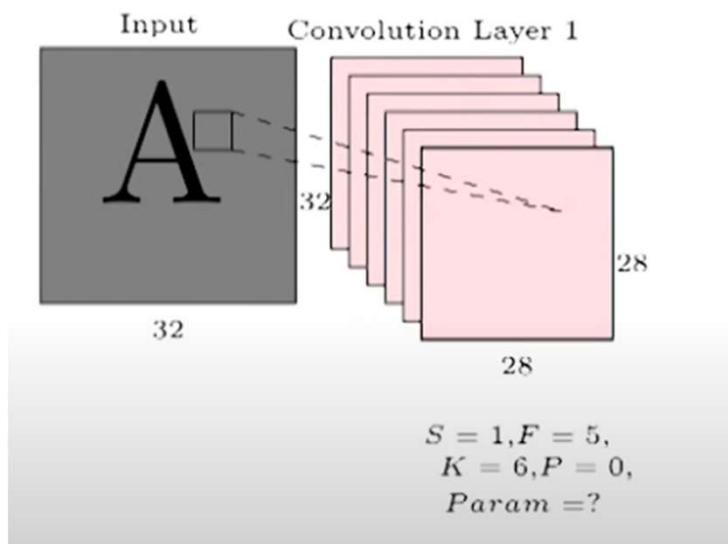
[LeCun et al., 1998]

LeNet-5



LeNet-5 consists of **seven layers**, with the layers responsible for learning hierarchical features of the input image

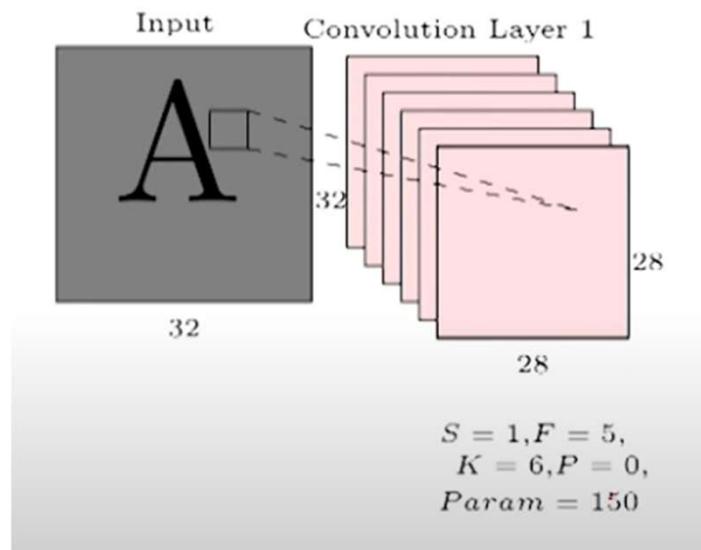
LeNet-5



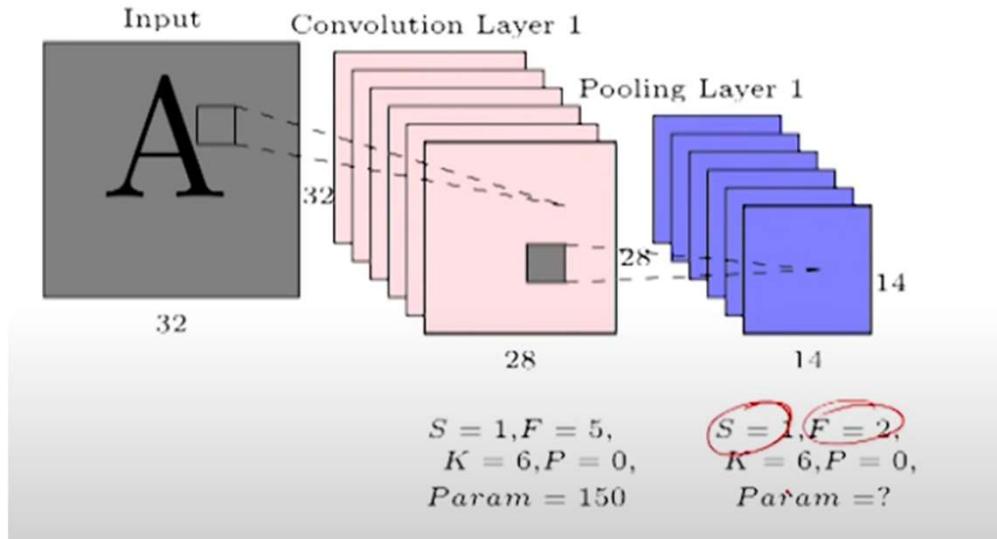
$$\frac{W_1 - F + 2P}{S} + 1$$

The number of filters K
The spatial extent F of each filter
S stride
P padding

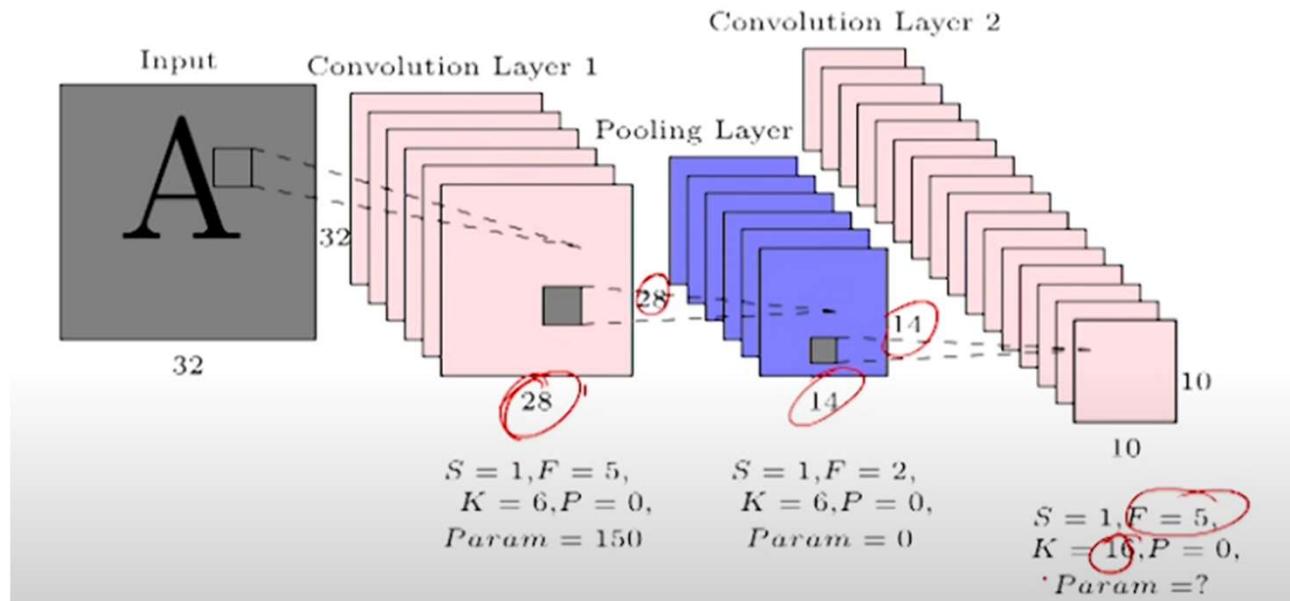
LeNet-5



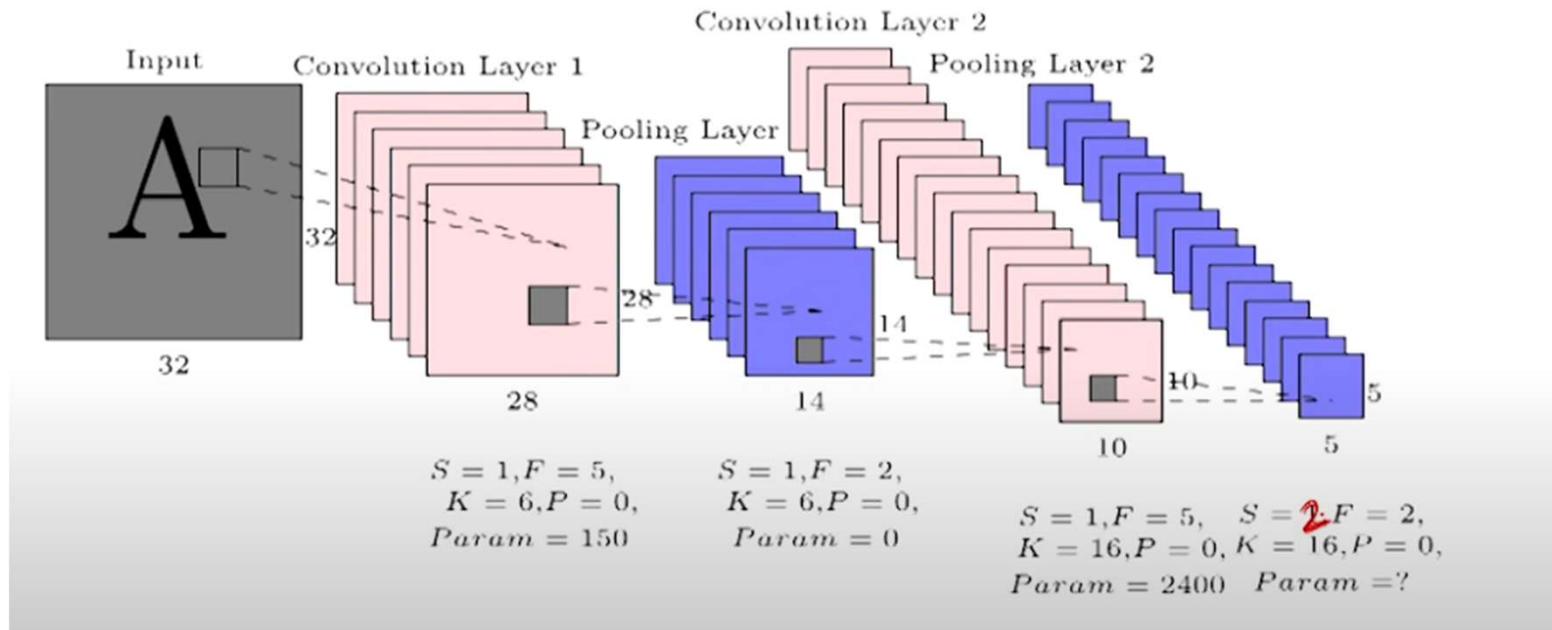
LeNet-5



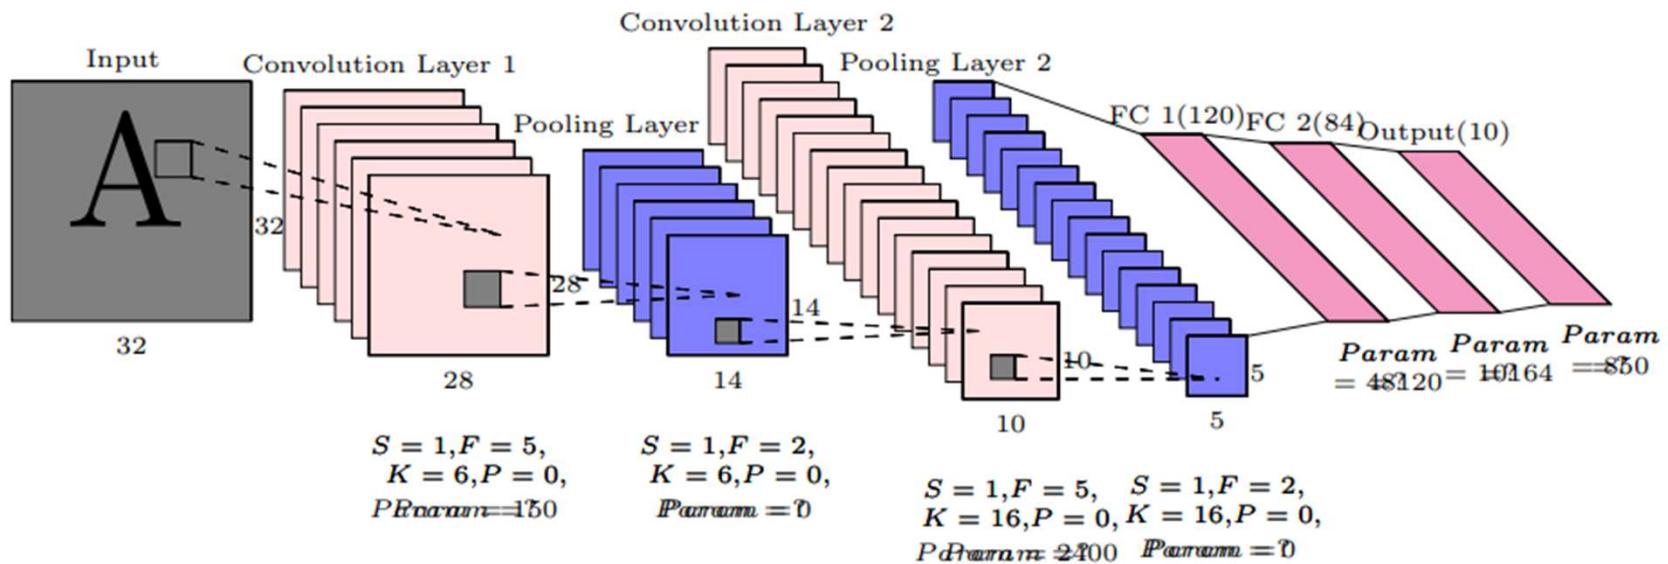
LeNet-5



LeNet-5



LeNet-5



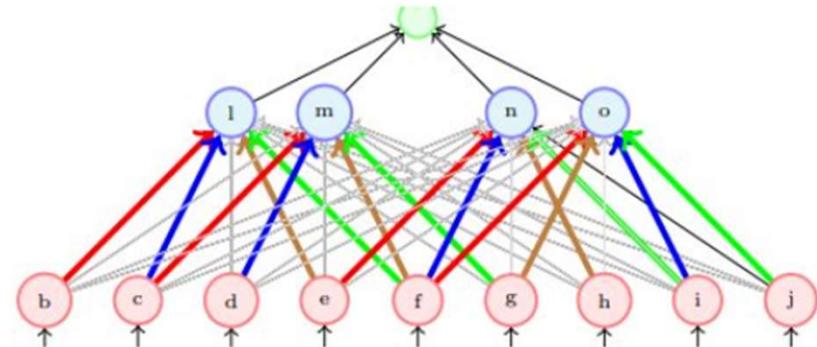
LeNet-5 Summary

Layer	Type	Number of Filters	Filter Size	Stride	Output Size	Parameters
Input	-	-	-	-	32×32	0
C1	Convolutional	6	5×5	1	$28 \times 28 \times 6$	$6 \times (5 \times 5) = 150$
S2	Average Pooling	-	2×2	2	$14 \times 14 \times 6$	0
C3	Convolutional	16	$5 \times 5 \times 6$	1	$10 \times 10 \times 16$	$16 \times (5 \times 5 \times 6) = 2,400$
S4	Average Pooling	-	2×2	2	$5 \times 5 \times 16$	0
C5	Fully Connected Conv.	120	$5 \times 5 \times 16$	1	$1 \times 1 \times 120$	$120 \times (5 \times 5 \times 16) = 48,000$
F6	Fully Connected	84	-	-	84	$84 \times 120 = 10,080$
Output	Fully Connected	10	-	-	10	$10 \times 84 = 840$

How do we train a CNN?

$$\begin{array}{|c|c|} \hline O_{11} & O_{12} \\ \hline O_{21} & O_{22} \\ \hline \end{array} = \text{Convolution} \left(\begin{array}{|c|c|c|} \hline X_{11} & X_{12} & X_{13} \\ \hline X_{21} & X_{22} & X_{23} \\ \hline X_{31} & X_{32} & X_{33} \\ \hline \end{array}, \begin{array}{|c|c|} \hline F_{11} & F_{12} \\ \hline F_{21} & F_{22} \\ \hline \end{array} \right)$$

$$\begin{array}{|c|c|} \hline \partial E / \partial F_{11} & \partial E / \partial F_{12} \\ \hline \partial E / \partial F_{21} & \partial E / \partial F_{22} \\ \hline \end{array} = \text{Convolution} \left(\begin{array}{|c|c|c|} \hline X_{11} & X_{12} & X_{13} \\ \hline X_{21} & X_{22} & X_{23} \\ \hline X_{31} & X_{32} & X_{33} \\ \hline \end{array}, \begin{array}{|c|c|} \hline \partial E / \partial O_{11} & \partial E / \partial O_{12} \\ \hline \partial E / \partial O_{21} & \partial E / \partial O_{22} \\ \hline \end{array} \right)$$

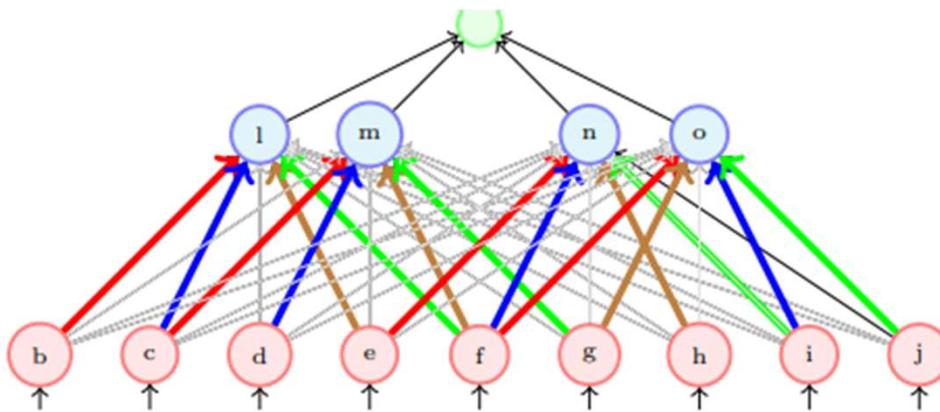


Slide taken from Forward And Backpropagation in Convolutional Neural Network. - Medium

How do we train a CNN?

Backpropagation

- **In PyTorch:** The nn.Conv2d layer performs 2D convolutions.
- **In TensorFlow/Keras:** The Conv2D layer performs 2D convolutions.

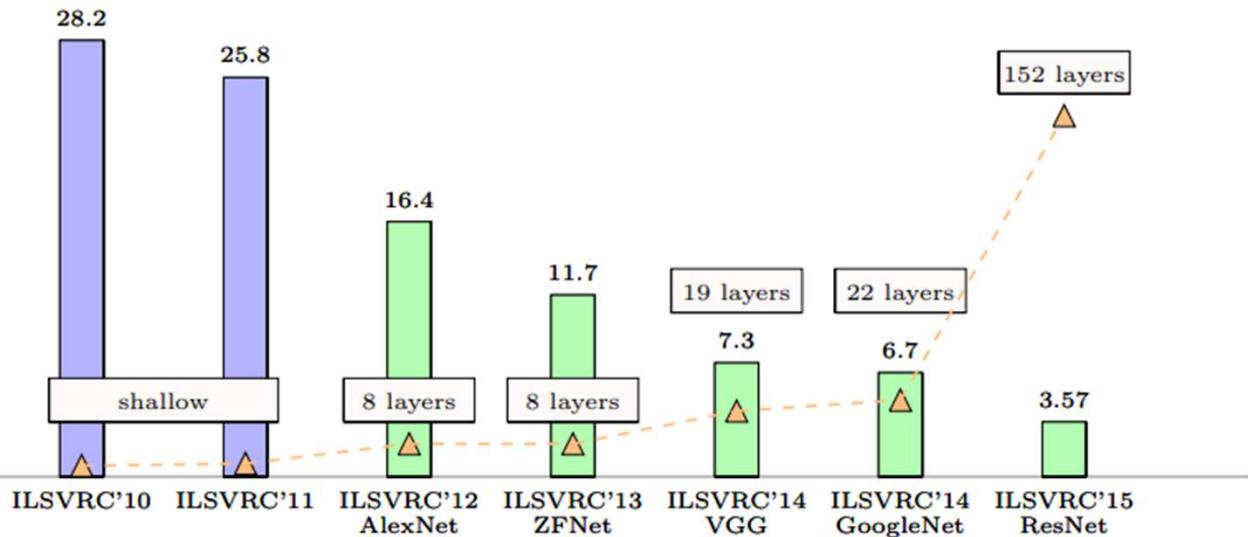


Packages like **PyTorch**, **TensorFlow**, and **Keras** simplify the process of convolution, gradient computation, and weight updates. They provide high-level abstractions, automatic differentiation, optimized performance, and scalability, allowing you to focus on model design and experimentation

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners

- The annual “Olympics” of computer vision.
 - Teams from across the world compete to see who has the best computer vision model for tasks such as classification, localization, detection, and more.
 - **2012 marked the first year where a CNN was used** to achieve a top 5 test error rate of 15.3%.
 - The next best entry achieved an error of 26.2%.
- **ImageNet** is a database of millions of labeled images
 - created by **Stanford's Vision Lab** under the leadership of **Fei-Fei Li** in 2009.
 - ImageNet contains over 14 million labeled images.
 - **Categories:** The dataset is organized into more than 20,000 categories .

ImageNet Solutions



- **Size:** Over 14 million images across more than 20,000 categories.

- **ImageNet Large Scale Visual Recognition Challenge (ILSVRC):** An annual competition that ran from 2010 to 2017, where researchers competed to create models that achieved the highest accuracy on the dataset. This challenge significantly advanced the development of deep learning architectures.

Architecture

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4

CONV5

Max POOL3

FC6

FC7

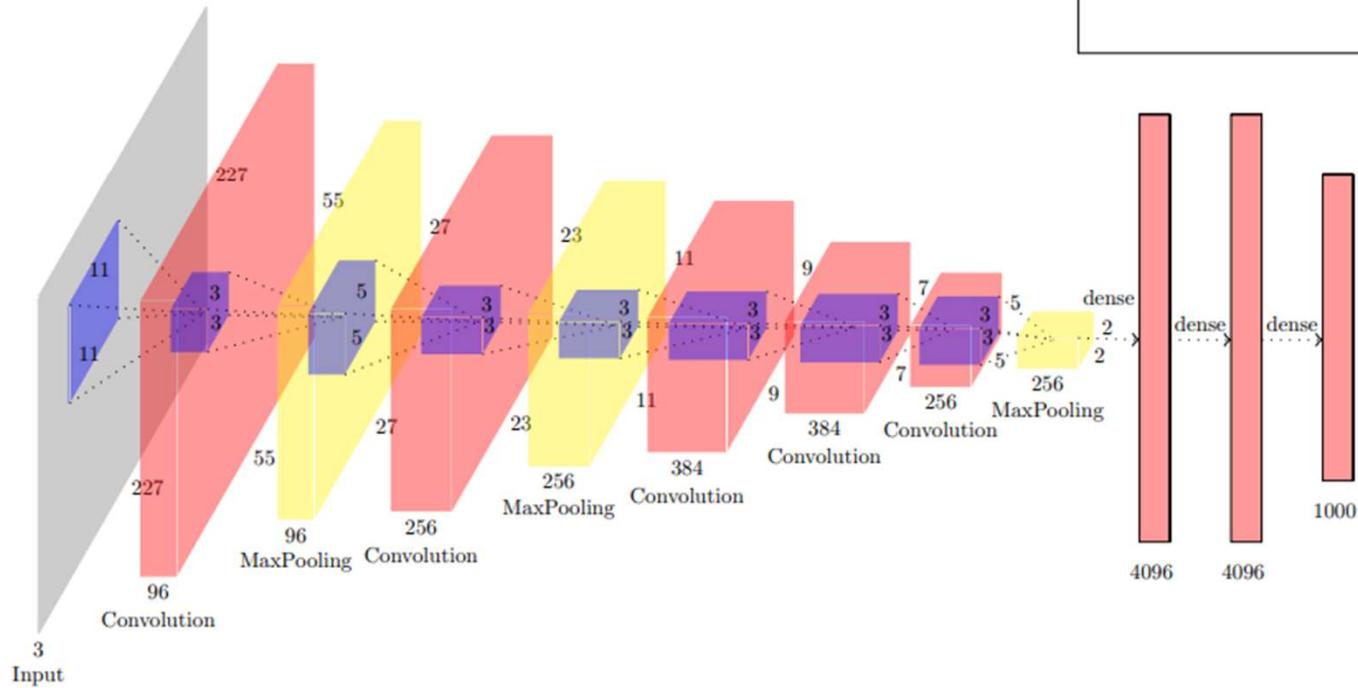
FC8

AlexNet

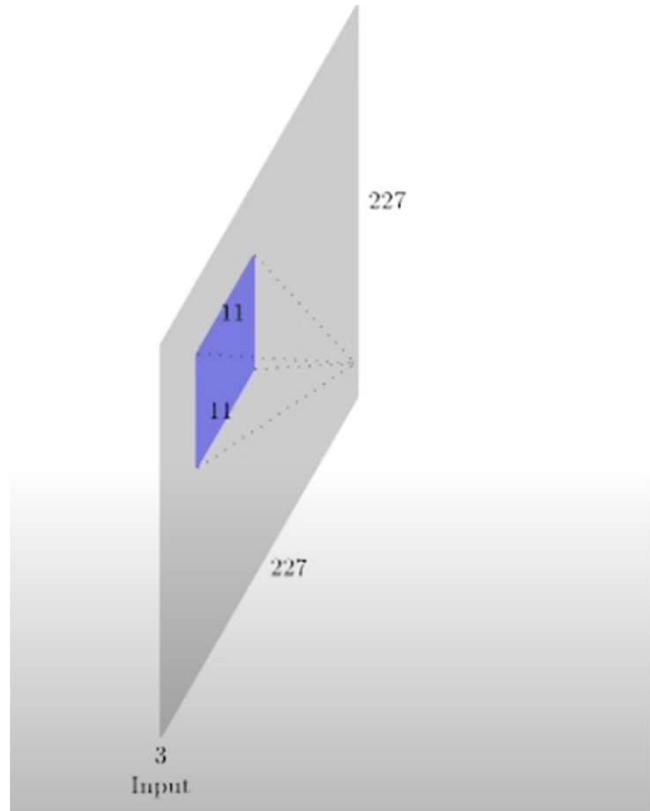
- Input: 227x227x3 images (224x224 before padding)
- First layer: 96 11x11 filters applied at stride 4
- **Output volume size?**
$$(N-F)/s+1 = (227-11)/4+1 = 55 \rightarrow [55x55x96]$$
- **Number of parameters in this layer?**
$$(11*11*3)*96 = 35K$$

AlexNet

Total Parameters: 27.55M

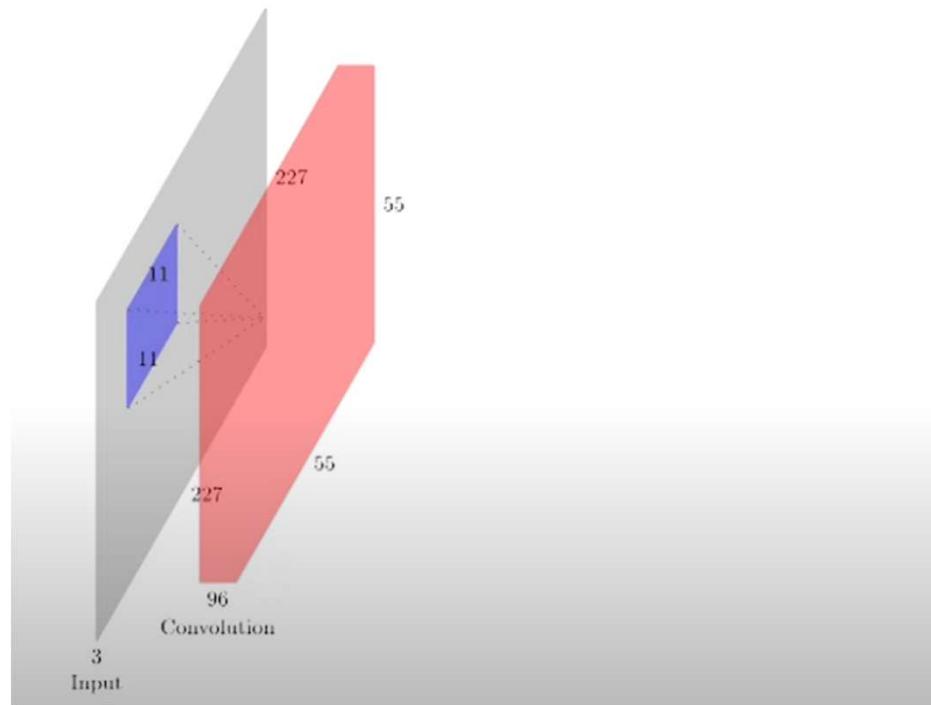


AlexNet



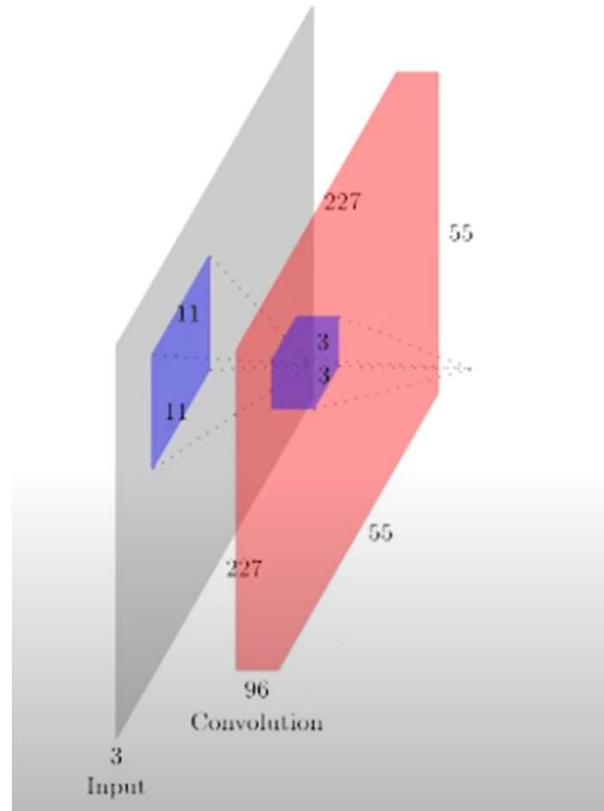
Input: $227 \times 227 \times 3$
Conv1: $K = 96, F = 11$
 $\underline{S = 4, P = 0}$
Output: $W_2 = ?, H_2 = ?$
Parameters: ?

AlexNet



Input: $227 \times 227 \times 3$
Conv1: $K = 96, F = 11$
 $S = 4, P = 0$
Output: $W_2 = 55, H_2 = 55$
Parameters: $(11 \times 11 \times 3) \times 96 = 34K$

AlexNet



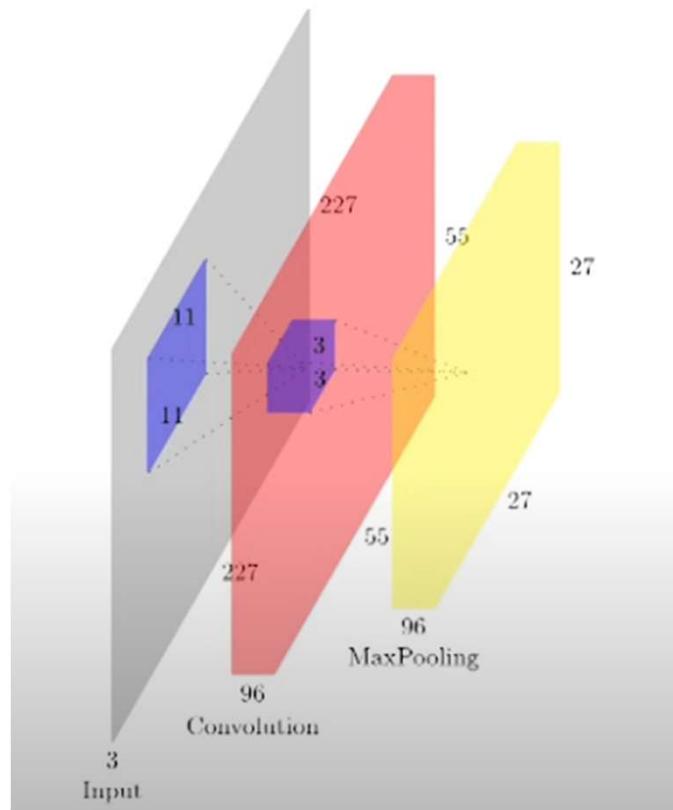
Max Pool Input: $55 \times 55 \times 96$

$$F = 3, S = 2$$

Output: $W_2 = ?, H_2 = ?$

Parameters: ?

AlexNet



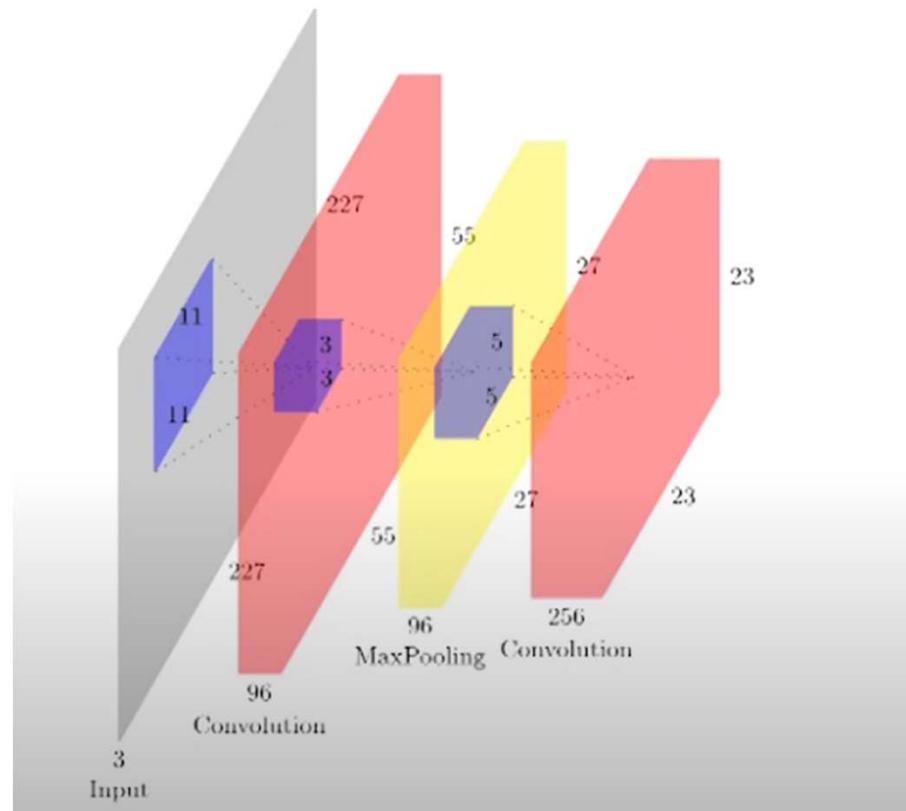
Max Pool Input: $55 \times 55 \times 96$

$$F = 3, S = 2$$

Output: $W_2 = 27, H_2 = 27$

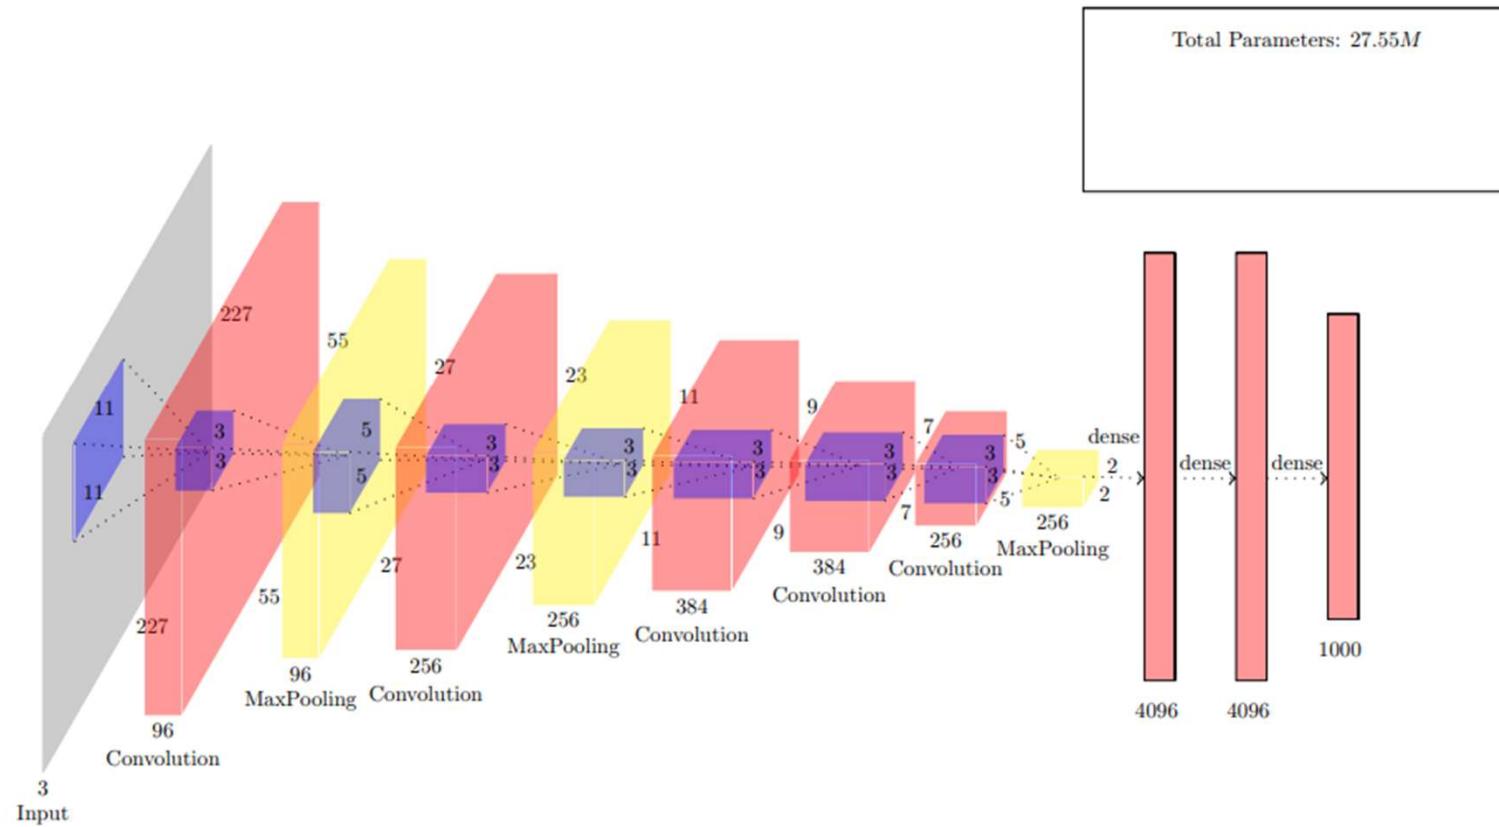
Parameters: 0

AlexNet



Input: $27 \times 27 \times 96$
Conv1: $K = 256, F = 5$
 $S = 1, P = 0$
Output: $W_2 = 23, H_2 = 23$
Parameters: $(5 \times 5 \times 96) \times 256 = 0.6M$

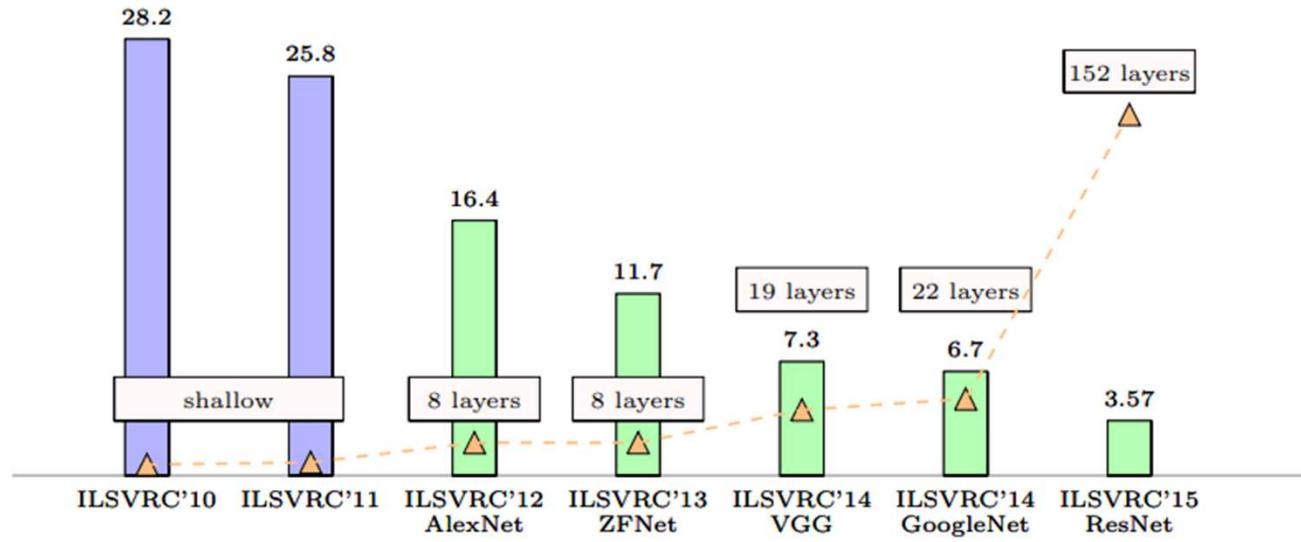
AlexNet



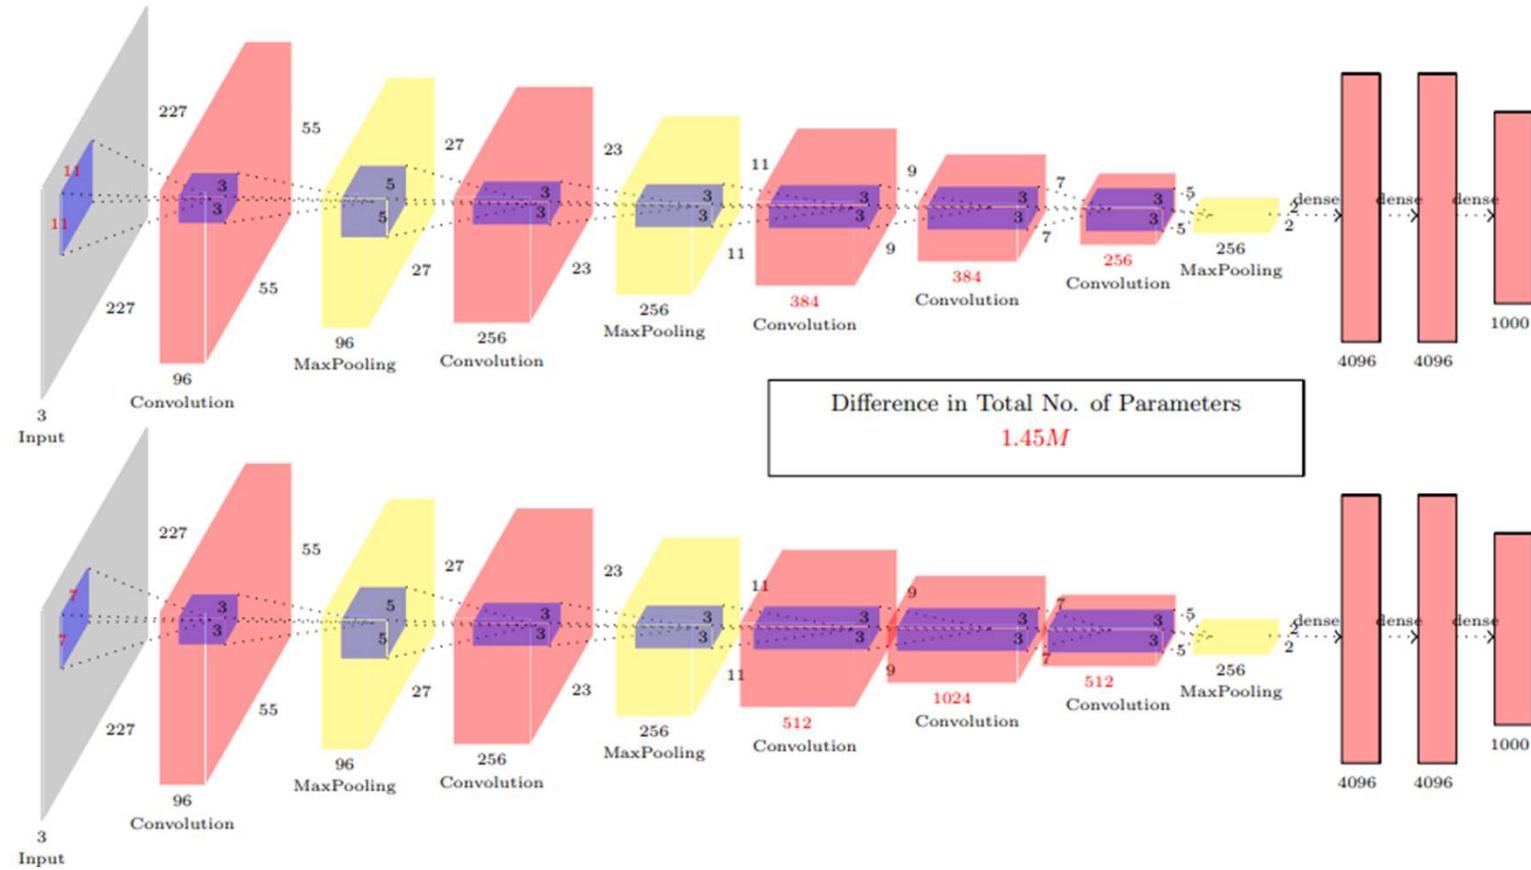
AlexNet Summary

Layer	Output Size	#Filters	F Size	Stride	Padding	Parameters
Input Layer	224 x 224 x 3	N/A	N/A	N/A	N/A	N/A
Conv1	55 x 55 x 96	96	11 x 11	4	0	34,944
Max Pool 1	27 x 27 x 96	N/A	3 x 3	2	0	0
Conv2	27 x 27 x 256	256	5 x 5	1	2	614,656
Max Pool 2	13 x 13 x 256	N/A	3 x 3	2	0	0
Conv3	13 x 13 x 384	384	3 x 3	1	1	884,736
Conv4	13 x 13 x 384	384	3 x 3	1	1	1,327,488
Conv5	13 x 13 x 256	256	3 x 3	1	1	884,992
Max Pool 3	6 x 6 x 256	N/A	3 x 3	2	0	0
Flatten	1 x 1 x 9216	N/A	N/A	N/A	N/A	0
FC1	1 x 1 x 4096	4096	N/A	N/A	N/A	37,752,832
FC2	1 x 1 x 4096	4096	N/A	N/A	N/A	16,781,312
FC3 (Output)	1 x 1 x 1000	1000	N/A	N/A	N/A	4,096,000

ImageNet Solutions



ZFNet

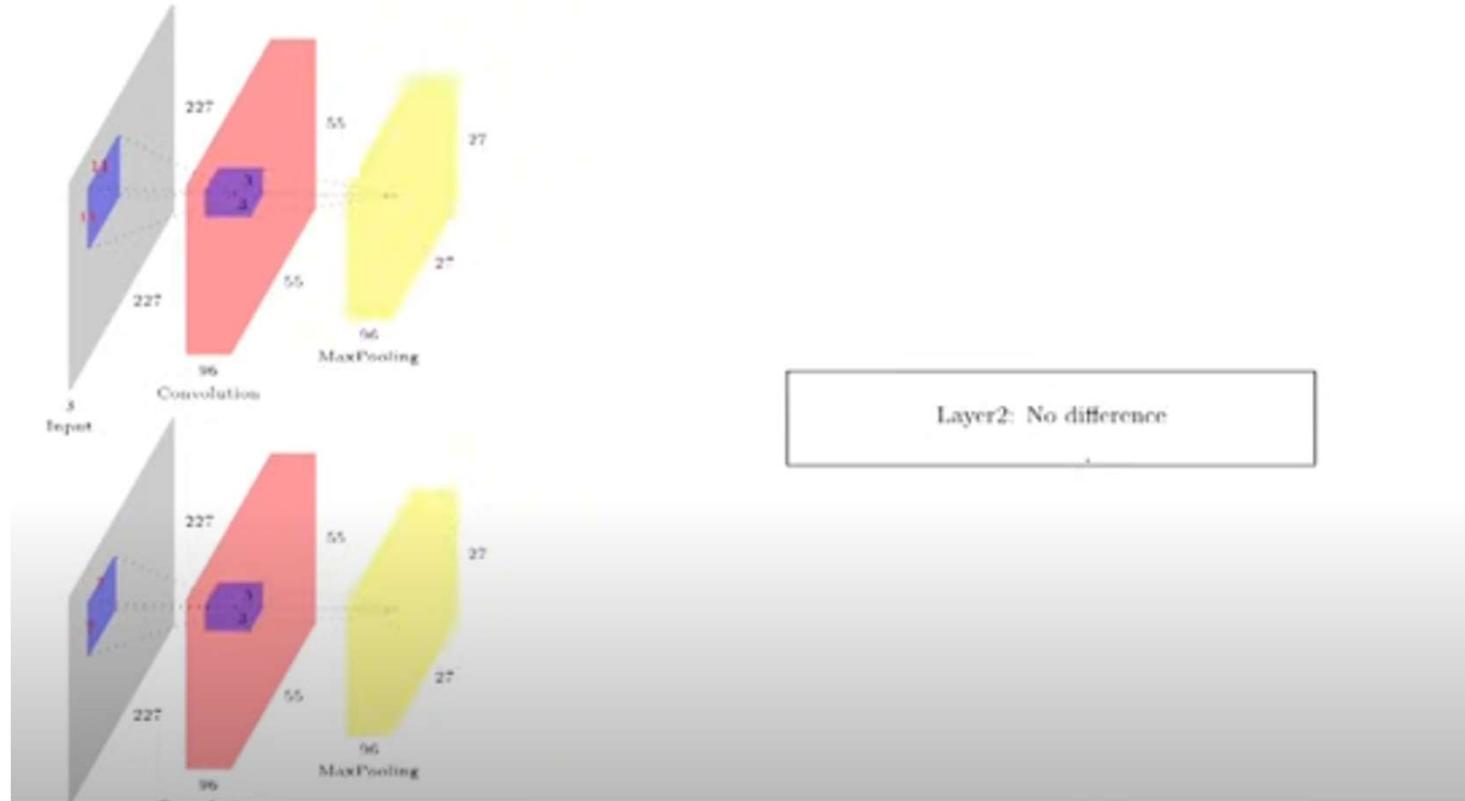


ZFNET vs AlexNet



Layer1: $F = 11 \rightarrow 7$
Difference in Parameters
 $((11^2 - 7^2) \times 3) \times 96 = 20.7K$

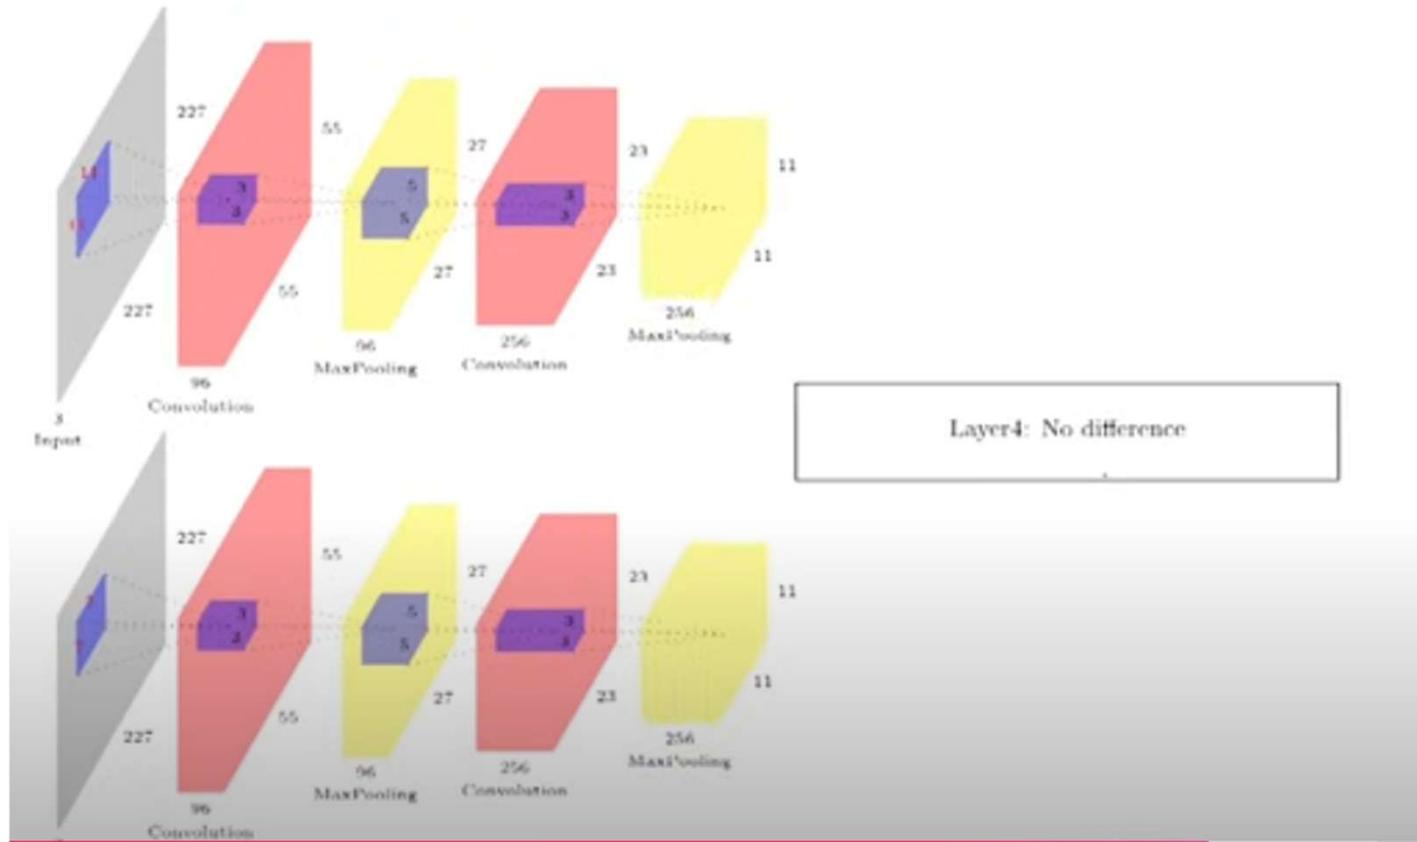
ZFNET vs AlexNet



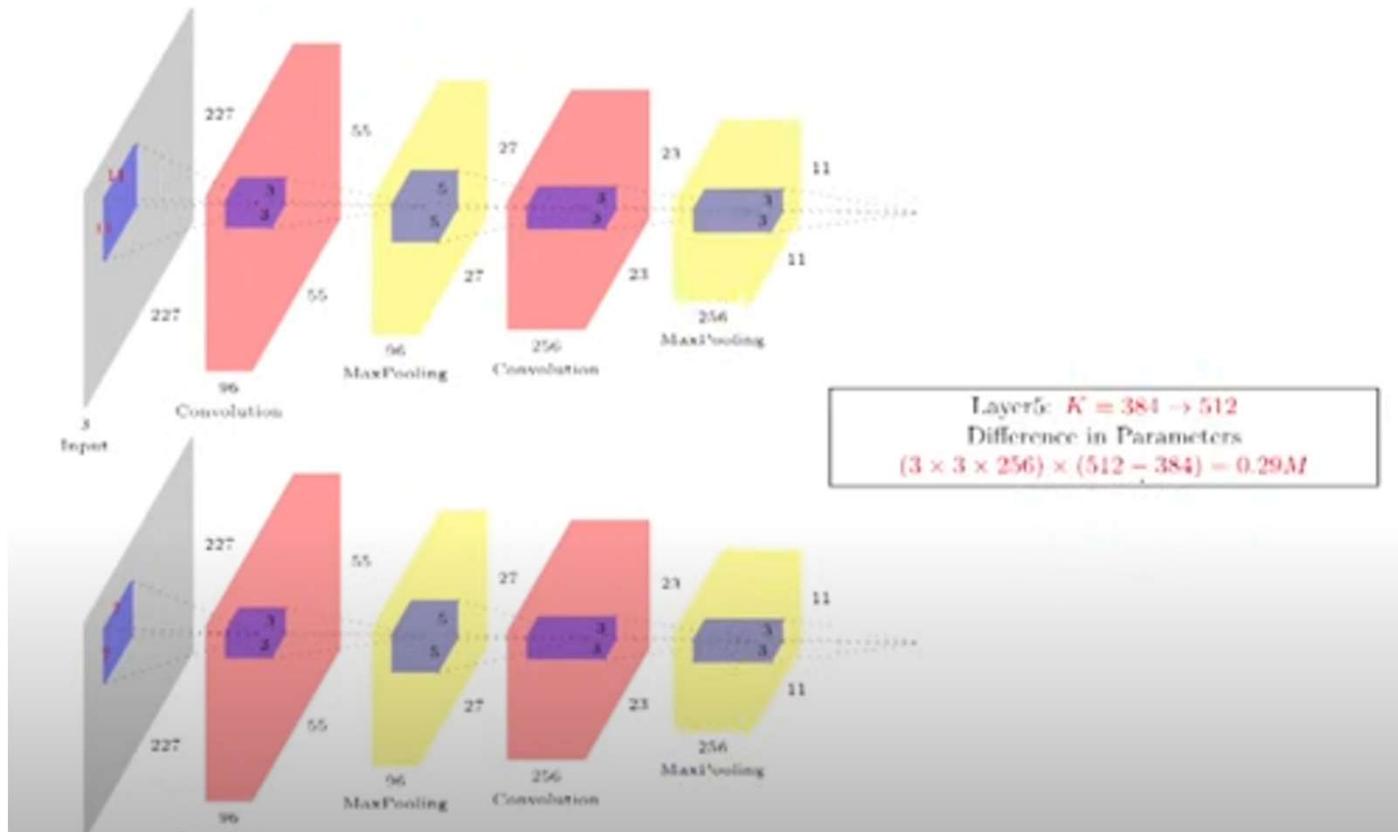
ZFNET vs AlexNet



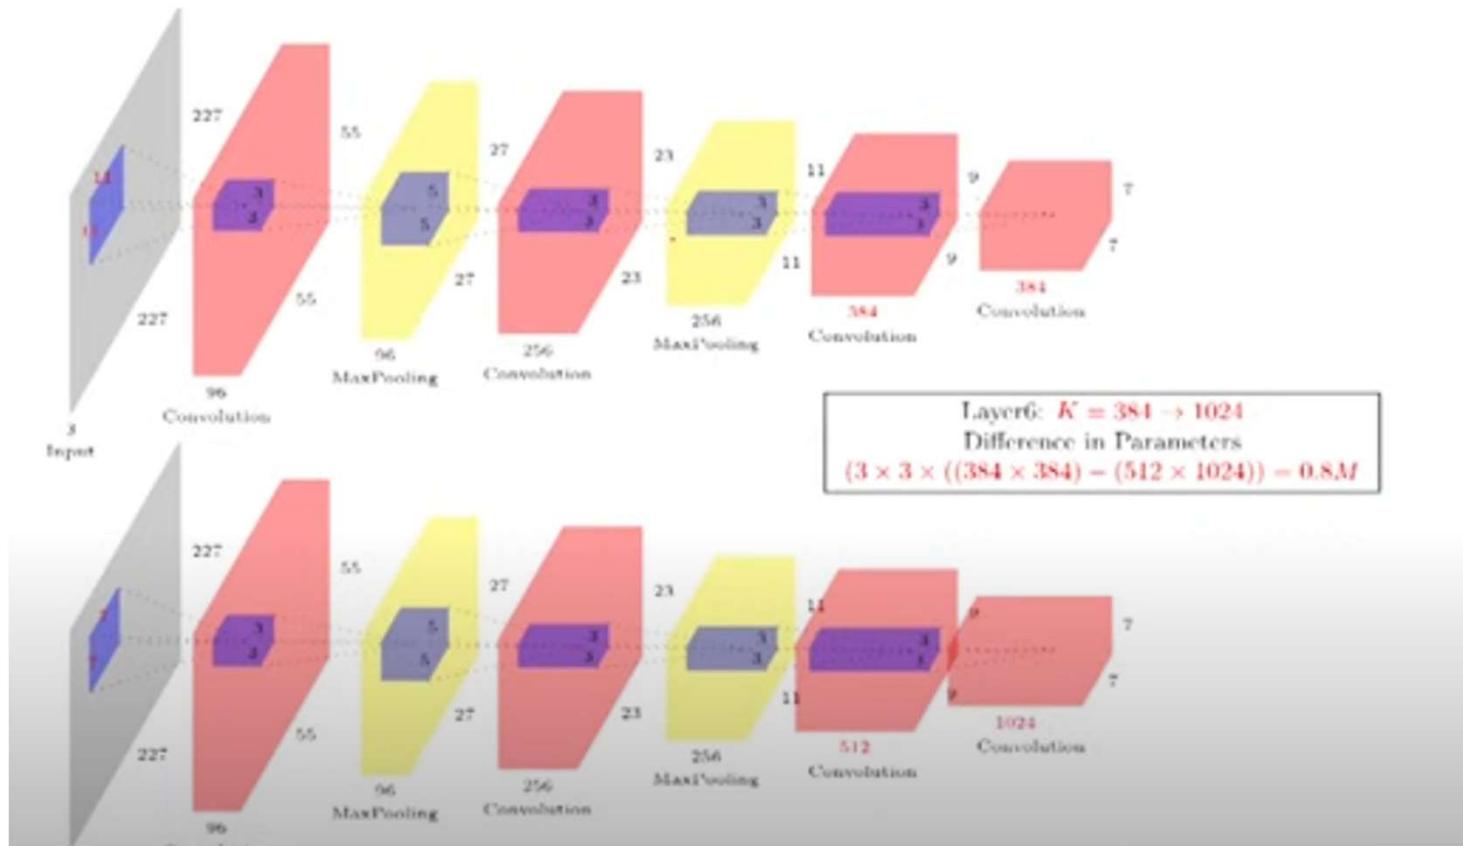
ZFNET vs AlexNet



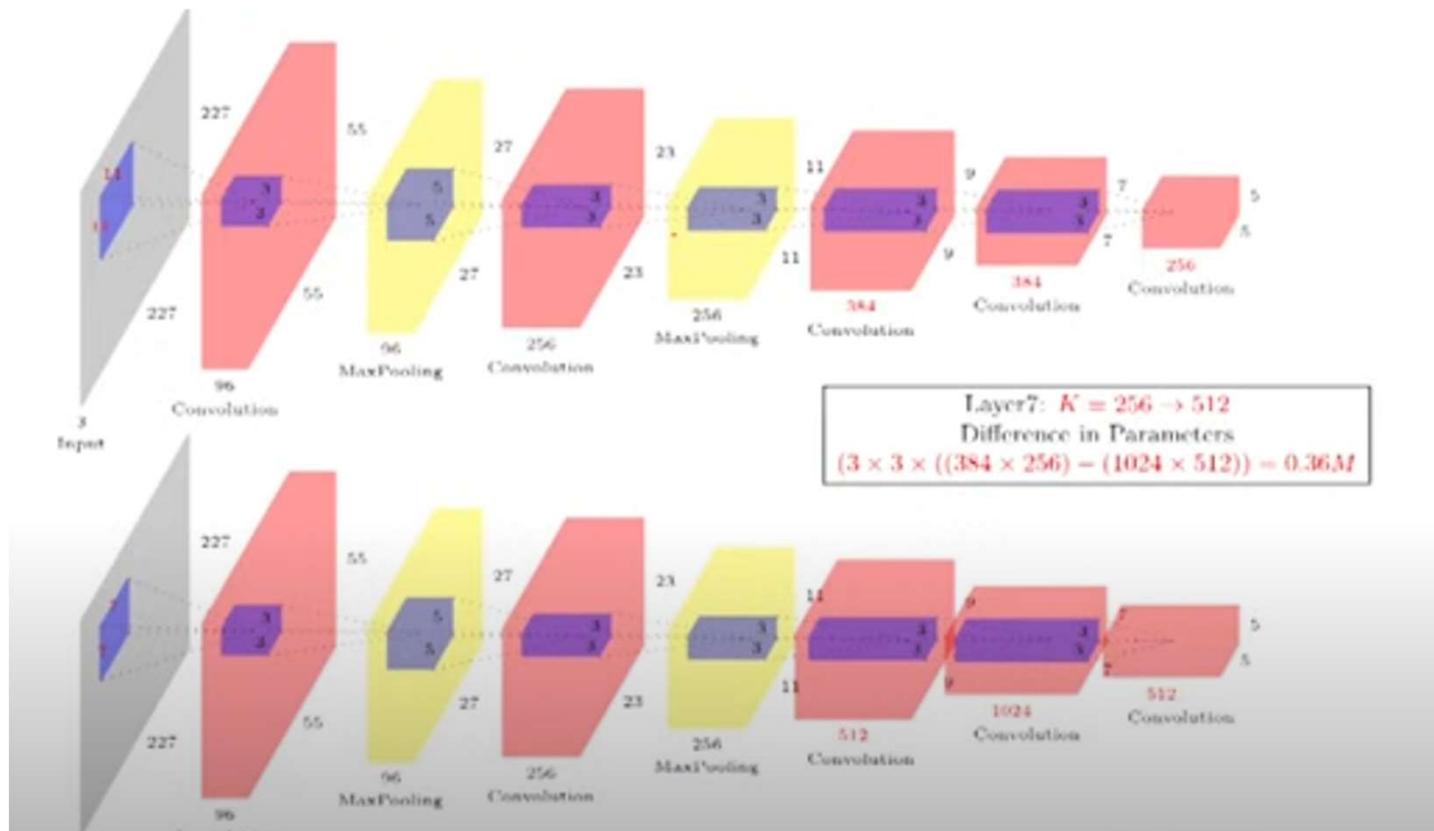
ZFNET vs AlexNet



ZFNET vs AlexNet



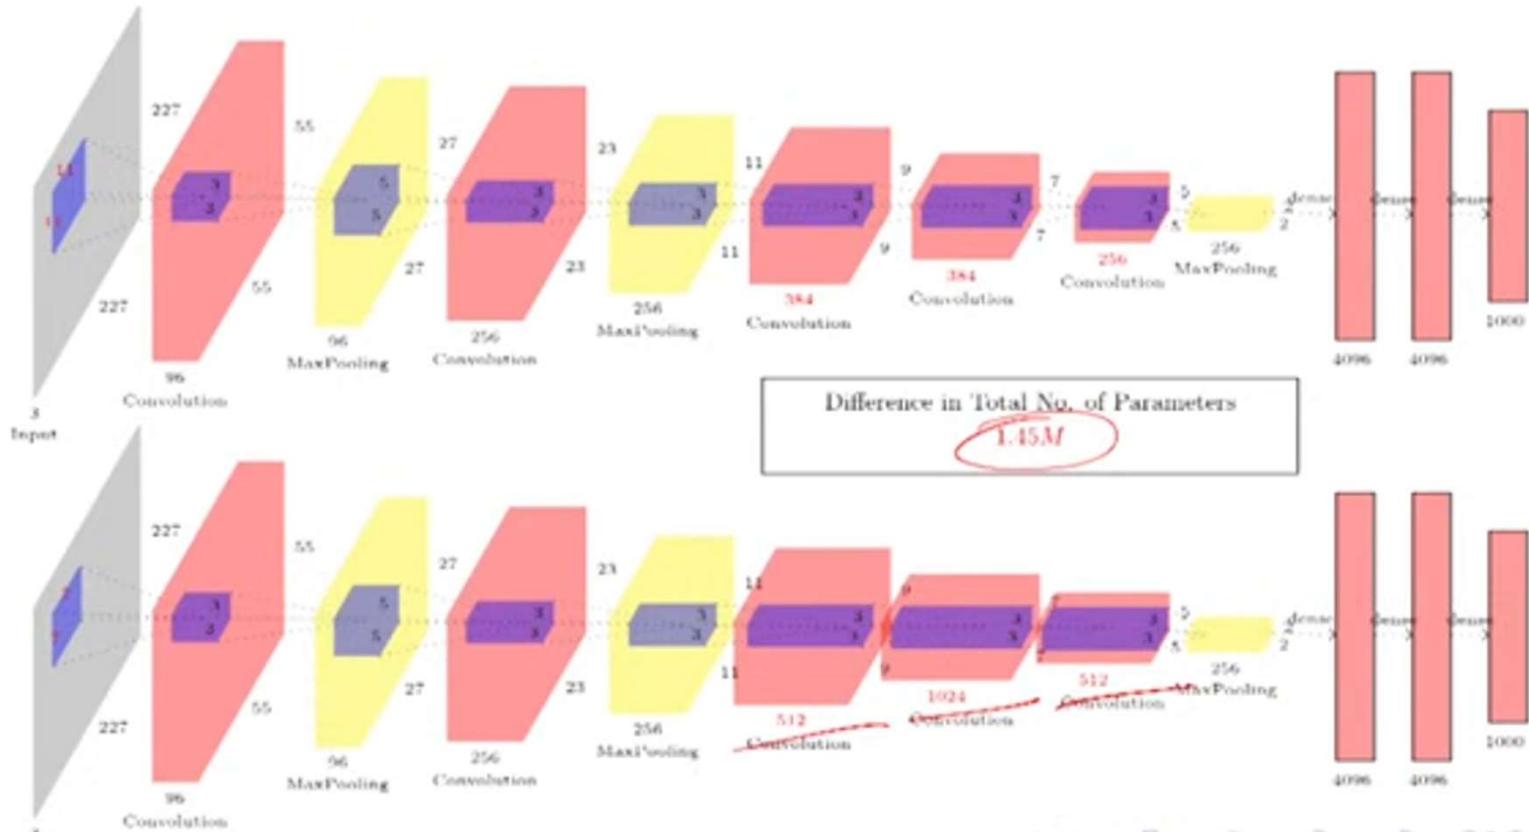
ZFNET vs AlexNet



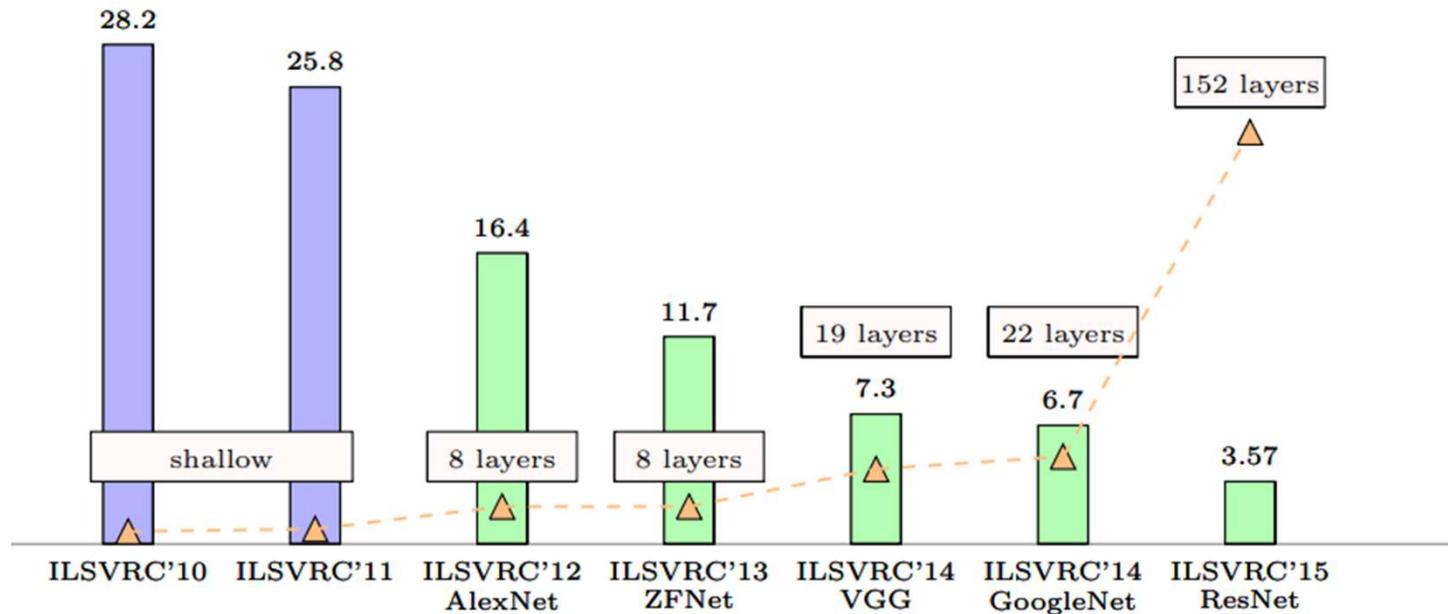
ZFNET vs AlexNet



ZFNET vs AlexNet



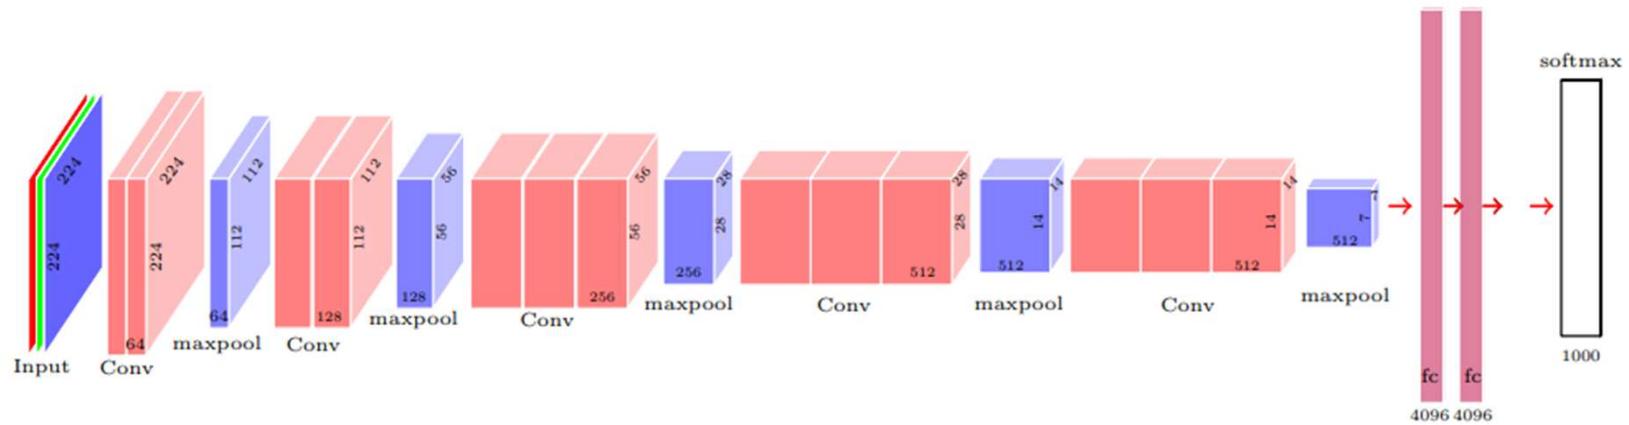
ImageNet Solutions



VGGNet

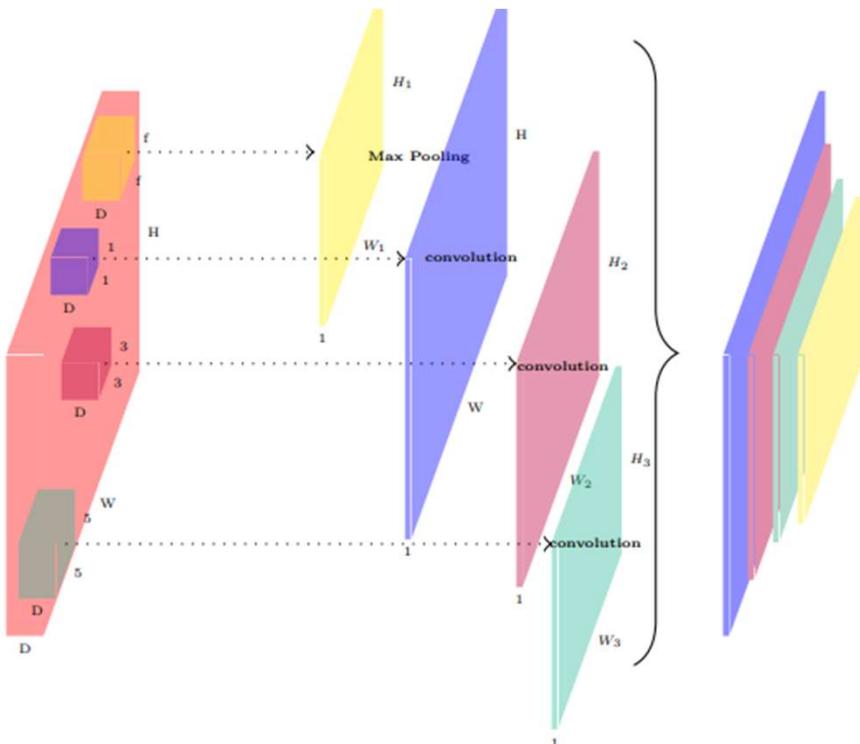
Visual Geometry Group (VGG) at the University of Oxford

16 Layers: VGG-16 has 16 weighted layers (13 convolutional layers and 3 fully connected layers), hence the name "VGG-16.



- Kernel size is 3×3 throughout
- Total parameters in non FC layers = $\sim 16M$
- Total Parameters in FC layers = $(512 \times 7 \times 7 \times 4096) + (4096 \times 4096) + (4096 \times 1024) = \sim 122M$
- Most parameters are in the first FC layer ($\sim 102M$)

GoogLeNet - Inception v1



- Consider the output at a certain layer of a convolutional neural network
- After this layer we could apply a max-pooling layer
- Or a 1×1 convolution
- Or a 3×3 convolution
- Or a 5×5 convolution
- **Question:** Why choose between these options (convolution, maxpooling, filter sizes)?
- **Idea:** Why not apply all of them at the same time and then concatenate the feature maps?

GoogLeNet

introduced a novel module called the **Inception module** to improve computational efficiency and reduce the model's number of parameters without compromising performance.

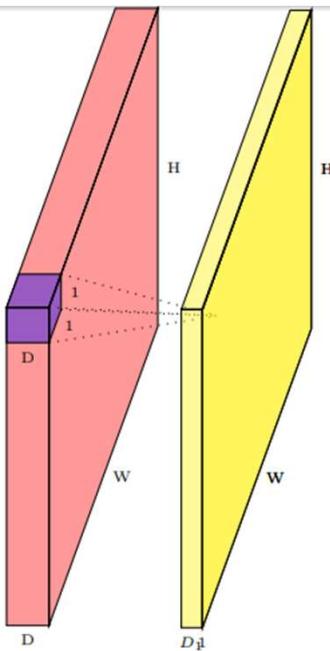
- **Inception Modules:**

- The core innovation in GoogLeNet is the Inception module, which processes input at multiple scales (1x1, 3x3, and 5x5 convolutions) and combines the outputs.
- Each module performs a mix of 1x1, 3x3, and 5x5 convolutions, along with 3x3 max-pooling, which are concatenated to capture both local and global features effectively.
- Using 1x1 convolutions reduces dimensionality and computational cost, making the network deeper without increasing the number of parameters excessively.

- **Depth and Architecture:**

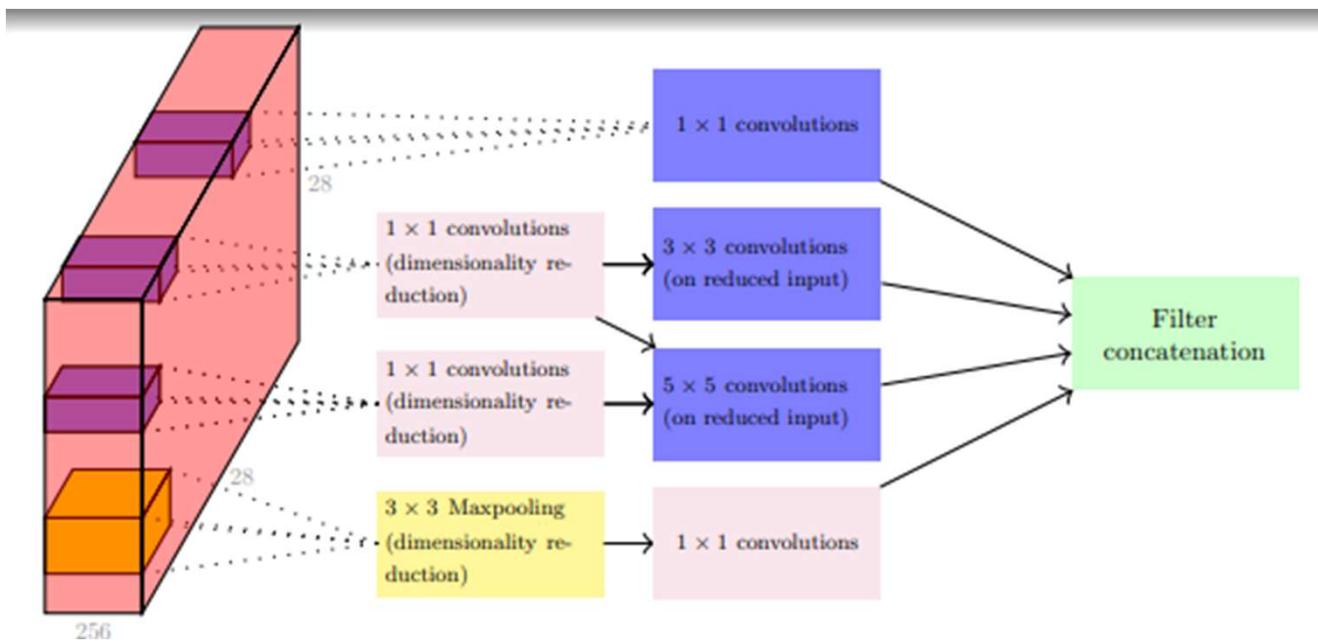
- GoogLeNet has **22 layers** (excluding pooling and softmax layers) with only **5 million parameters**, which is much smaller than VGG-16's 138 million parameters, despite being deeper.

GoogLeNet

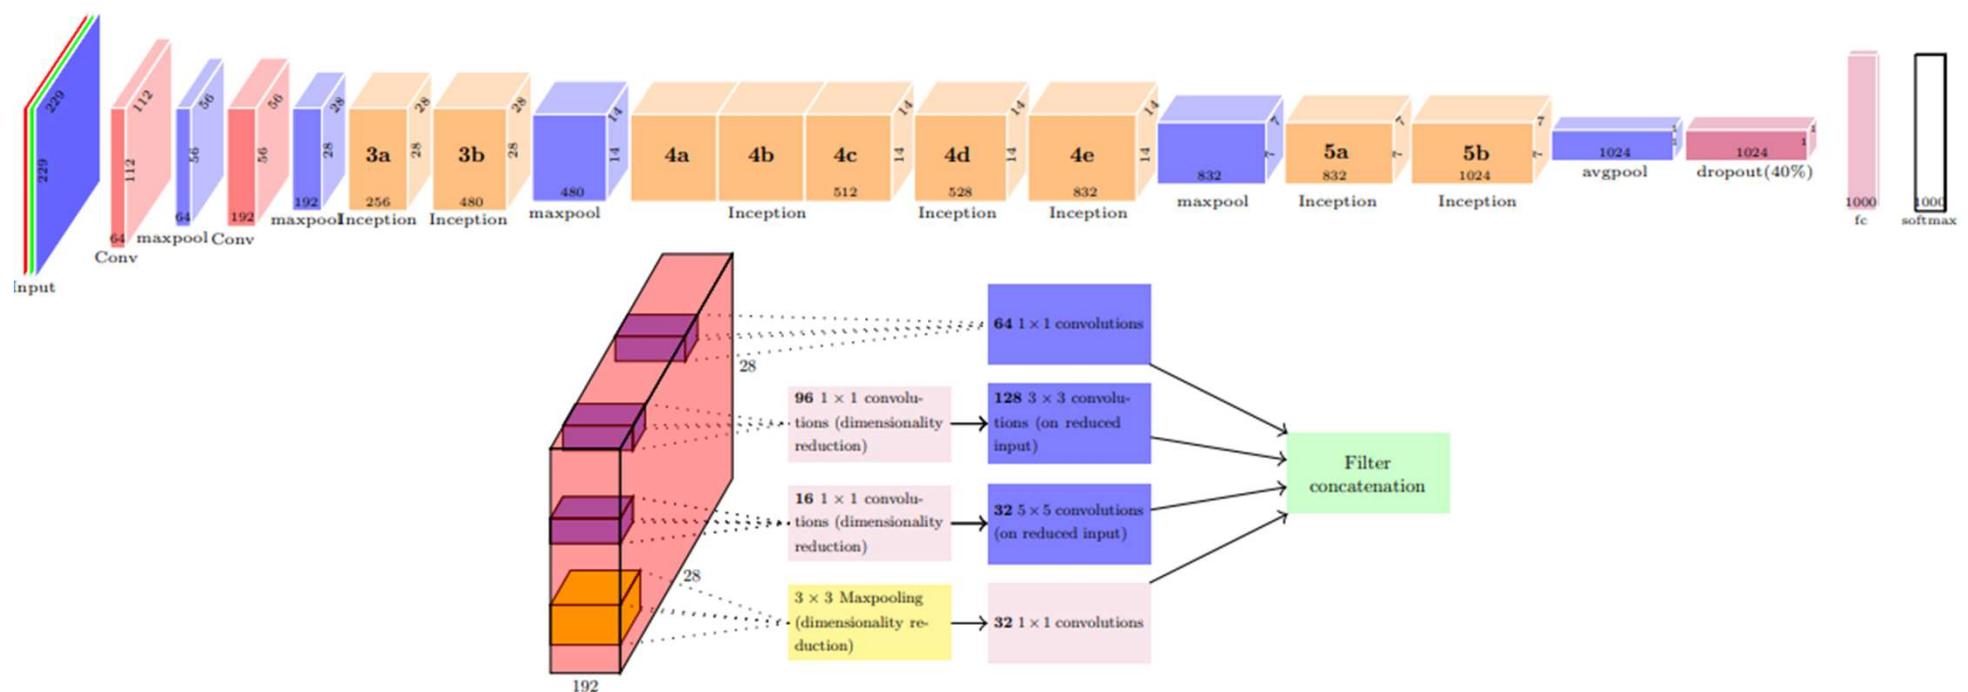


- Yes, by using 1×1 convolutions
- Huh?? What does a 1×1 convolution do ?
- It aggregates along the depth
- So convolving a $D \times W \times H$ input with $D_1 1 \times 1$ ($D_1 < D$) filters will result in a $D_1 \times W \times H$ output ($S = 1, P = 0$)
- If $D_1 < D$ then this effectively reduces the dimension of the input and hence the computations
- Specifically instead of $O(F \times F \times D)$ we will need $O(F \times F \times D_1)$ computations
- We could then apply subsequent 3×3 , 5×5 filter on this reduced output

GoogLeNet



GoogLeNet



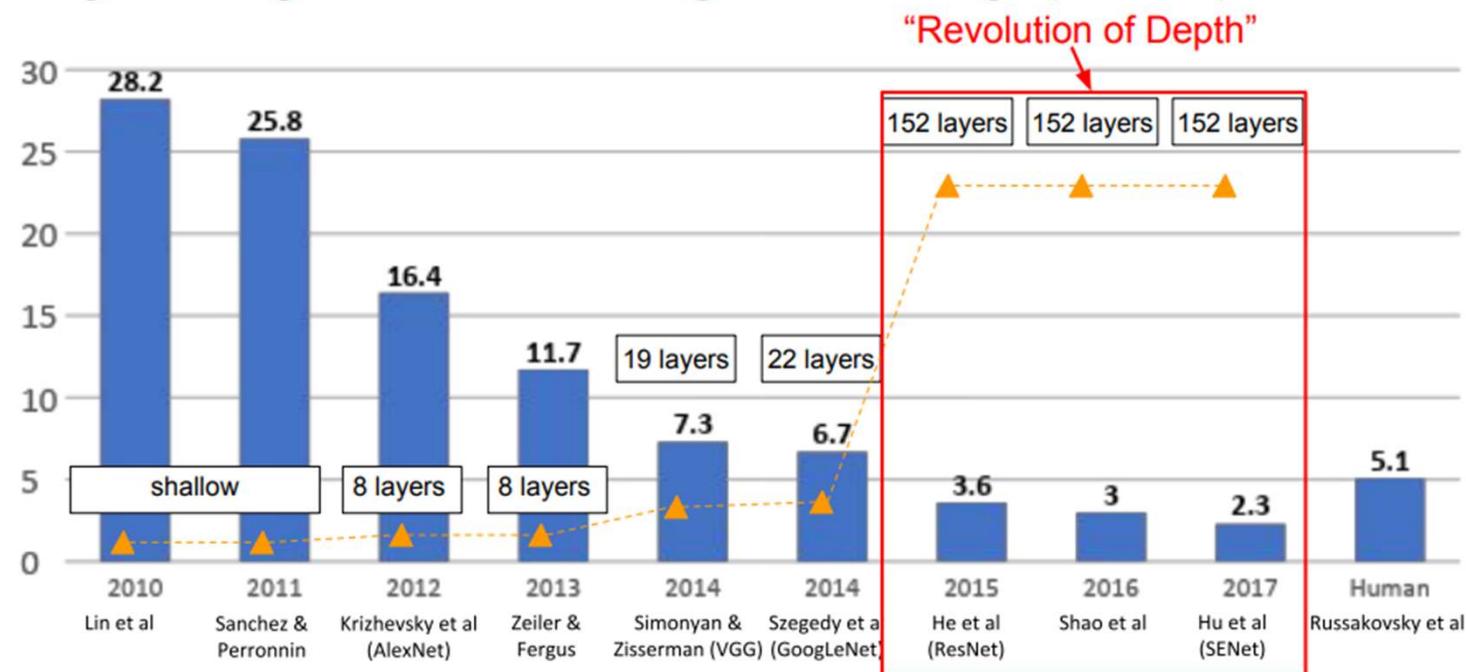
GoogLeNet

Successors

Following the success of GoogLeNet, Google continued to improve the Inception architecture with the development of:

- **Inception v2:** Introduced factorization techniques like reducing 5x5 convolutions into two smaller 3x3 convolutions.
- **Inception v3:** Introduced optimizations like the use of batch normalization, improved factorization strategies, and the RMSProp optimizer.
- **Inception v4 and Inception-ResNet:** Further improvements in both accuracy and efficiency, with the inclusion of residual connections.

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



ResNet

Deep Residual Learning for Image Recognition

Kaiming He

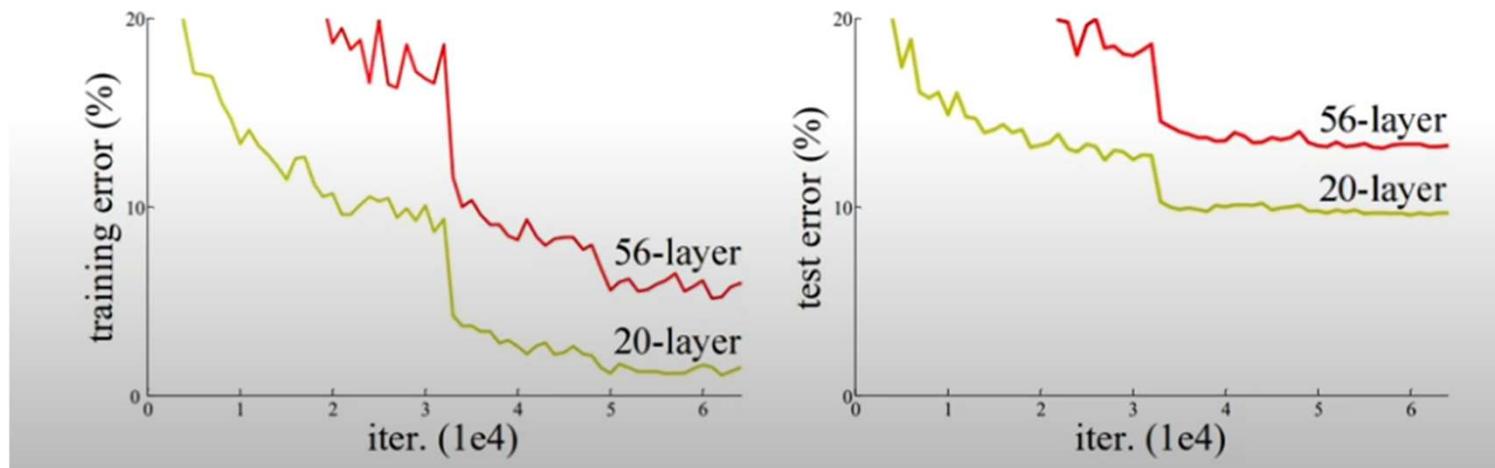
Xiangyu Zhang

Shaoqing Ren

Jian Sun

Microsoft Research

{kahe, v-xiangz, v-shren, jiansun}@microsoft.com



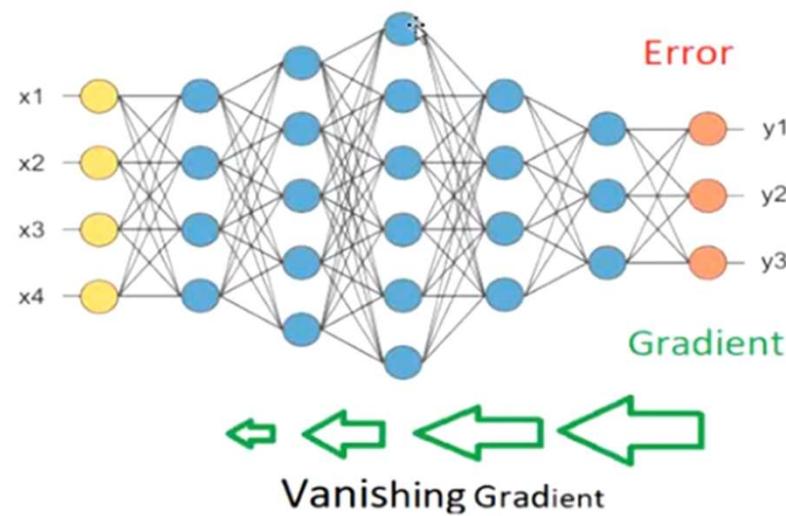
56-layer model performs worse on both training and test error -> The deeper model performs worse, but it's not caused by overfitting!

Hypothesis: the problem is an optimization problem, deeper models are harder to optimize

Is learning better networks as easy as stacking more layers?

- vanishing/exploding gradients
 - normalized initialization and intermediate normalization layers
- degradation problem : with the network depth increasing, accuracy gets saturated and then degrades rapidly.
- not caused by overfitting
- Addressed by deep residual learning

Vanishing Gradient Problem



The Degradation Problem

- When the network depth increases, accuracy gets saturated, and then **degrades rapidly**
- Adding more layers leads to higher training error

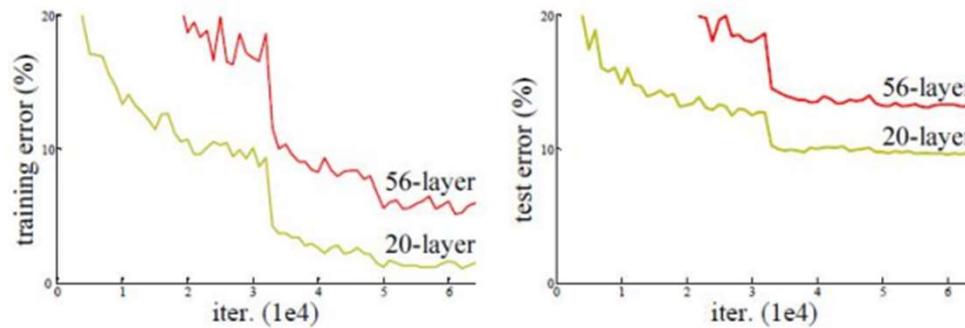
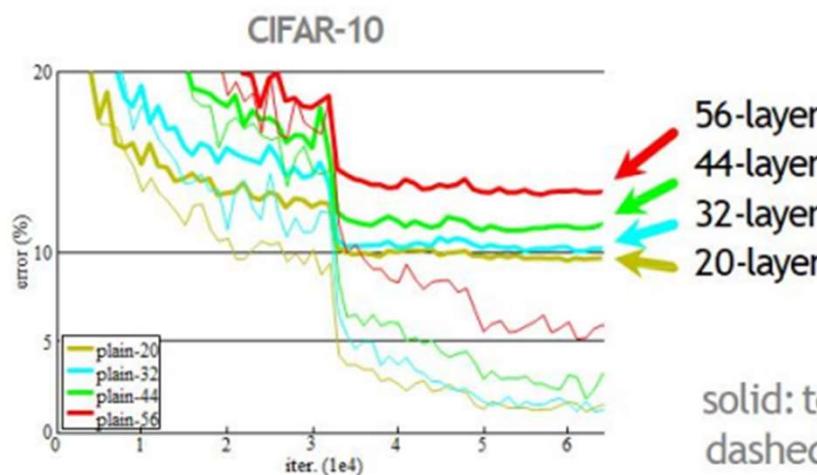
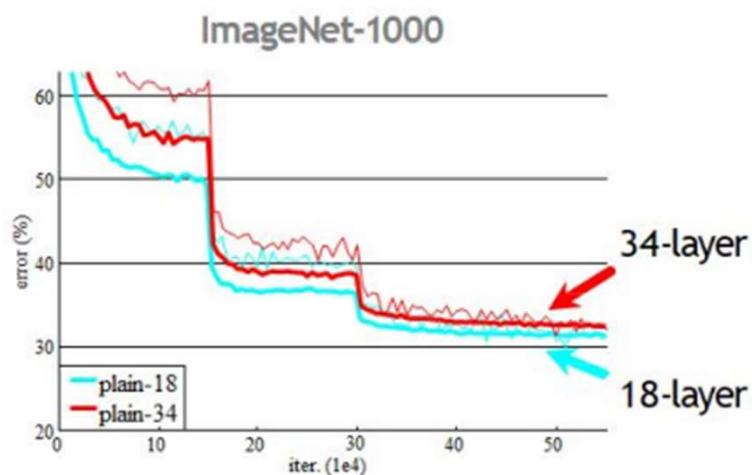


Figure 1: Training Error (left) and Test error (right) on CIFAR-10 with “Plain” networks

The Degradation Problem



solid: test/val
dashed: train



- Overly deep plain networks have higher training error
- A general phenomenon, observed in many datasets

Deep Residual Learning

1. **Traditional Mapping:** A deep neural network tries to learn a function $F(x)$, which maps the input x to the desired output.
2. **Residual Mapping:** Instead of learning $F(x)$, ResNet learns the **residual**:

$$H(x) = F(x) + x$$

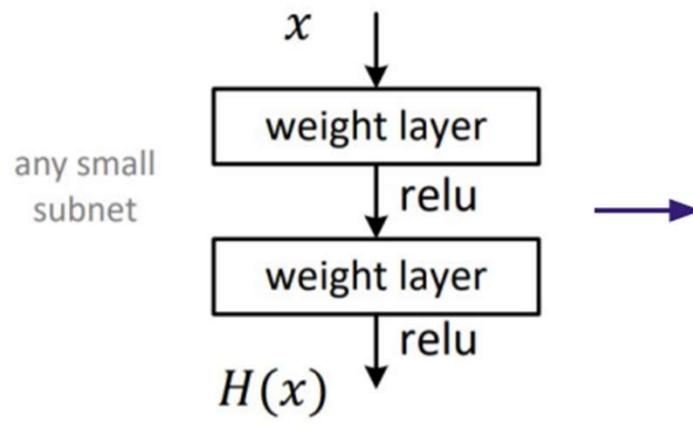
where:

- $F(x)$ is the residual function to be learned by the network.
- x is the input to the layer (or block).
- $H(x)$ is the final output after adding $F(x)$ and x .

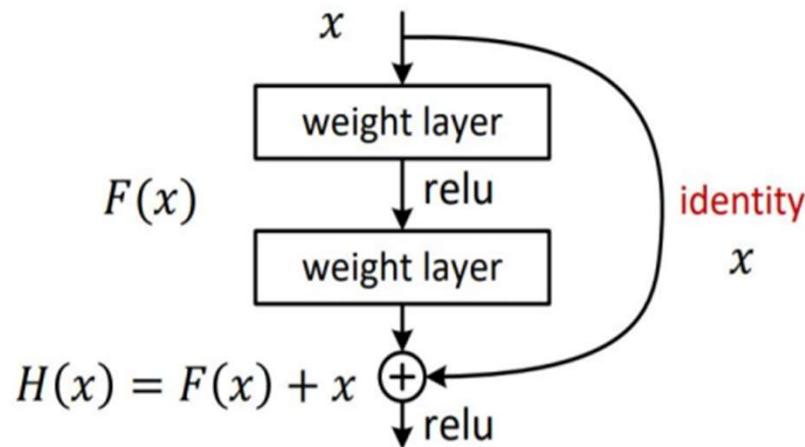
By reformulating the problem this way, the network focuses on learning the difference (residual) between the input and output, which is often easier than learning the entire transformation.

Deep Residual Learning

- Plain net



- Residual net



Learn the residual mapping $F(x)$ rather than unreferenced $H(x)$

Skip Connection: bypasses the convolutional layers and adds the input directly to the output.

Types of Skip Connections

- **Identity Connection:** The input x is directly added to the output without any transformation.
- **Projection Connection:** When dimensions do not match (e.g., due to different feature map sizes), a linear transformation (e.g., 1×1 convolution) is applied to x .

1. Identity Connection

- **Definition:**

- The input (x) is directly added to the output of the residual block ($F(x)$) without any transformation.
- Used when the input and output have the same dimensions.

- **Mathematical Representation:**

$$y = F(x) + x$$

- x : Input to the residual block.
- $F(x)$: Output of the block's transformations (e.g., convolution, ReLU, etc.).
- y : Final output after addition.

- **When to Use:**

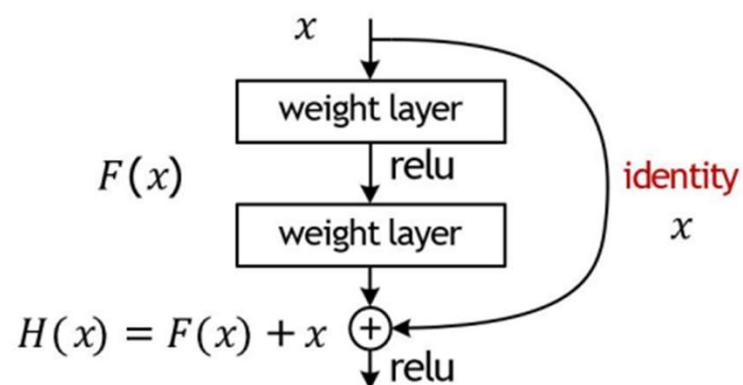
- When the input (x) and the residual block output ($F(x)$) have the same spatial dimensions and the same number of channels.
- No additional computation is needed for alignment.

2. Projection Connection

- **Definition:**
 - The input (x) is transformed (via a convolutional layer or other operations) before being added to the output ($F(x)$).
 - Used when the dimensions of the input and output differ.
- **Mathematical Representation:**
$$y = F(x) + W_s \cdot x$$
 - W_s : A projection layer (e.g., a 1×1 convolution) that transforms x to match the dimensions of $F(x)$.
 - $F(x)$: Output of the residual block.
 - y : Final output after addition.
- **When to Use:**
 - When the spatial size of the input and output differs (e.g., due to downsampling).
 - When the number of channels in x and $F(x)$ do not match.

Identity shortcut

- $F(x)$ is a **residual** mapping w.r.t. **identity**



- If optimal mapping is closer to identity, easier to capture small fluctuations
- If identity is optimal, easy to set weights as 0

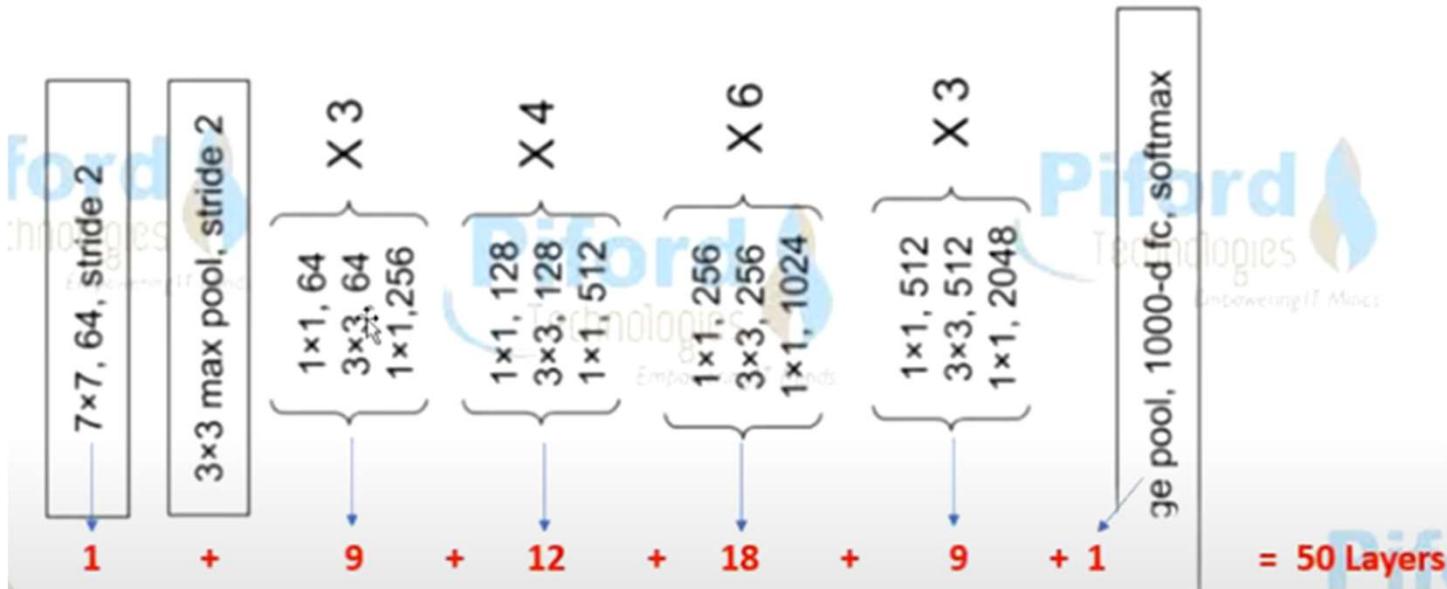
- Residual learning reformulates the target function $H(x)$ to:

$$H(x) = F(x) + x$$

where $F(x)$ is the residual function, and x is the input to the layer.

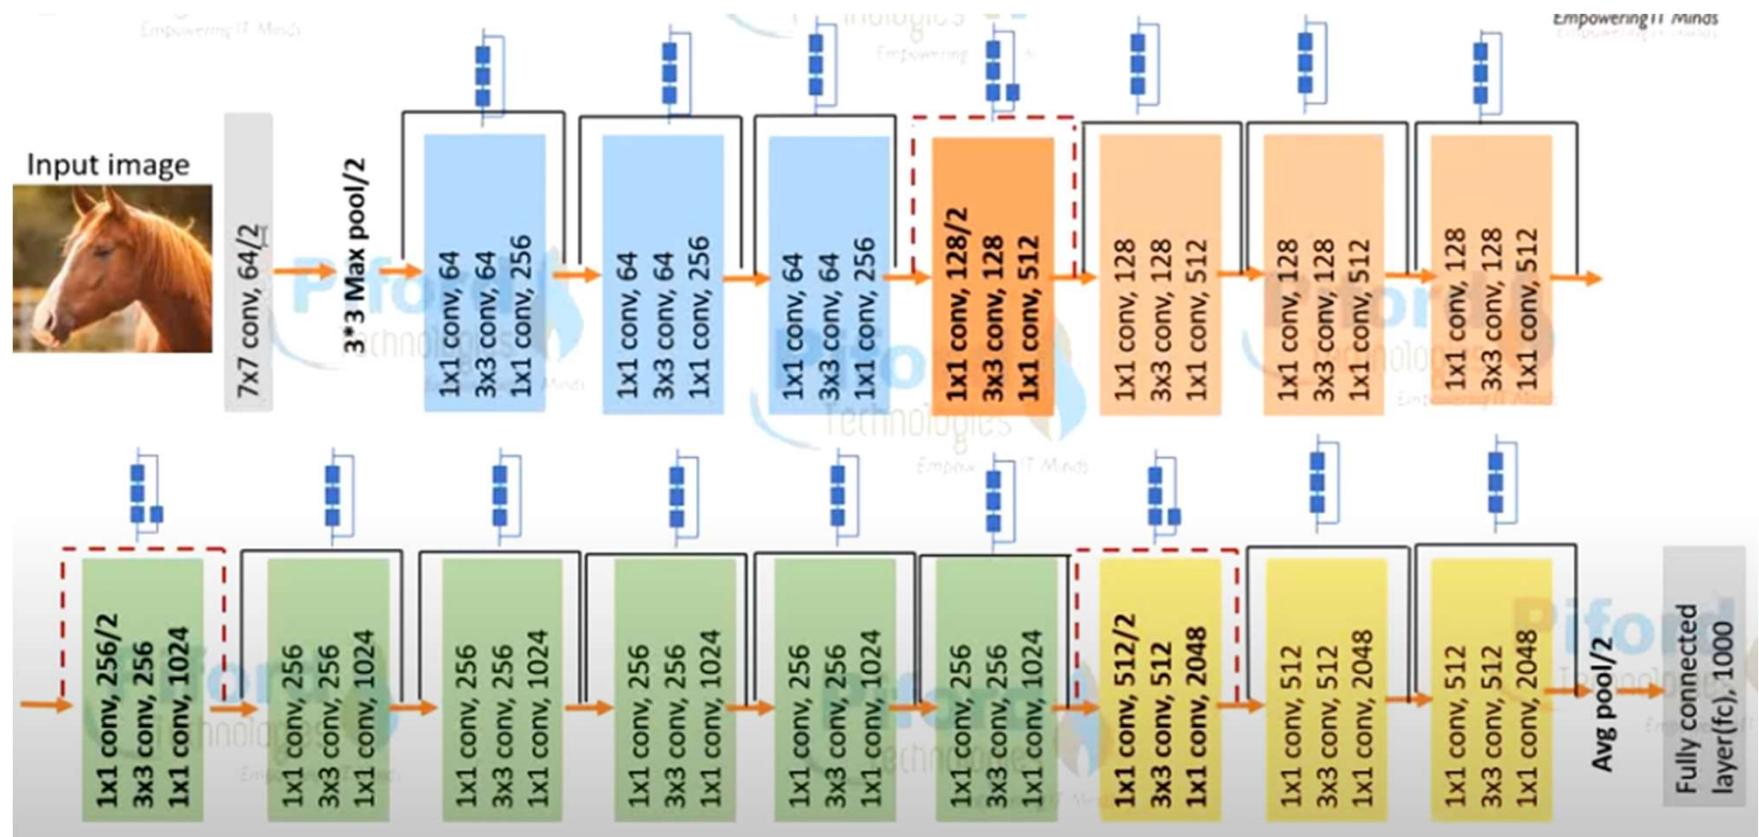
- If the desired function $H(x)$ is close to x (the identity mapping), the residual $F(x)$ will be small. Thus, the network only needs to learn these small fluctuations $F(x) = H(x) - x$.
- Learning small residuals $F(x)$ is simpler and faster because:
 - The gradients during backpropagation focus on fine-tuning the small differences.
 - Small residuals are less likely to cause numerical instability during training.

ResNet 50

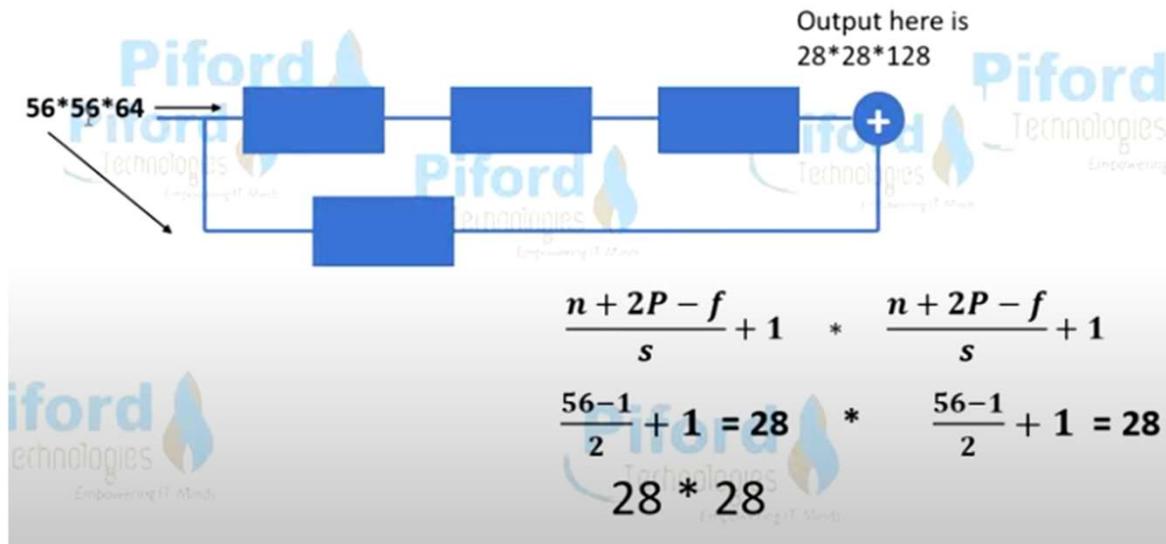


ResNet-18
ResNet-34
ResNet-50,
ResNet-101,
ResNet-152

ResNet 50

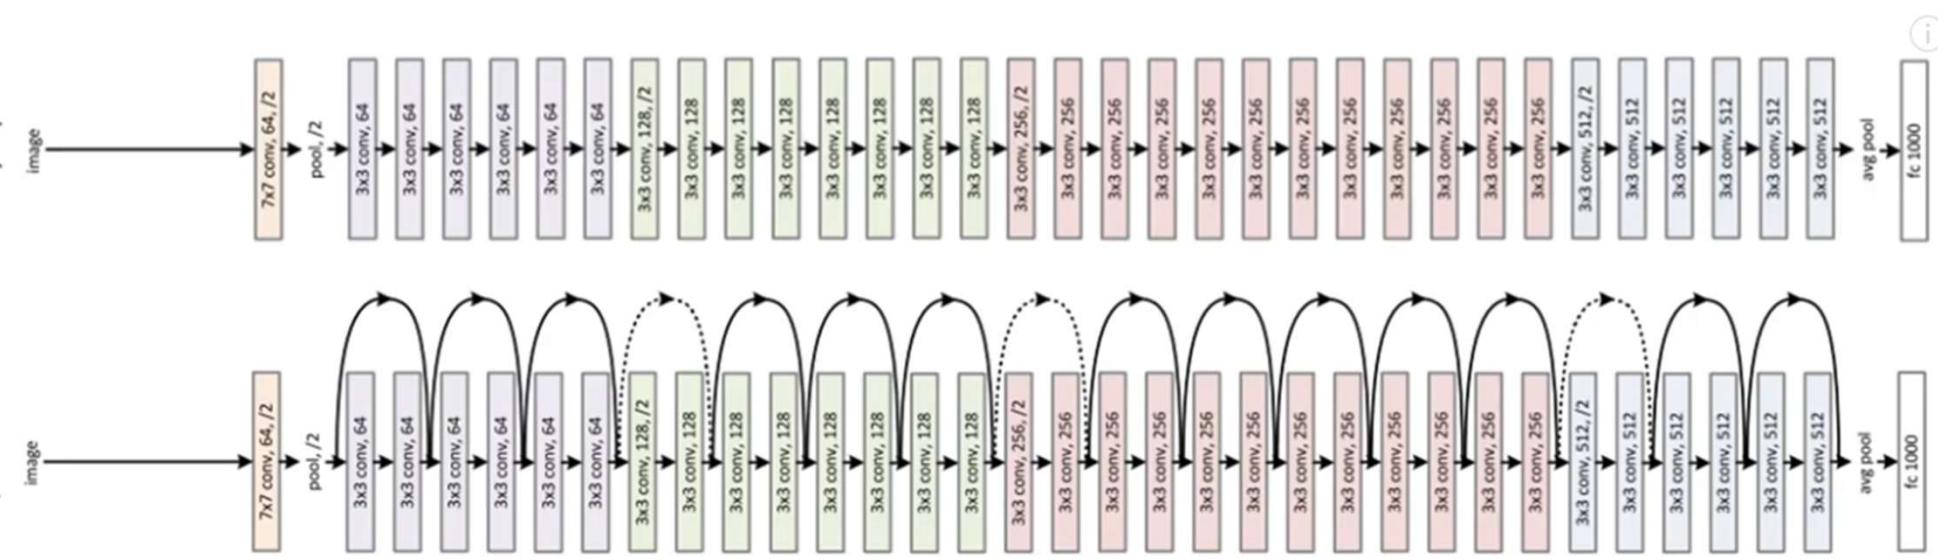


1 x 1 Convolutions



34-layer plain

34-layer residual



Summary

- Key Concept:
 - Learns residual mappings ($F(x) = H(x) - x$) instead of direct mappings ($H(x)$).
 - Introduced skip connections to alleviate **vanishing gradient problem**.
- Architecture:
 - Residual Block: $y = F(x) + x$ (Identity Connection).
 - Bottleneck Block: $1 \times 1 \rightarrow 3 \times 3 \rightarrow 1 \times 1$ convolutions for deeper networks.
- Variants:
 - ResNet-18/34: Basic residual blocks.
 - ResNet-50/101/152: Bottleneck blocks for efficiency.
- Advantages:
 - Enables training very deep networks (up to 152 layers).
 - Backbone for modern architectures like Faster R-CNN and Mask R-CNN.
- Applications:
 - Image Classification, Object Detection, Medical Imaging, etc.

Complexity Comparison

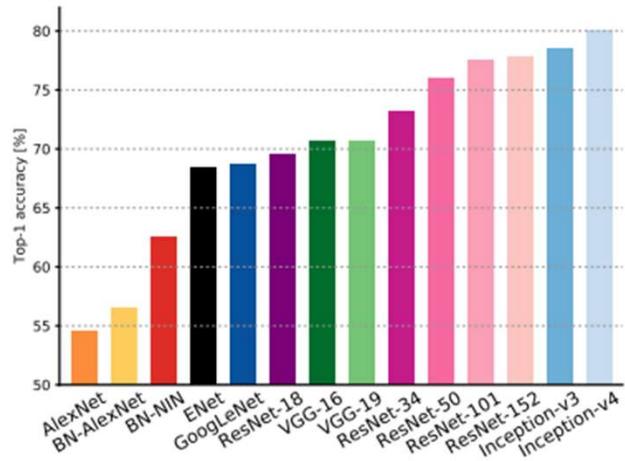


Figure 1: **Top1 vs. network.** Single-crop top-1 validation accuracies for top scoring single-model architectures. We introduce with this chart our choice of colour scheme, which will be used throughout this publication to distinguish effectively different architectures and their correspondent authors. Notice that networks of the same group share the same hue, for example ResNet are all variations of pink.

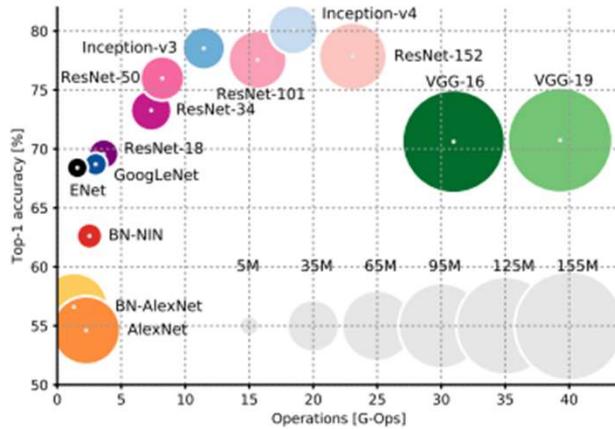
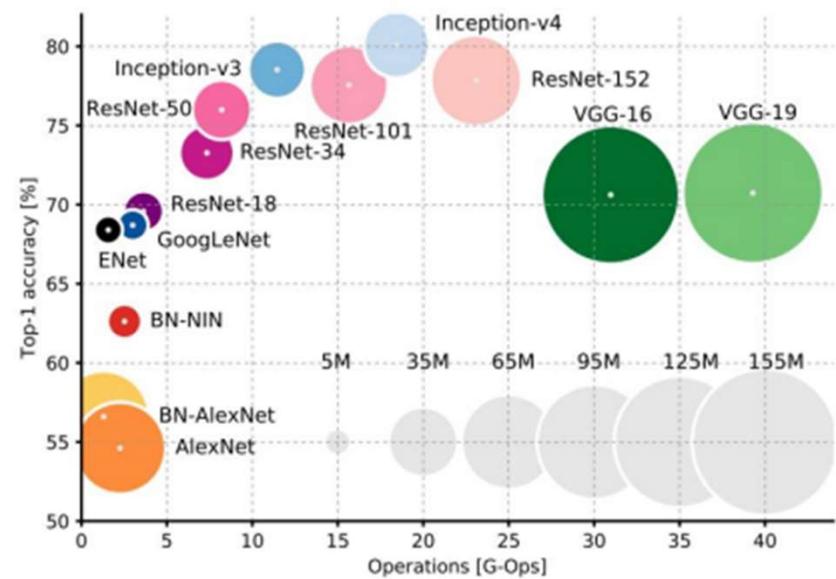
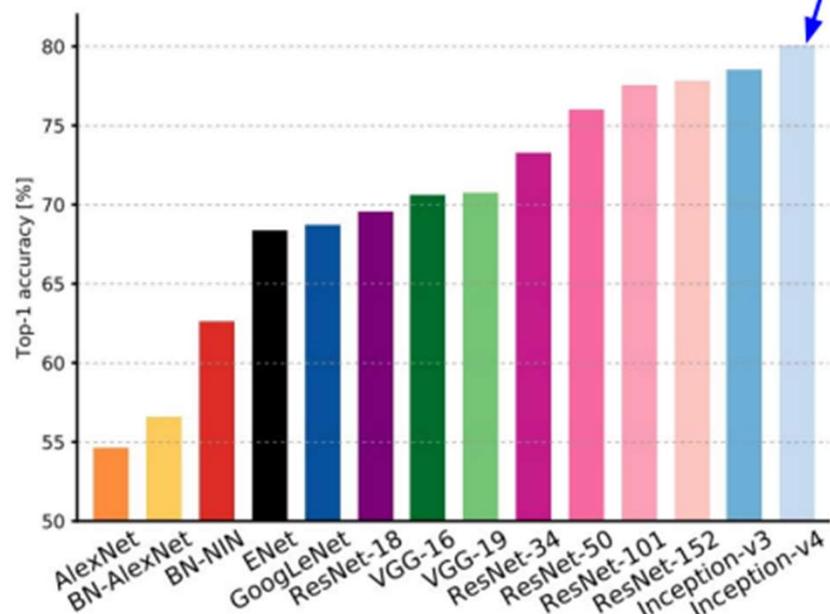


Figure 2: **Top1 vs. operations, size \propto parameters.** Top-1 one-crop accuracy versus amount of operations required for a single forward pass. The size of the blobs is proportional to the number of network parameters; a legend is reported in the bottom right corner, spanning from 5×10^6 to 155×10^6 params. Both these figures share the same y-axis, and the grey dots highlight the centre of the blobs.

Canziani, Alfredo, Adam Paszke, and Eugenio Culurciello. "An analysis of deep neural network models for practical applications." *arXiv preprint arXiv:1605.07678* (2016).

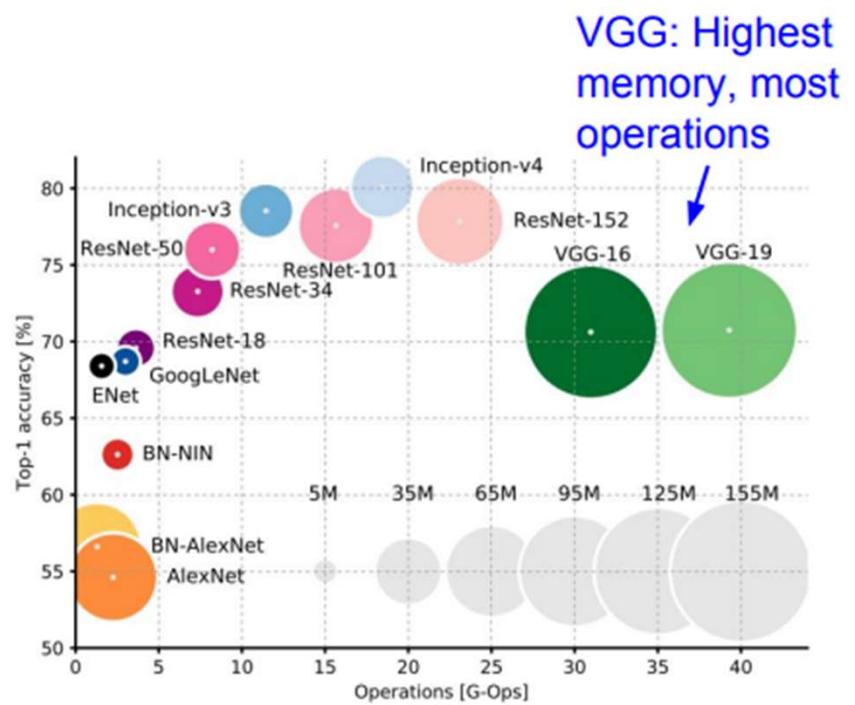
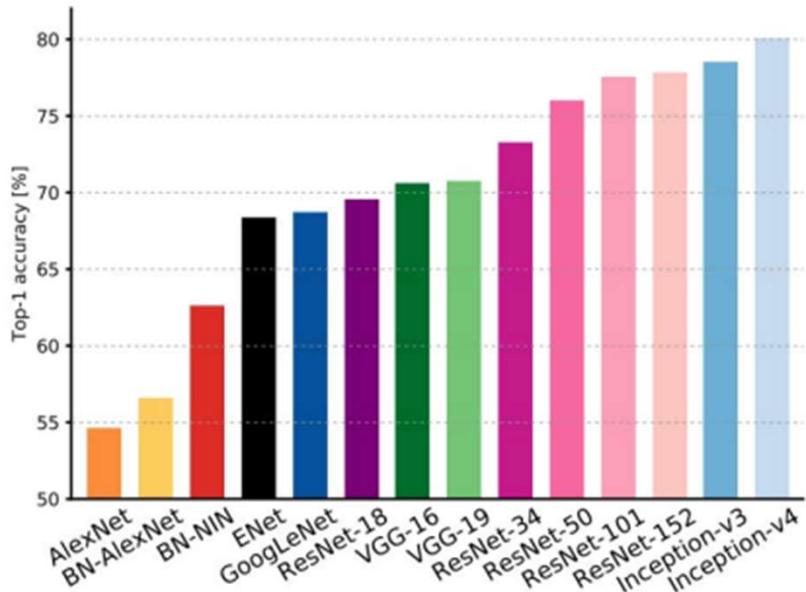
Comparing complexity...

Inception-v4: Resnet + Inception!



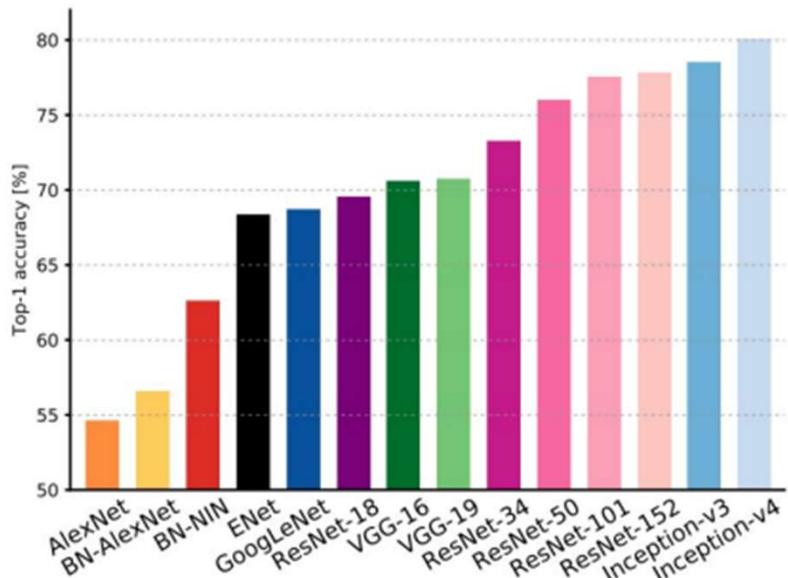
An Analysis of Deep Neural Network Models for Practical Applications, 2017.

Comparing complexity...

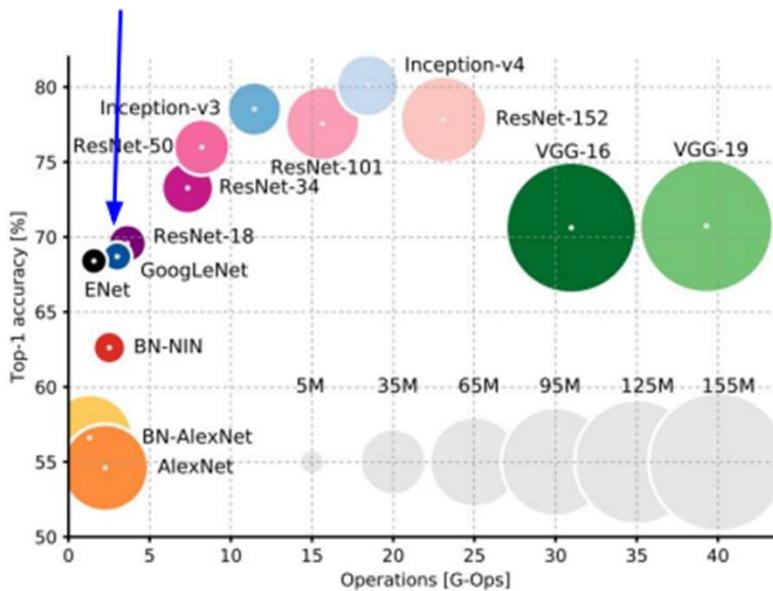


An Analysis of Deep Neural Network Models for Practical Applications, 2017.

Comparing complexity...

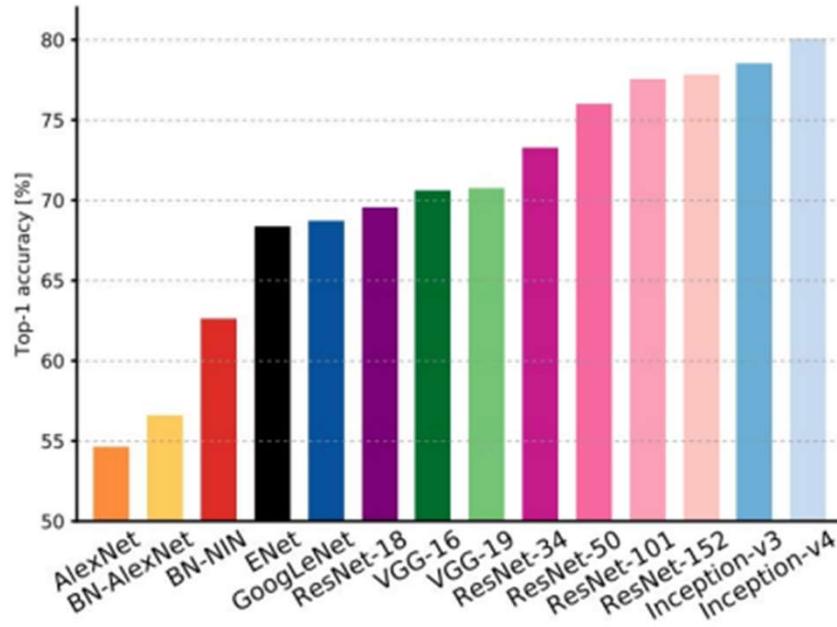


GoogLeNet:
most efficient

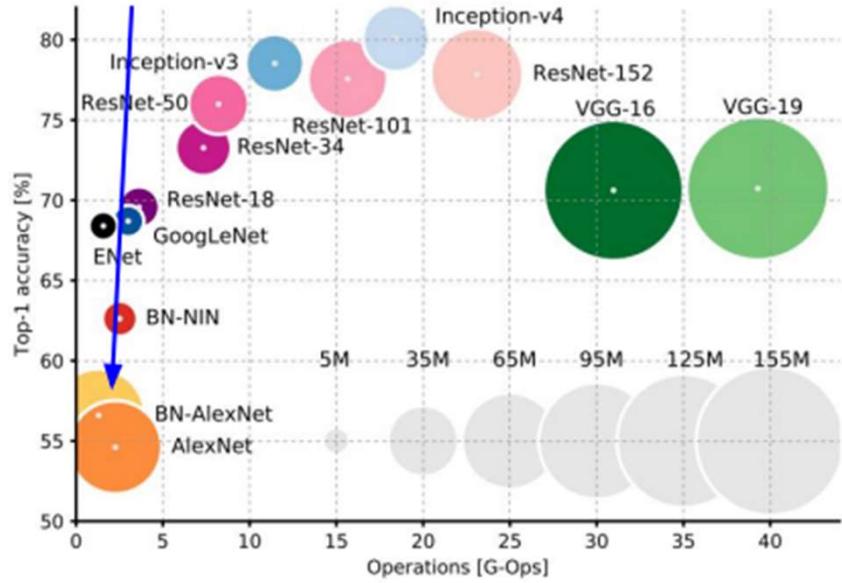


An Analysis of Deep Neural Network Models for Practical Applications, 2017.

Comparing complexity...

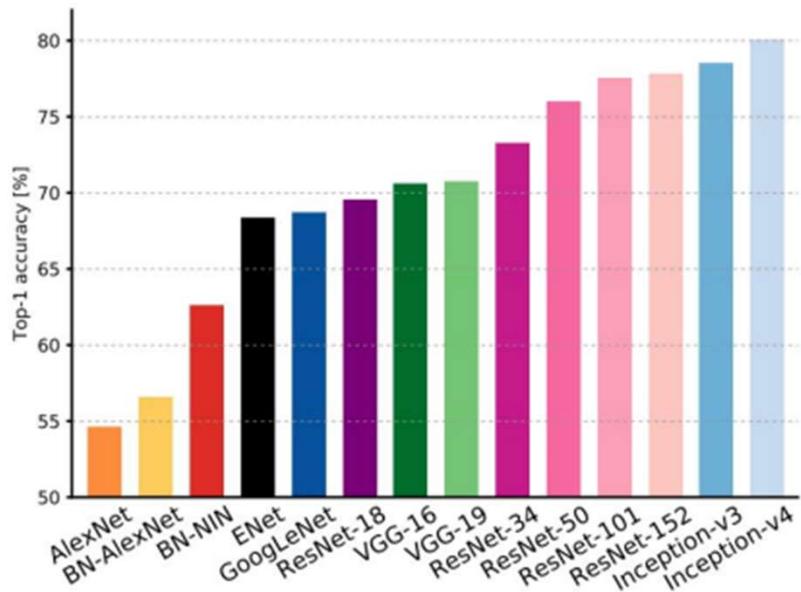


AlexNet:
Smaller compute, still memory heavy, lower accuracy

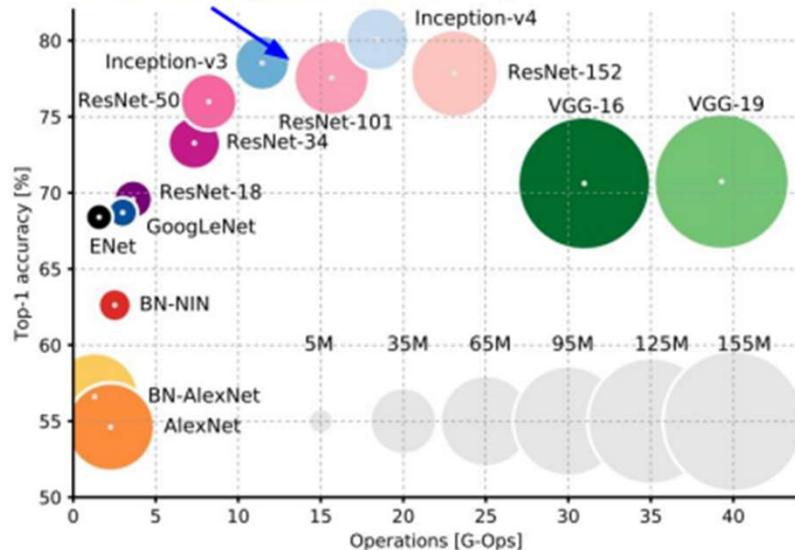


An Analysis of Deep Neural Network Models for Practical Applications, 2017.

Comparing complexity...



ResNet:
Moderate efficiency depending on
model, highest accuracy



An Analysis of Deep Neural Network Models for Practical Applications, 2017.

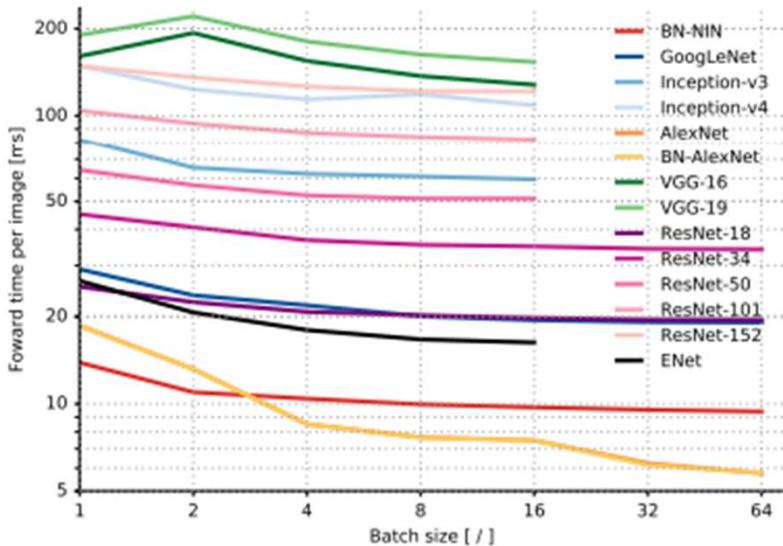


Figure 3: Inference time vs. batch size. This chart shows inference time across different batch sizes with a logarithmic ordinate and logarithmic abscissa. Missing data points are due to lack of enough system memory required to process larger batches. A speed up of 3× is achieved by AlexNet due to better optimisation of its fully connected layers for larger batches.

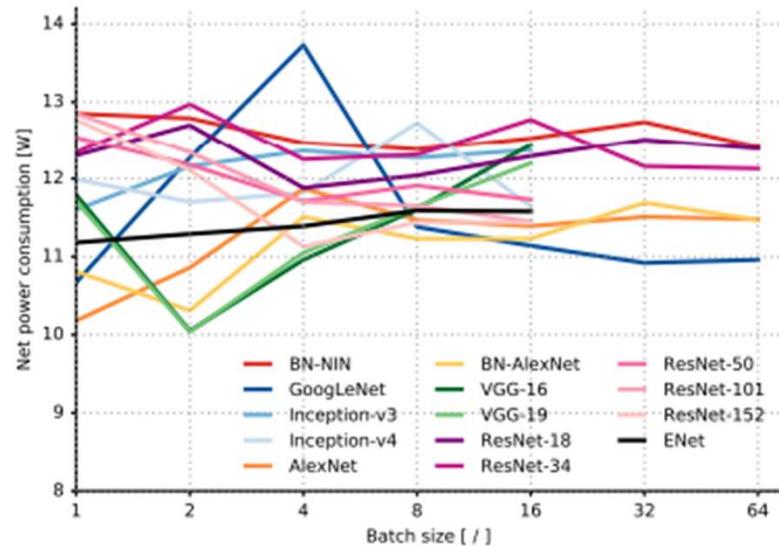
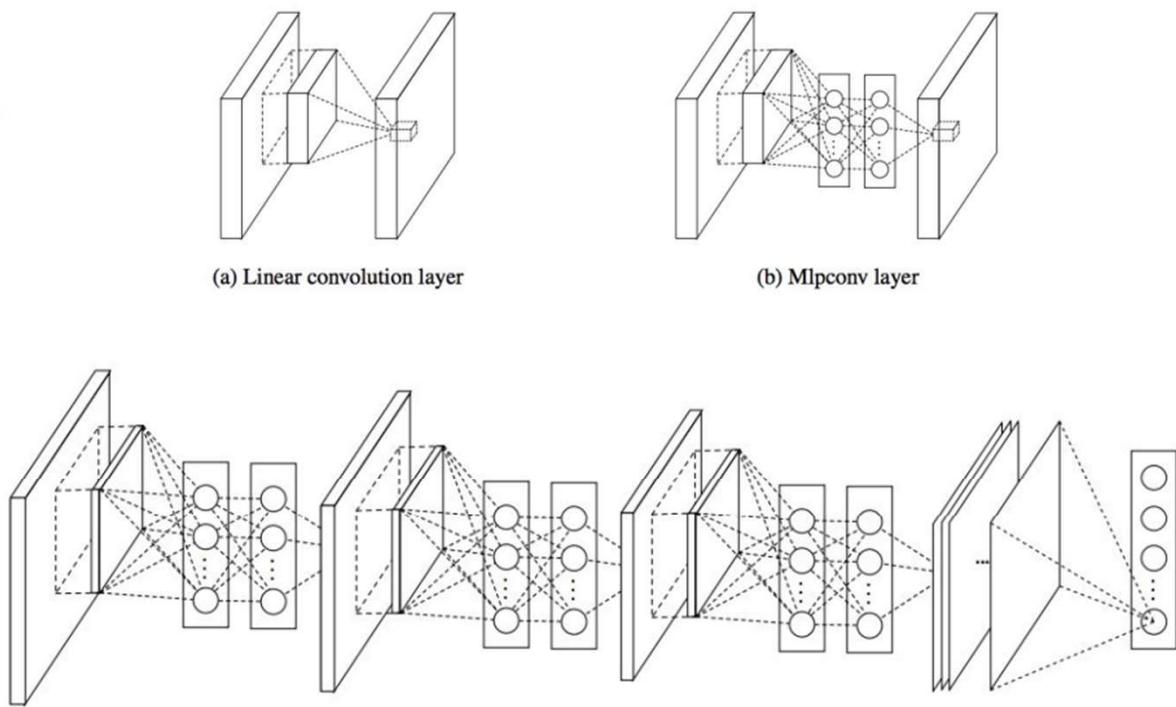


Figure 4: Power vs. batch size. Net power consumption (due only to the forward processing of several DNNs) for different batch sizes. The idle power of the TX1 board, with no HDMI screen connected, was 1.30 W on average. The max frequency component of power supply current was 1.4 kHz, corresponding to a Nyquist sampling frequency of 2.8 kHz.

Network in Network (NiN)

[Lin et al. 2014]

- Mlpconv layer with “micronetwork” within each conv layer to compute more abstract features for local patches
- Micronetwork uses multilayer perceptron (FC, i.e. 1x1 conv layers)
- Precursor to GoogLeNet and ResNet “bottleneck” layers
- Philosophical inspiration for GoogLeNet



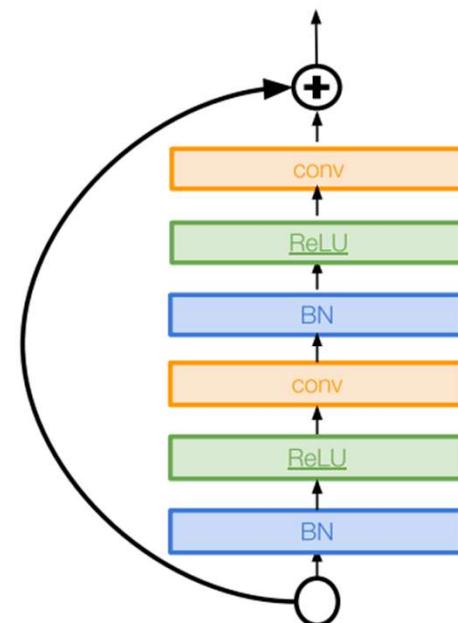
Figures copyright Lin et al., 2014. Reproduced with permission.

Improving ResNets...

Identity Mappings in Deep Residual Networks

[He et al. 2016]

- Improved ResNet block design from creators of ResNet
- Creates a more direct path for propagating information throughout network (moves activation to residual mapping pathway)
- Gives better performance

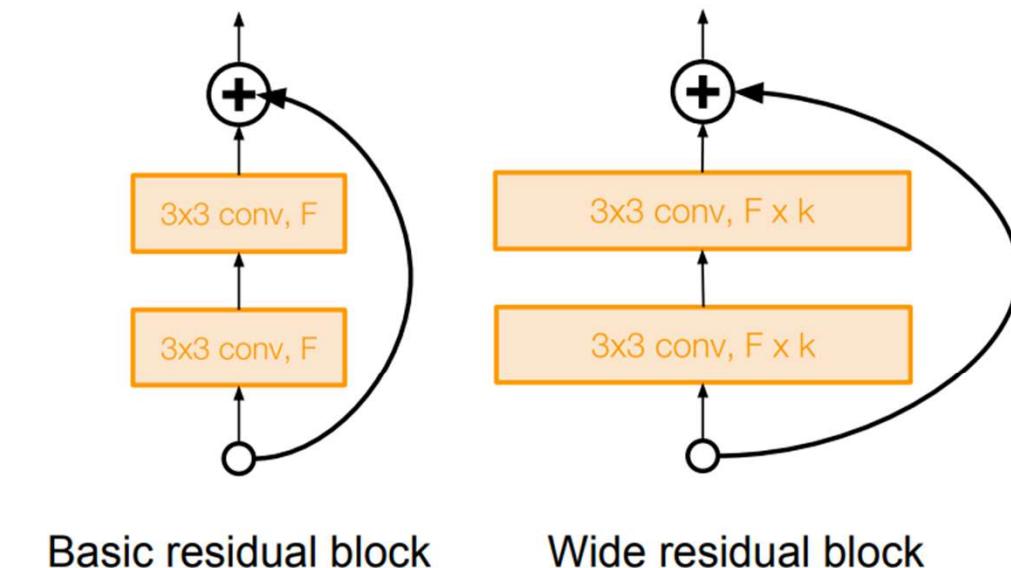


Improving ResNets...

Wide Residual Networks

[Zagoruyko et al. 2016]

- Argues that residuals are the important factor, not depth
- Use wider residual blocks ($F \times k$ filters instead of F filters in each layer)
- 50-layer wide ResNet outperforms 152-layer original ResNet
- Increasing width instead of depth more computationally efficient (parallelizable)

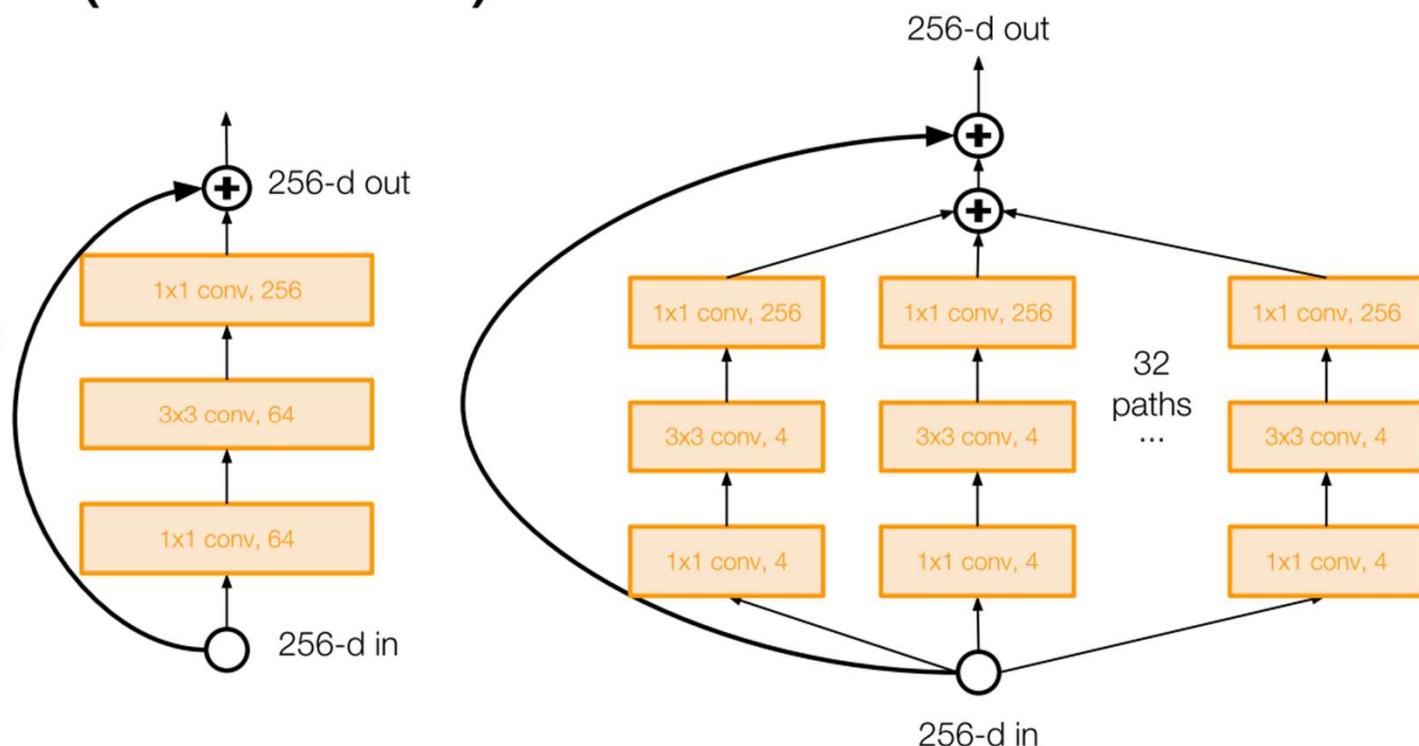


Improving ResNets...

Aggregated Residual Transformations for Deep Neural Networks (ResNeXt)

[Xie et al. 2016]

- Also from creators of ResNet
- Increases width of residual block through multiple parallel pathways (“cardinality”)
- Parallel pathways similar in spirit to Inception module

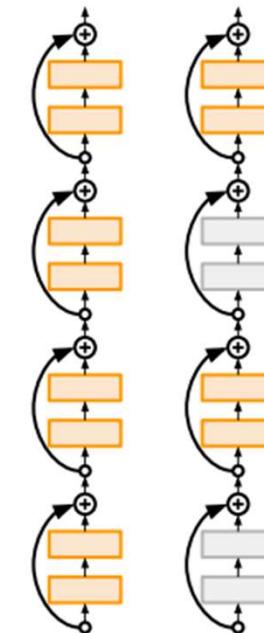


Improving ResNets...

Deep Networks with Stochastic Depth

[Huang et al. 2016]

- Motivation: reduce vanishing gradients and training time through short networks during training
- Randomly drop a subset of layers during each training pass
- Bypass with identity function
- Use full deep network at test time

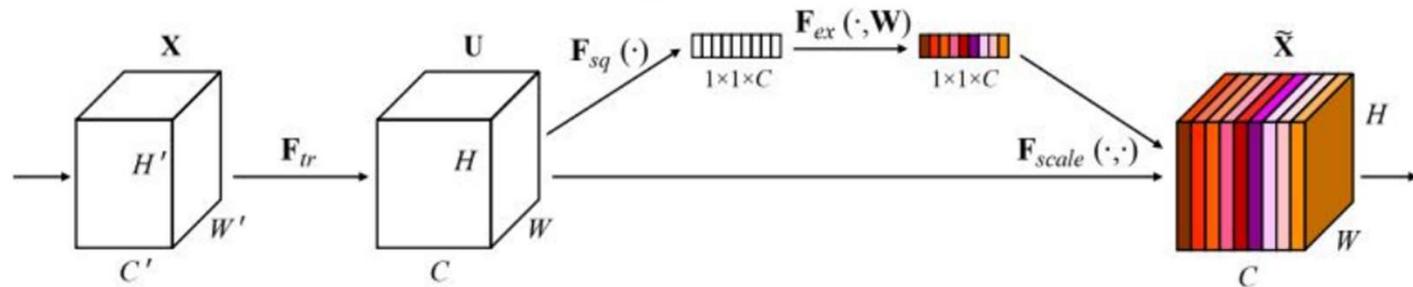
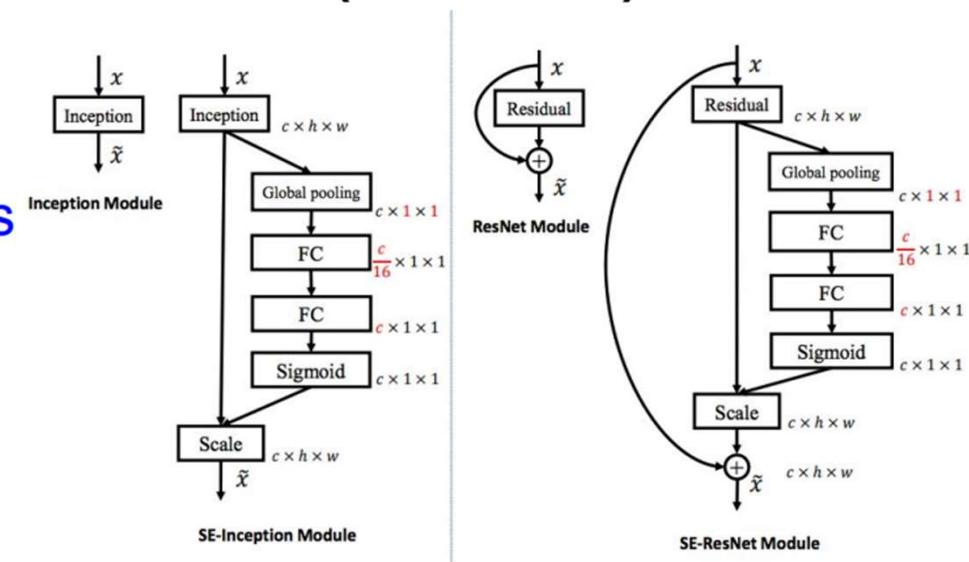


Improving ResNets...

Squeeze-and-Excitation Networks (SENet)

[Hu et al. 2017]

- Add a “feature recalibration” module that learns to adaptively reweight feature maps
- Global information (global avg. pooling layer) + 2 FC layers used to determine feature map weights
- ILSVRC’17 classification winner (using ResNeXt-152 as a base architecture)

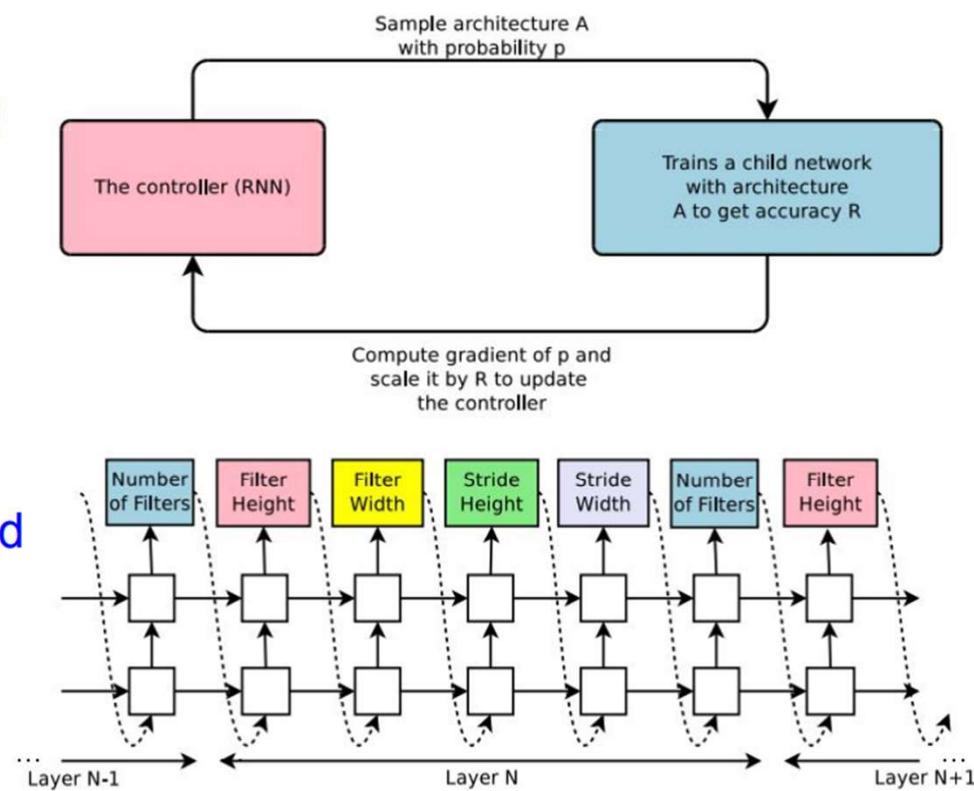


Meta-learning: Learning to learn network architectures...

Neural Architecture Search with Reinforcement Learning (NAS)

[Zoph et al. 2016]

- “Controller” network that learns to design a good network architecture (output a string corresponding to network design)
- Iterate:
 - 1) Sample an architecture from search space
 - 2) Train the architecture to get a “reward” R corresponding to accuracy
 - 3) Compute gradient of sample probability, and scale by R to perform controller parameter update (i.e. increase likelihood of good architecture being sampled, decrease likelihood of bad architecture)

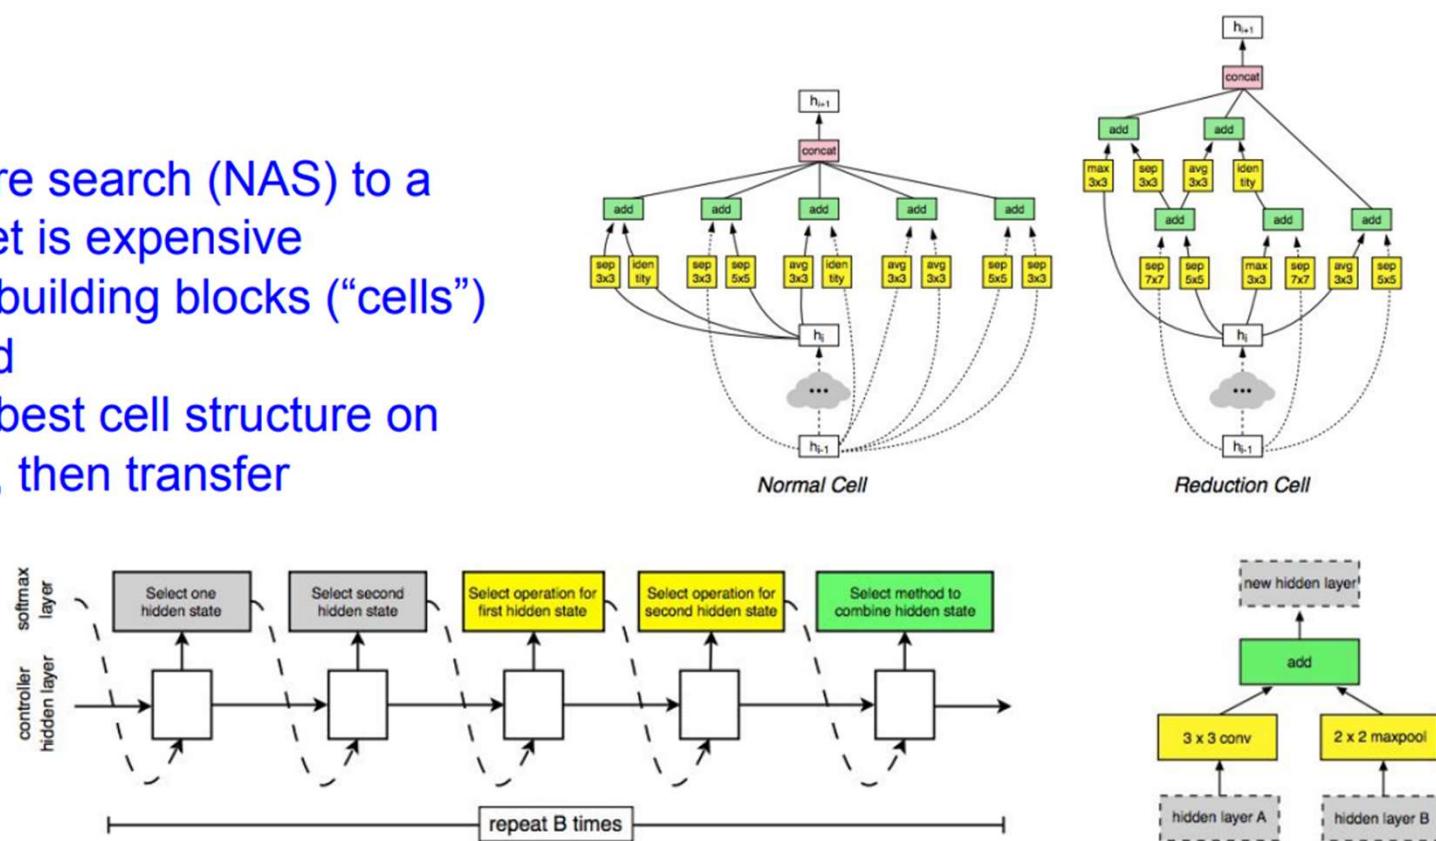


Meta-learning: Learning to learn network architectures...

Learning Transferable Architectures for Scalable Image Recognition

[Zoph et al. 2017]

- Applying neural architecture search (NAS) to a large dataset like ImageNet is expensive
- Design a search space of building blocks (“cells”) that can be flexibly stacked
- NASNet: Use NAS to find best cell structure on smaller CIFAR-10 dataset, then transfer architecture to ImageNet



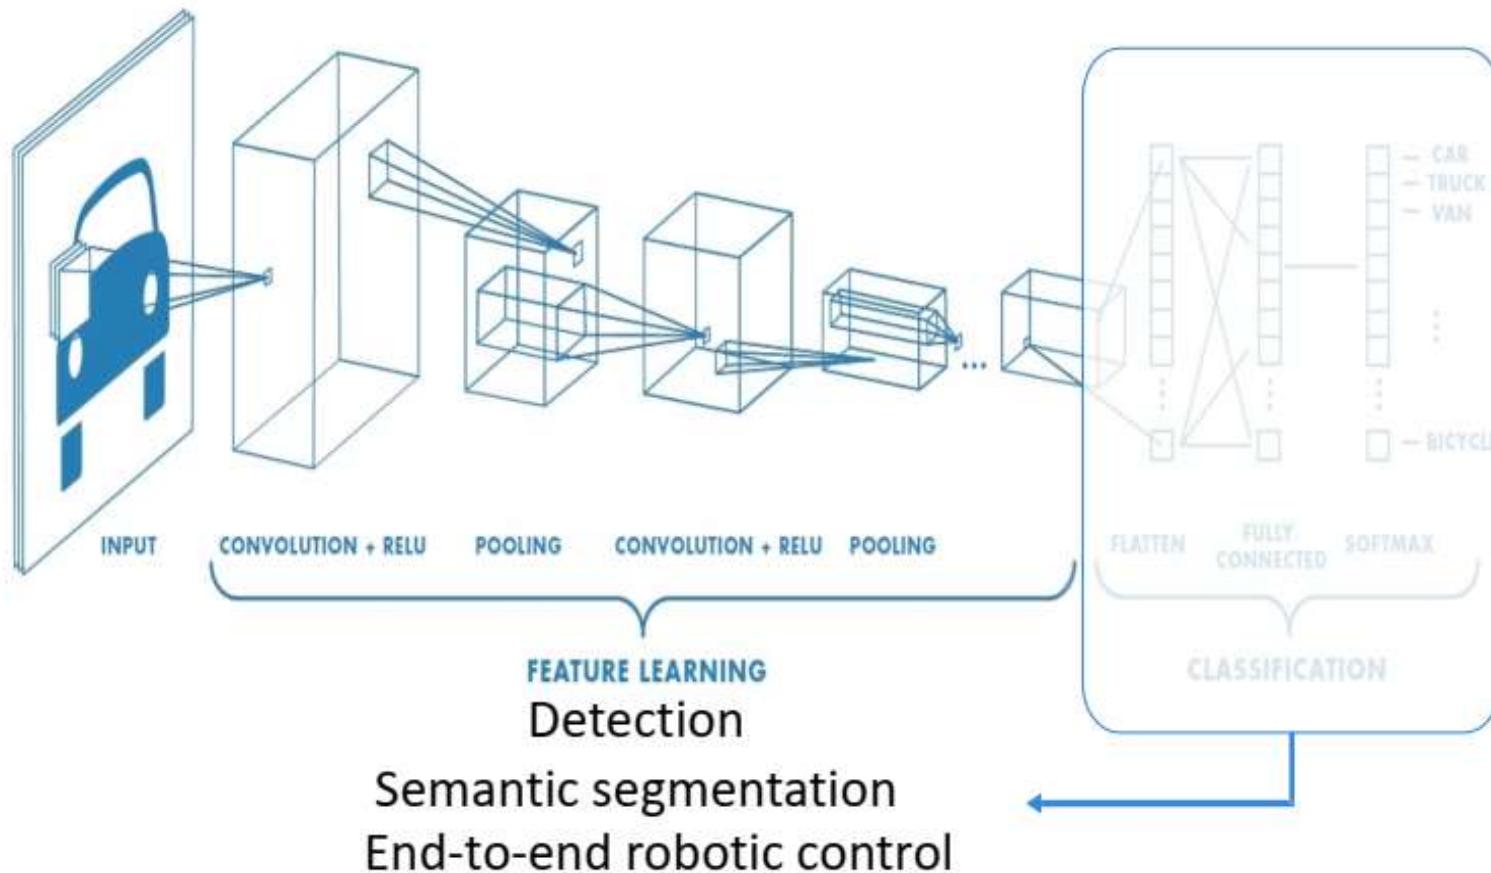
Summary: CNN Architectures

- VGG, GoogLeNet, ResNet all in wide use, available in model zoos
- ResNet current best default, also consider SENet when available
- Trend towards extremely deep networks
- Significant research centers around design of layer / skip connections and improving gradient flow
- Efforts to investigate necessity of depth vs. width and residual connections
- Even more recent trend towards meta-learning

Applications

https://cs231n.stanford.edu/slides/2018/cs231n_2018_lecture09.pdf

An Architecture for Many Applications



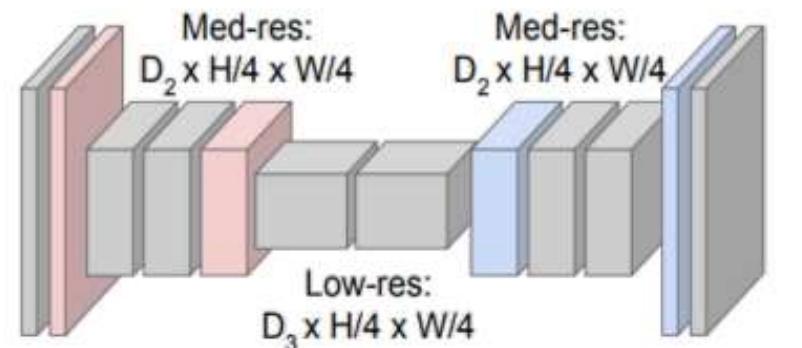
Semantic Segmentation: Fully Convolutional Networks

FCN: Fully Convolutional Network.

Network designed with all convolutional layers, with **downsampling** and **upsampling** operations



Input:
 $3 \times H \times W$



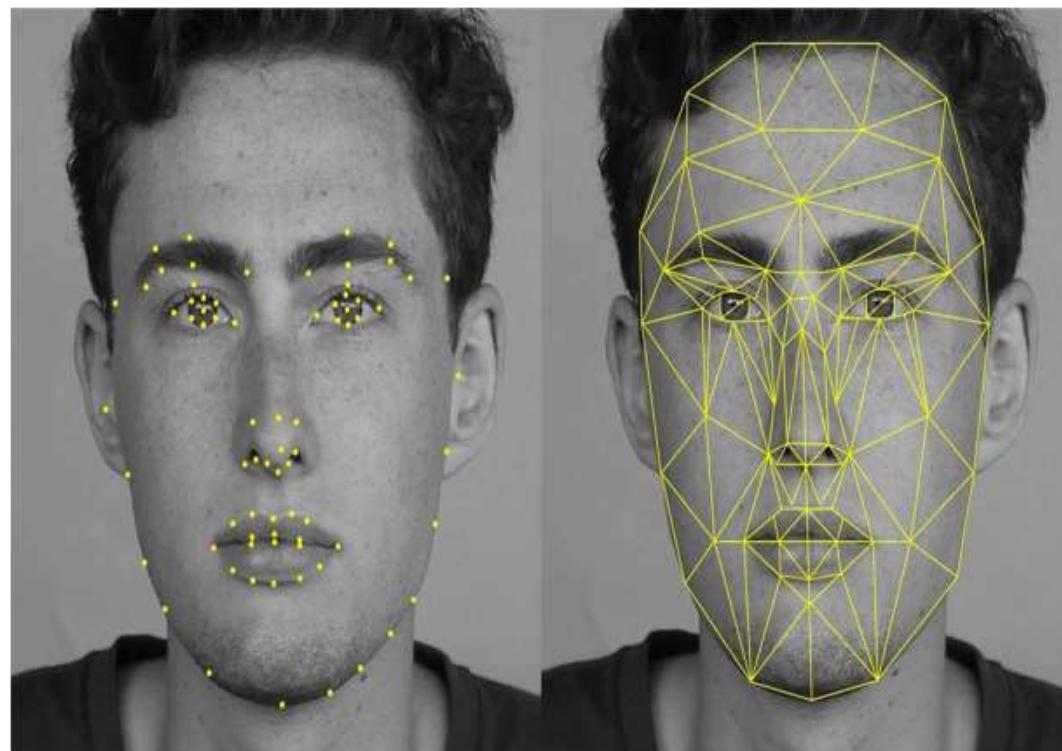
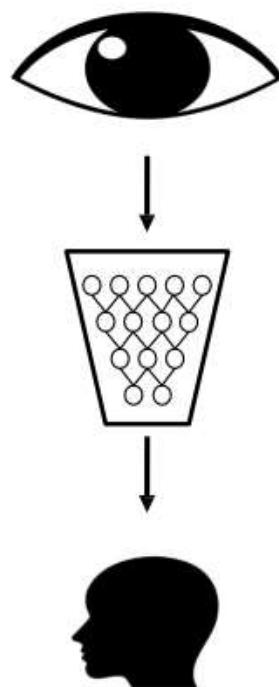
High-res:
 $D_1 \times H/2 \times W/2$

High-res:
 $D_1 \times H/2 \times W/2$

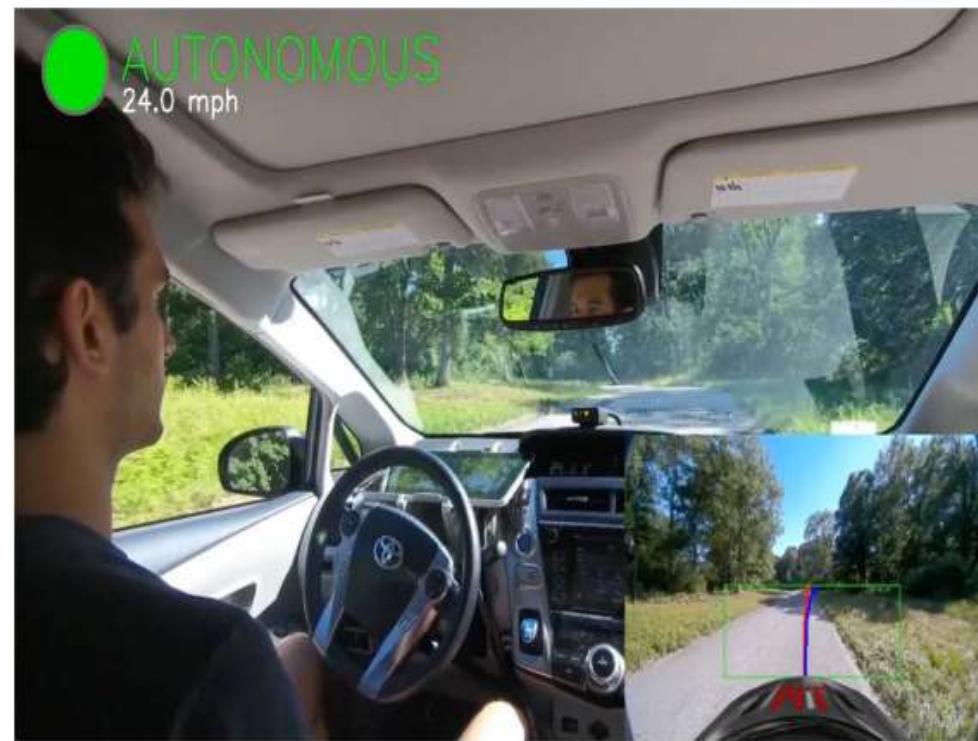
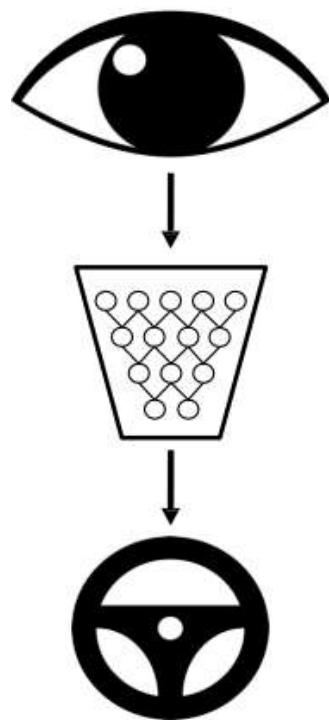


Predictions:
 $H \times W$

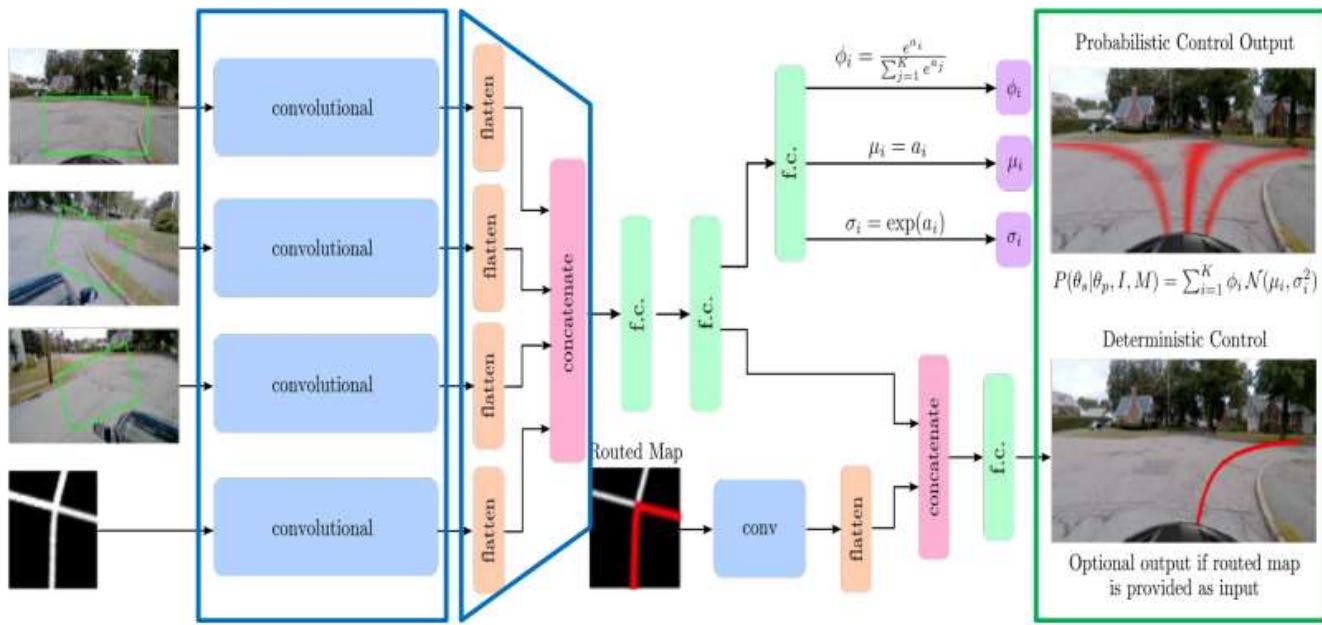
Facial Detection & Recognition



Self-Driving Cars



End-to-End Framework for Autonomous Navigation

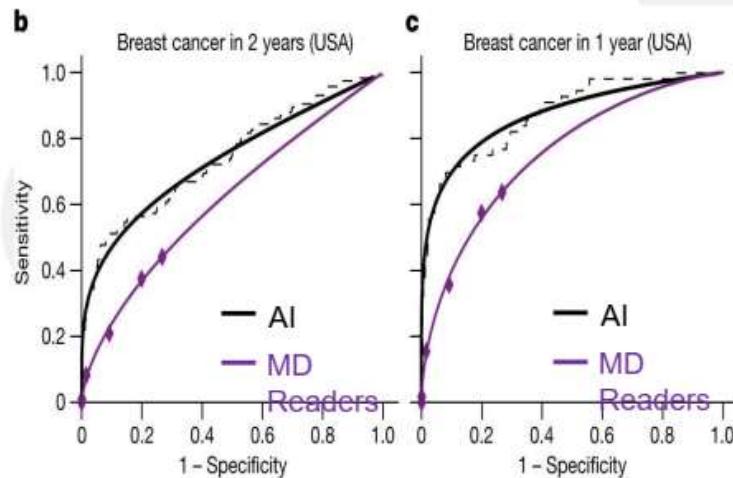


Entire model trained end-to-end
without any human labelling or annotations

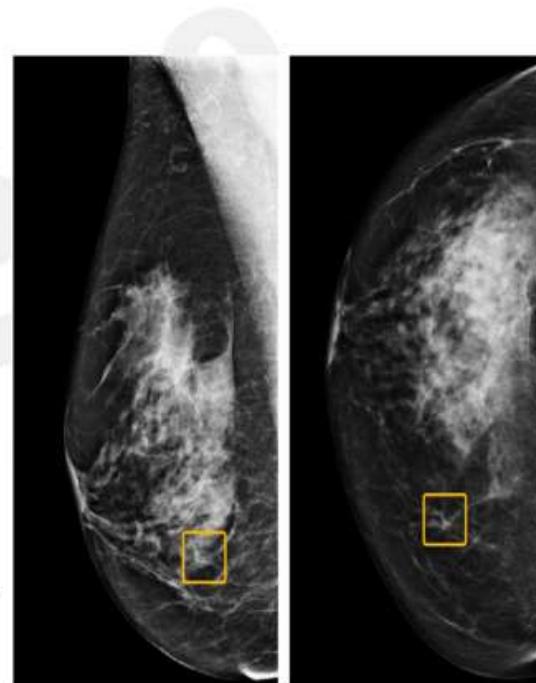
Breast Cancer Screening

International evaluation of an AI system for breast cancer screening

nature



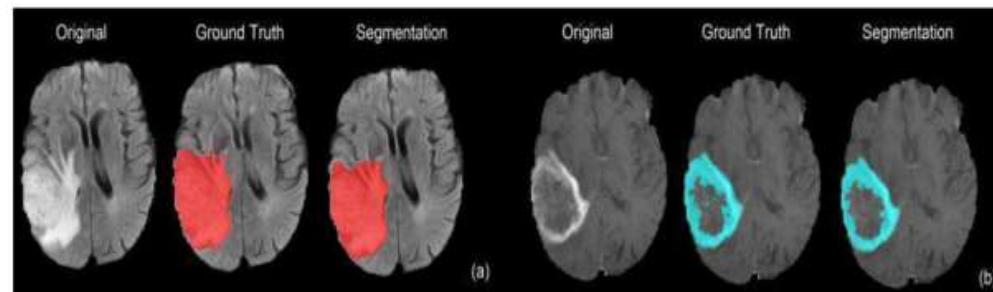
CNN-based system outperformed expert radiologists at detecting breast cancer from mammograms



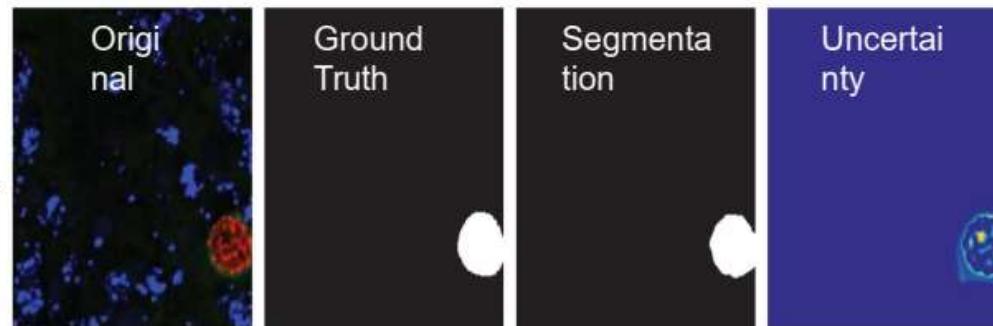
Breast cancer case missed by radiologist but detected by AI

Semantic Segmentation: Biomedical Image Analysis

Brain Tumors
Dong+ MIUA
2017.



Malaria Infection
Soleimany+ arXiv
2019.

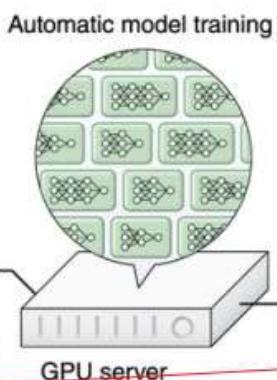


DeepBind

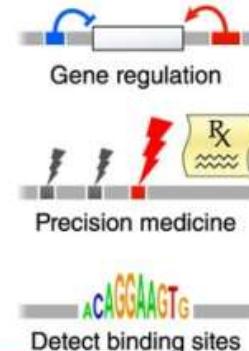
1. High-throughput experiments



2. Massively parallel deep learning



3. Community needs



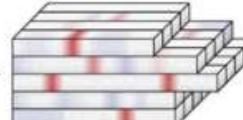
Current batch of inputs

CTAACGACCCGTCCT
CTAGGGGACCAAGTACT
TAGCACCCTCTATTCGACCC
CTCGGGGCCCTGCATC
TACAAATGAGCACCAA

Convolve

Motif detectors

Motif scans



Thresholds

Features

Outputs

Weights

Targets

Prediction errors

Current model parameters

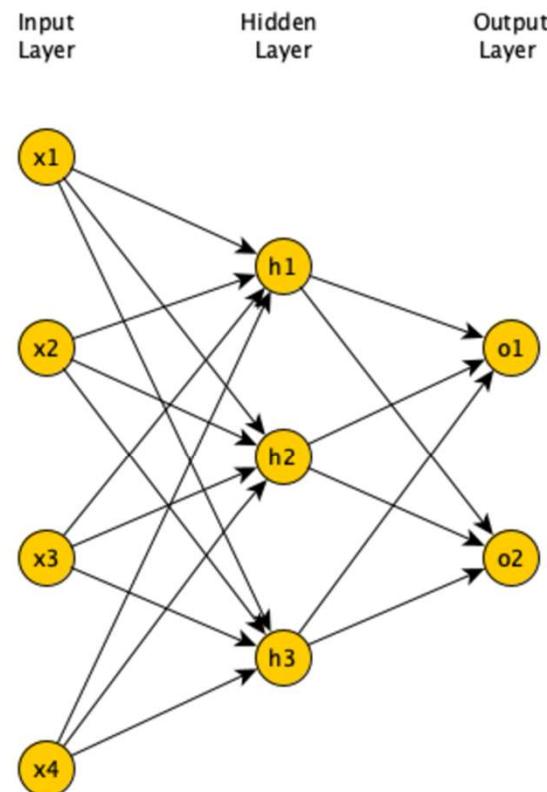
Parameter updates

Update +

Backprop

Feedforward neural networks (FNN)

- In FNN all nodes in a layer connected to all nodes in next layer
 - Each connection has a weight, must be learned by learning algorithm



Limitations of FNN

- If input is 64 pixel by 64 pixel grayscale image
 - Each grayscale pixel represented by 1 value, usually between 0 to 255
 - Where 0 is black, 255 is white, and values in between are shades of gray
- Since each grayscale pixel represented by 1 value, we say *channel* size is 1
- Image represented by $64 \times 64 \times 1 = 4,096$ values (rows x columns x channels)
 - Hence, input layer of FNN has 4096 nodes
- Lets assume next layer has 500 nodes
 - Since FNN fully connected, we have $4,096 \times 500 = 2,048,000$ weights

Limitations of FNN

- For complex problems, we need multiple hidden layers in our FNN
 - Compunds the problem of having many weights
- Having too many weights
 - Makes learning more difficult as dimension of search space is increased
 - Makes training more time/resource consuming
 - Increases the likelihood of overfitting
- Problem is further compounded for color images
 - Each pixel in color image represented by 3 values (RGB color mode)
 - Since each pixel represented by 3 values, we say *channel* size is 3
 - Image represented by $64 \times 64 \times 3 = 12,288$ values (rows x columns x channels)
 - Number of weights is now $12,288 \times 500 = 6,144,000$

Inspiration for CNN

- In 1959 Hubel & Wiesel did an experiment to understand how visual cortex of brain processes visual info
 - Recorded activity of neurons in visual cortex of a cat
 - While moving a bright line in front of the cat
- Some cells fired when bright line is shown at a particular angle/location
 - Called these *simple* cells
- Other cells fired when bright line was shown regardless of angle/location
 - Seemed to detect movement
 - Called these *complex* cells
- Seemed complex cells receive inputs from multiple simple cells
 - Have an hierarchical structure
- Hubel and Wiesel won Noble prize in 1981

Inspiration for CNN

- Inspired by complex/simple cells, Fukushima proposed *Neocognitron* (1980)
 - Hierarchical neural network used for handwritten Japanese character recognition
 - First CNN, had its own training algorithm
- In 1989, LeCun proposed CNN that was trained by backpropagation
- CNN got popular when outperformed other models at ImageNet Challenge
 - Competition in object classification/detection
 - On hundreds of object categories and millions of images
 - Run annually from 2010 to present
- Notable CNN architectures that won ImageNet challenge
 - AlexNet (2012), ZFNet (2013), GoogLeNet & VGG (2014), ResNet (2015)

