

Chapter 3

Memory

These slides support chapter 3 of the book

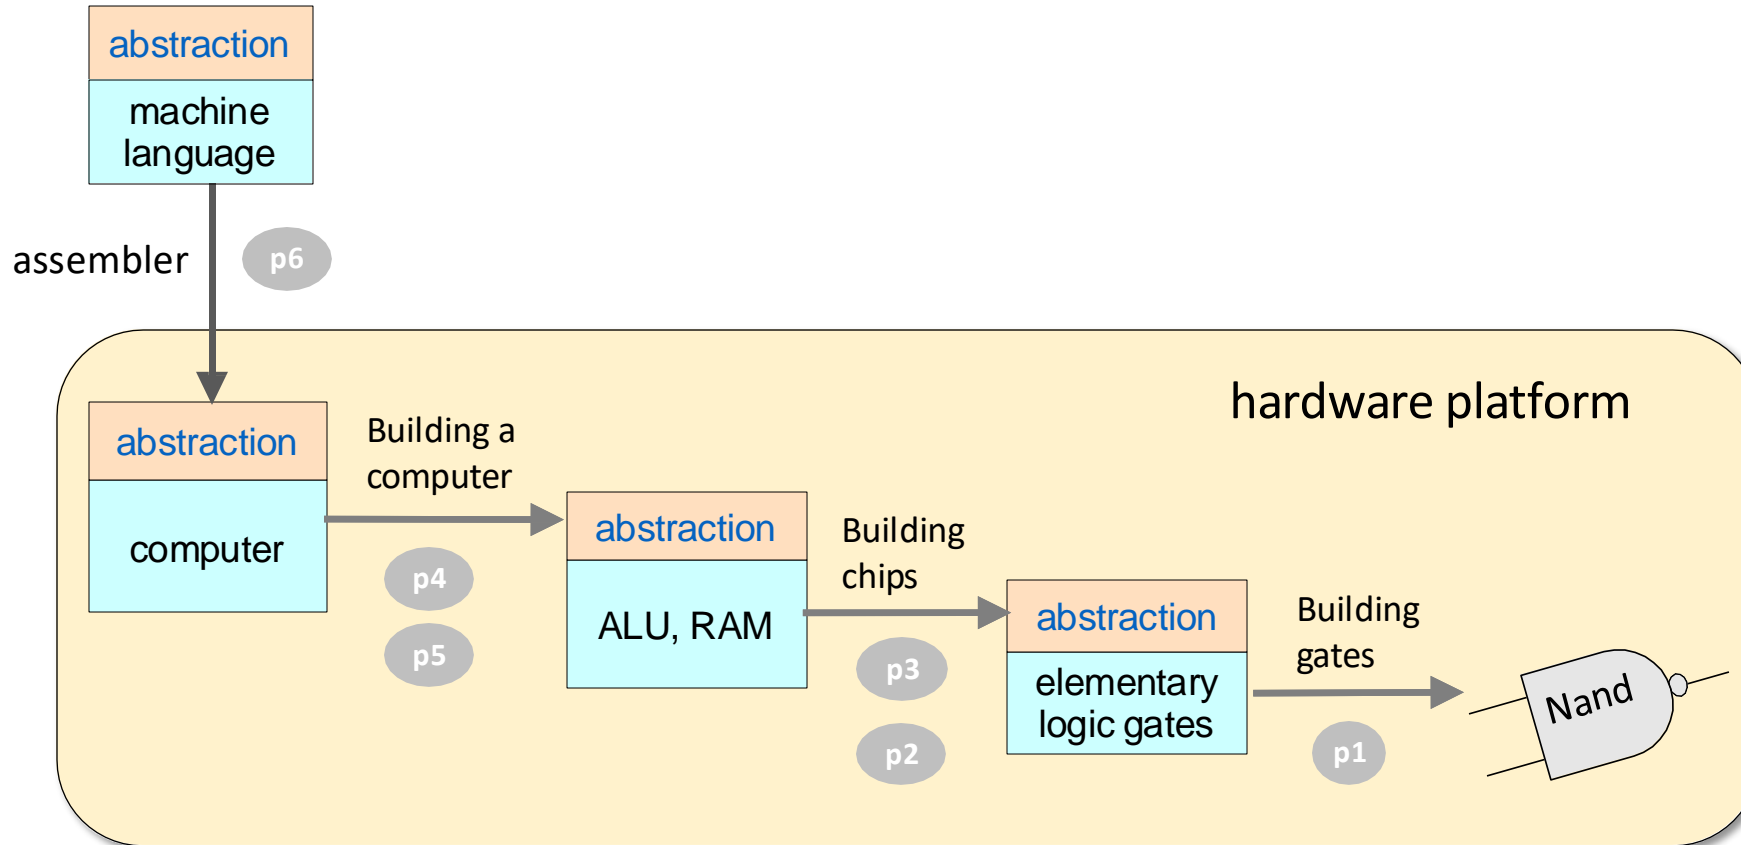
The Elements of Computing Systems

(1st and 2nd editions)

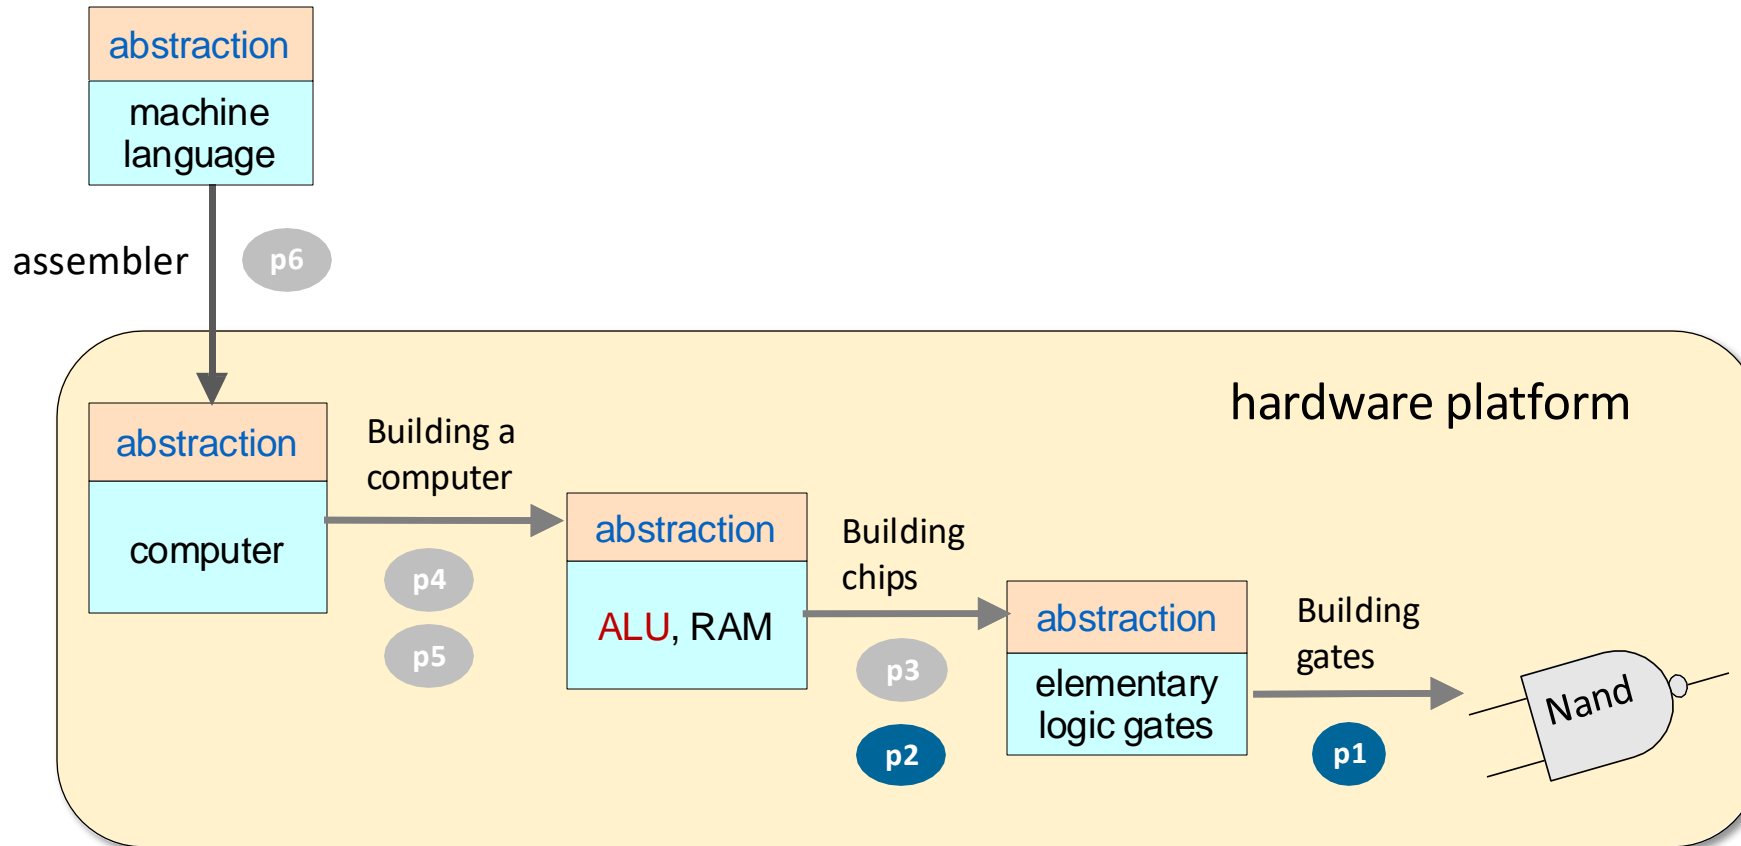
By Noam Nisan and Shimon Schocken

MIT Press

Nand to Tetris Roadmap: Hardware



Nand to Tetris Roadmap: Hardware

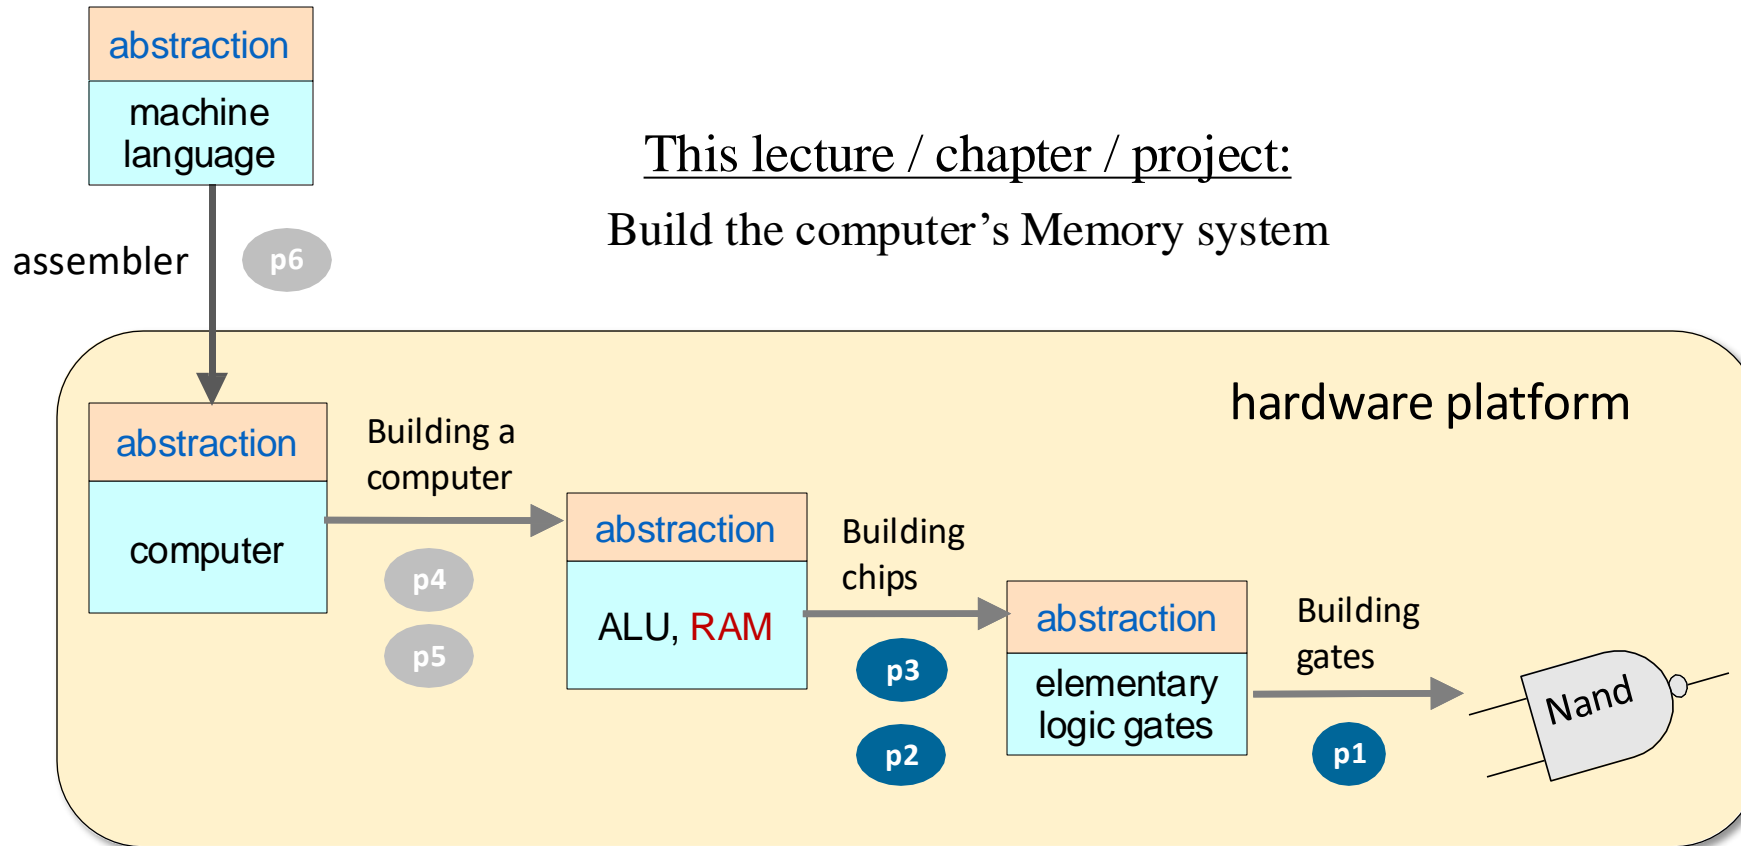


Project 1: Build basic logic gates

Project 2: Build the ALU

Nand to Tetris Roadmap: Hardware

This lecture / chapter / project:
Build the computer's Memory system



Project 1: Build basic logic gates

Project 2: Build the ALU

A common theme in computer science

- We present a simple model (the simpler, the better)
- We explore the model's power:
 - What the model can do
 - What it cannot do
- We then extend the model, to make it more powerful

Case in point:

Logic gates.

Logic gates

Model: And, Or, Not, ...

- Simple, and powerful:
Logic gates can realize any Boolean function, and can be combined to form powerful chips, like an ALU
- But, as a *general model of computation*, logic gates fall short

Limitations


- Logic gates cannot store information (bits) over time
- Feedback loops are not allowed: A chip's output cannot serve as its input
- Logic gates can handle only inputs of a fixed size.
For example, we can build an Or3 gate, and an Or4 gate, and so on, but we cannot build a single gate that computes Or for any given number of inputs

Extension

Allow logic gates to be sensitive to the progression of *time*.

Time-independent logic

- So far we ignored *time*
- The chip's inputs were just “sitting there” – fixed and unchanging
- The chip's output was a function (“combination”) of the current inputs, and the current inputs only
- This style of gate logic is sometimes called:
 - *time-independent logic*
 - *combinational logic*
- All the chips that we discussed and developed so far were combinational



ALU: The “topmost”
combinational chip

Hello, time

Software needs:

- The hardware must be able to remember things, over time:
- The hardware must be able to do things, one at a time (sequentially):

Example (variables):


```
x = 17
```

Example (iteration):

```
for i in range(0, 10):  
    print(i)
```

Hardware needs:

- The hardware must handle the *physical time delays* associated with *computing* and *moving* data from one chip to another.



Hello, time

Software needs:

- The hardware must be able to remember things, over time:
- The hardware must be able to do things, one at a time (sequentially):

Example (variables):

```
x = 17
```


Example (iteration):

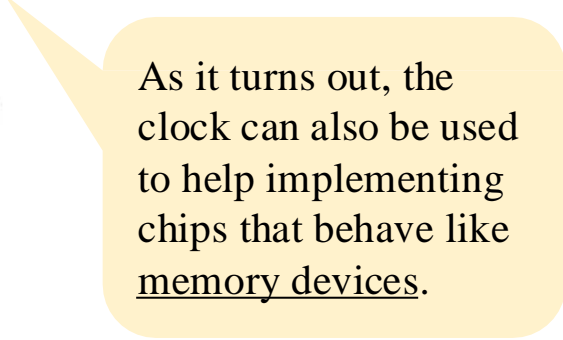
```
for i in range(0, 10):  
    print(i)
```

Hardware needs:

- The hardware must handle the *physical time delays* associated with *computing* and *moving* data from one chip to another.

It will take some time before 7 and 11 will settle down in the input ports, and before the sum $7 + 11$ will stabilize. Till then, the ALU will output nonsense.





Memory

Memory: The faculty of the brain by which data or information is encoded, stored, and retrieved when needed.

It is the *retention of information over time* for the purpose of influencing future action (Wikipedia)

Memory is time-based:


We remember *now* what was committed to memory *earlier*.



*It's a poor sort of memory
that only works backwards.*

-Lewis Carroll, through the White Queen

Memory



Basic abstractions:

- “Loading” a value
- “Storing” a value



loading

storing

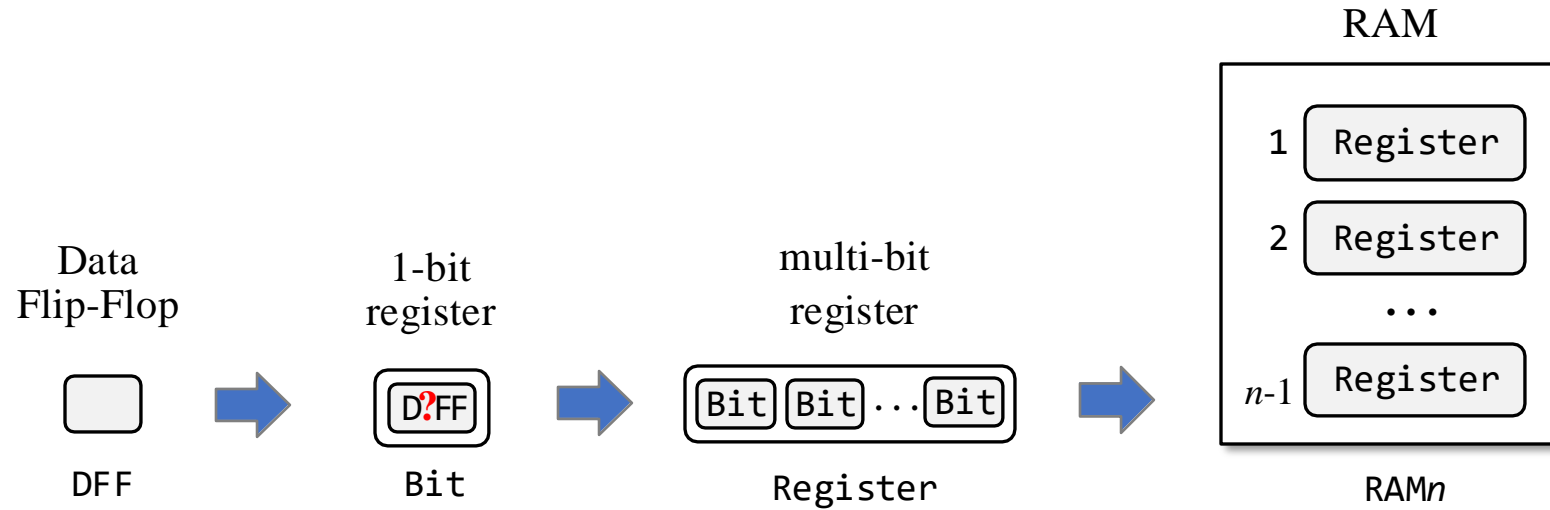
$x = 21, 21, 21, 21, 21, 21, 21, \dots$

loading

storing

The challenge: Building chips that realize this functionality.

Memory



The challenge: Building chips that realize this functionality.

Chapter 3: Memory

Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

Chapter 3: Memory

Abstraction

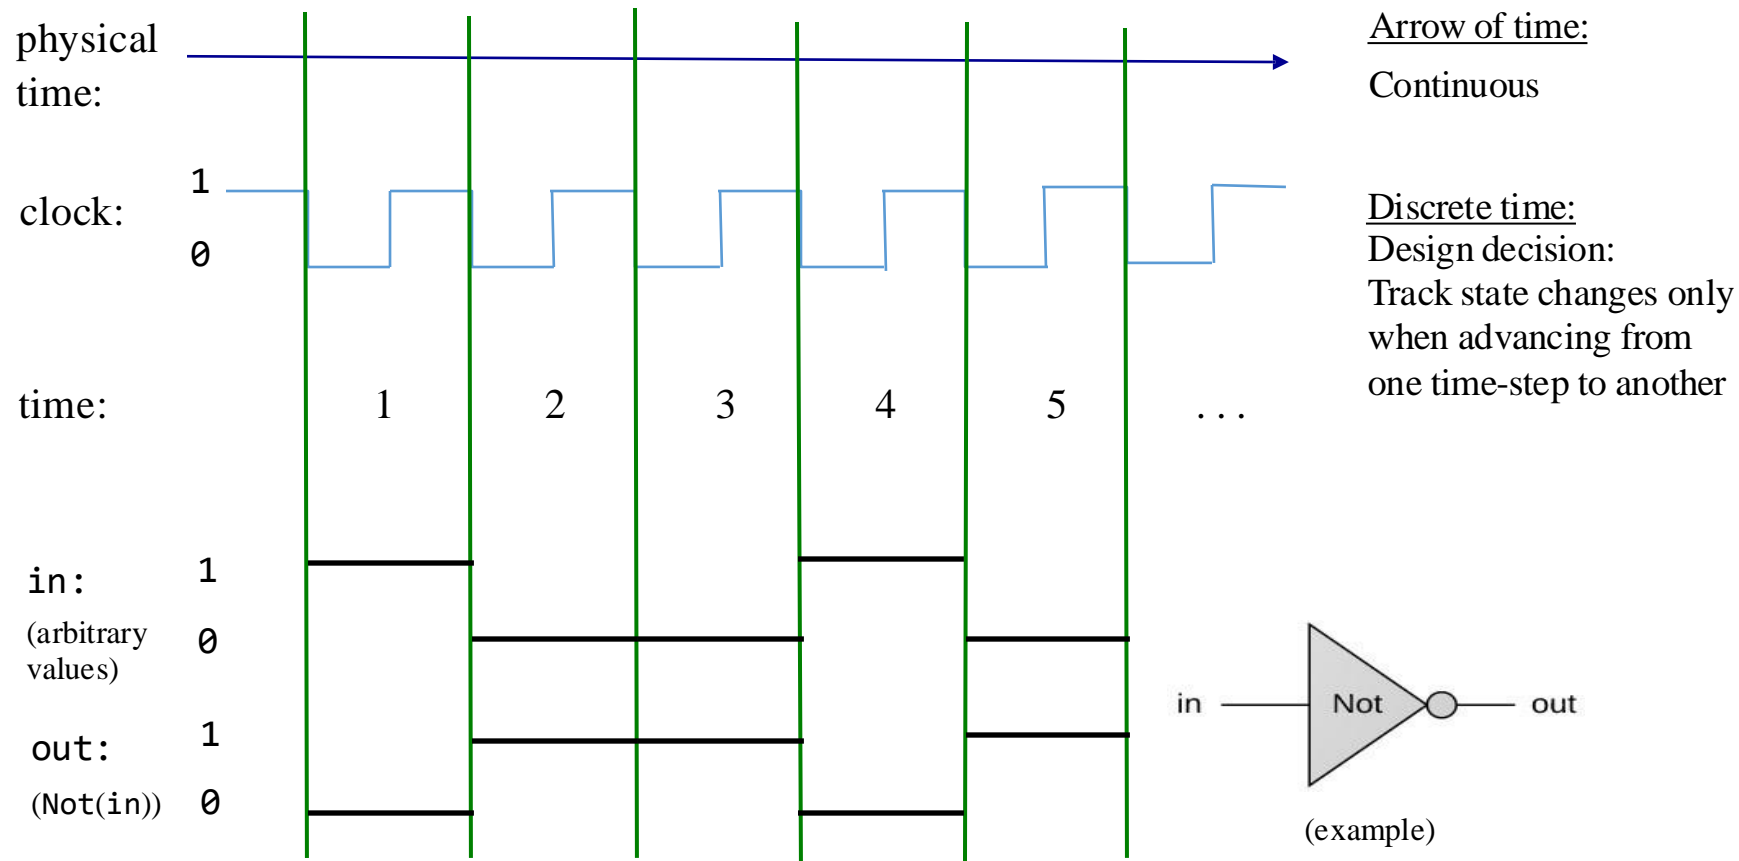
 Representing time

- Clock
- Registers
- RAM
- Counters

Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

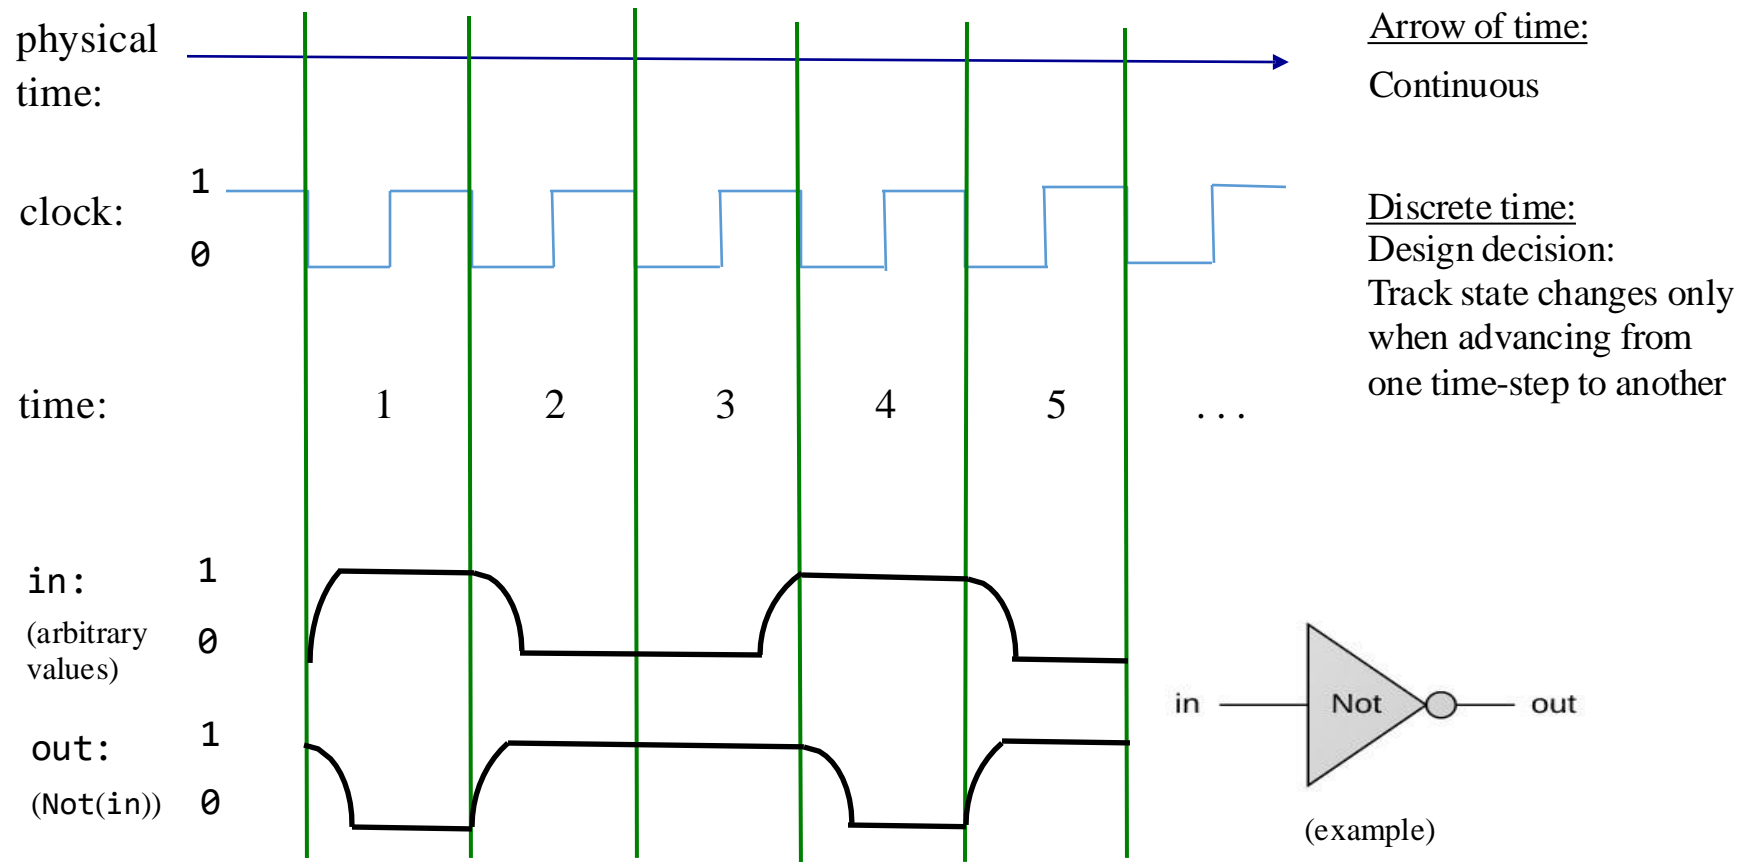
Chip behavior over time (example: Not gate)



Desired / idealized behavior of the in and out signals:

That's how we *want* the hardware to behave


Chip behavior over time (example: Not gate)



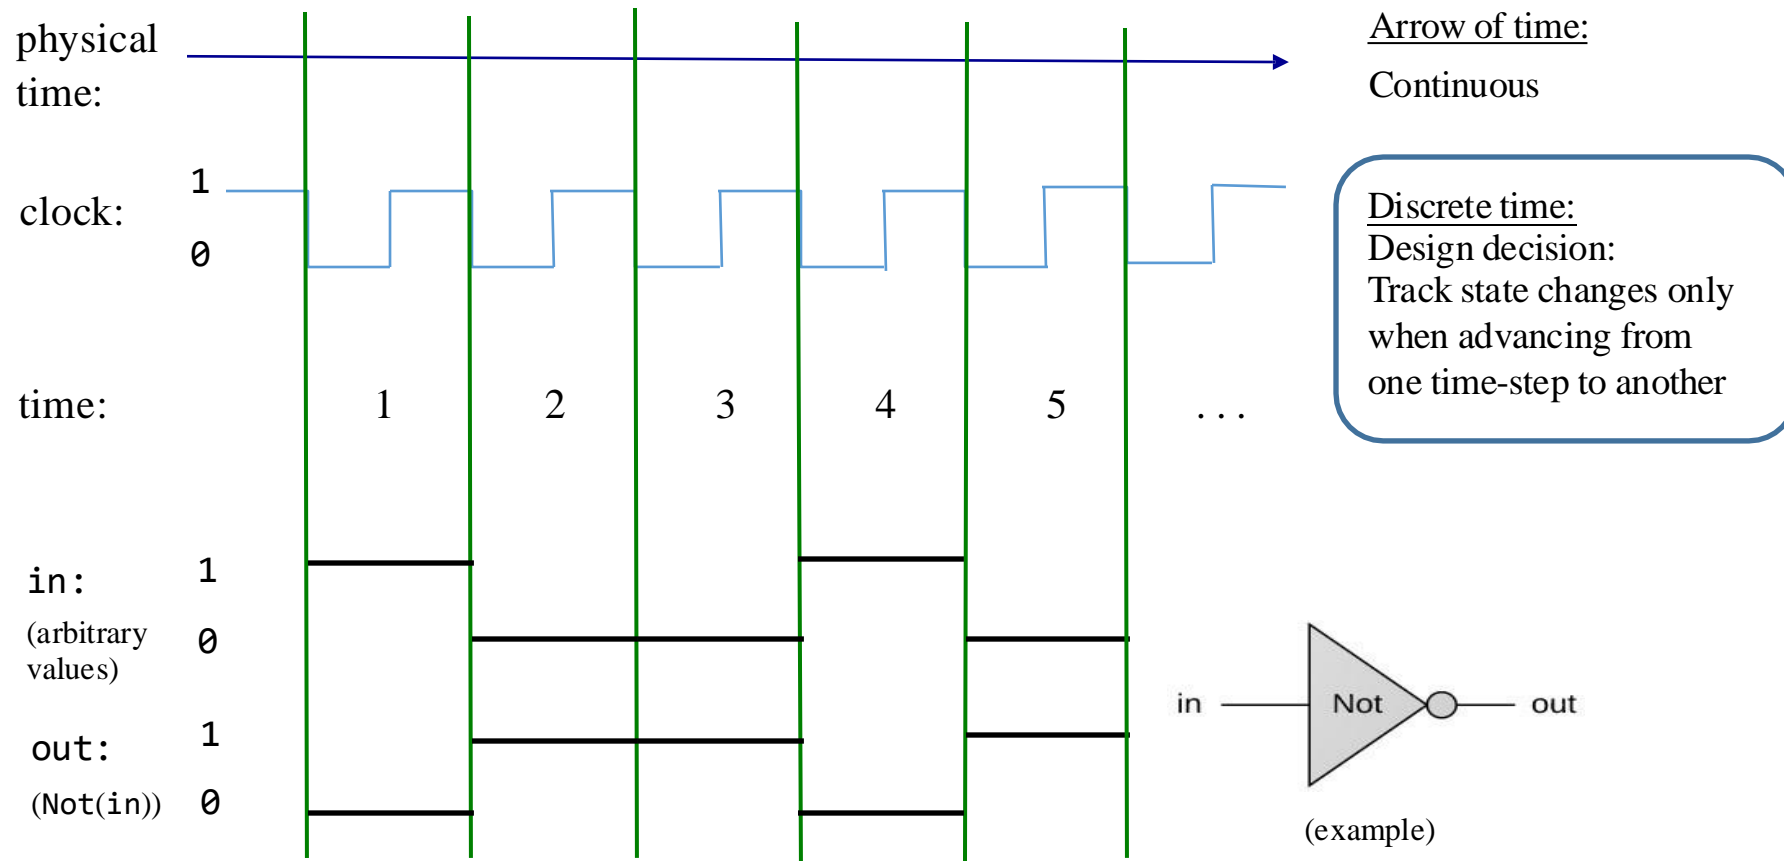
Actual behavior of the in and out signals:

Influenced by physical time delays

Chip behavior over time (example: Not gate)



Chip behavior over time (example: Not gate)



Resulting effect:

- Combinational chips react “immediately” to their inputs
- Facilitated by the decision to track changes only at cycle ends

Chapter 3: Memory

Abstraction



Representing time




Clock

- Registers
- RAM
- Counters

Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines


Clock: Simulated implementation



Interactive simulation

A clock icon can be used to generate a sequence of tick-tock signals:

0, 0+, 1, 1+, 2, 2+, 3, 3+, ...




Script-based simulation

“tick” and “tock” commands
can be used to advance the clock:

```
...  
// Sets inputs, advances the clock, and  
// writes output values as it goes along.  
set in 19,  
set load 1,  
tick,  
output,  
tock,  
output,  
tick, tock,  
output,  
...
```

Clock: Physical implementation




Physical clock

- An *oscillator* is used to deliver an ongoing train of “tick/tock” signals



“1 MHz electronic oscillator circuit which uses the resonant properties of an internal quartz crystal to control the frequency. Provides the clock signal for digital devices such as computers.” (Wikipedia)

- The oscillator’s output is connected to all the time-based (clocked) chips in the computer



Chip diagram convention:
A triangle icon represents a clock signal input

Chapter 3: Memory

Abstraction

✓ Representing time

✓ Clock

➡ Registers

- RAM
- Counters

Implementation


- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

Combinational logic / Sequential logic

Combinational logic:

The output depends on the current inputs


The clock is used to stabilize outputs



Sequential logic:

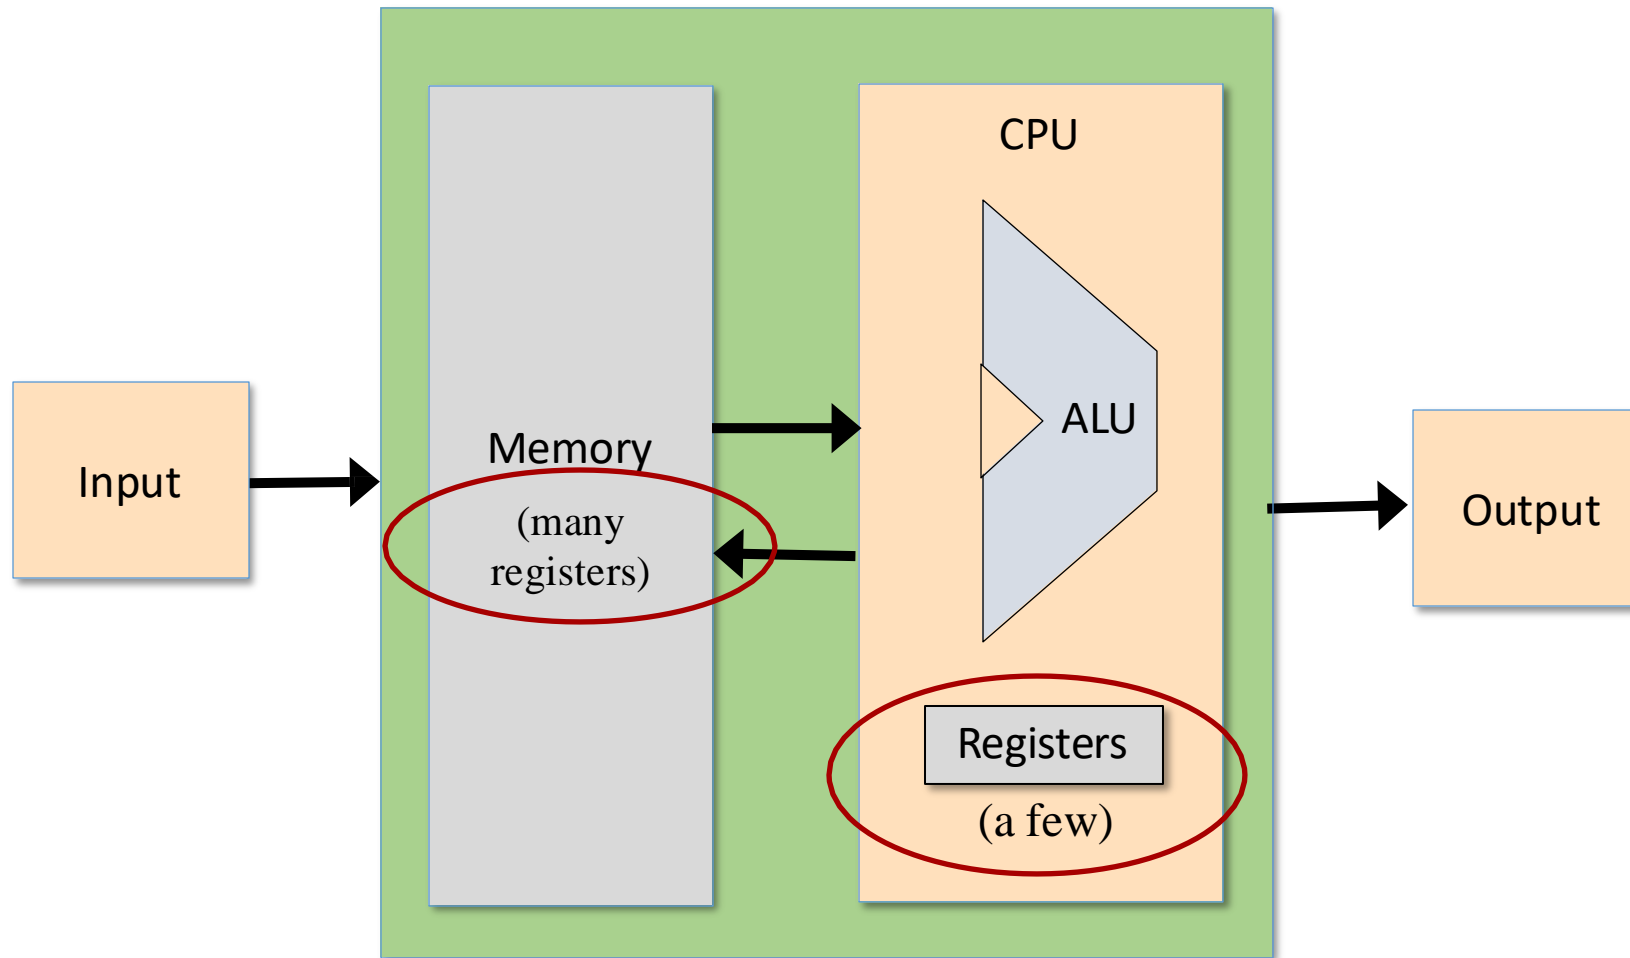
The output depends on:

- Previous inputs
- Current inputs (optionally)





This behavior can be used to build chips designed to maintain state: **Registers**.

Registers



Computer Architecture

Registers



Designed to:

- “Store” / “remember” / “maintain” / “persist” a value , until...
- “Instructed” to “load”, and then “store”, another value.

time:

$x = 17, 17, 17, 17, 17, 17, 17, \dots, 17$

loading


maintaining state

$x = 21, 21, 21, 21, 21, 21, \dots, 21$


loading

maintaining state


1-Bit Register



if $\text{load}(t-1)$ then $\text{out}(t) = \text{in}(t-1)$ else
 $\text{out}(t) = \text{out}(t-1)$



1-Bit Register



if $\text{load}(t-1)$ then $\text{out}(t) = \text{in}(t-1)$ else
 $\text{out}(t) = \text{out}(t-1)$

Usage:

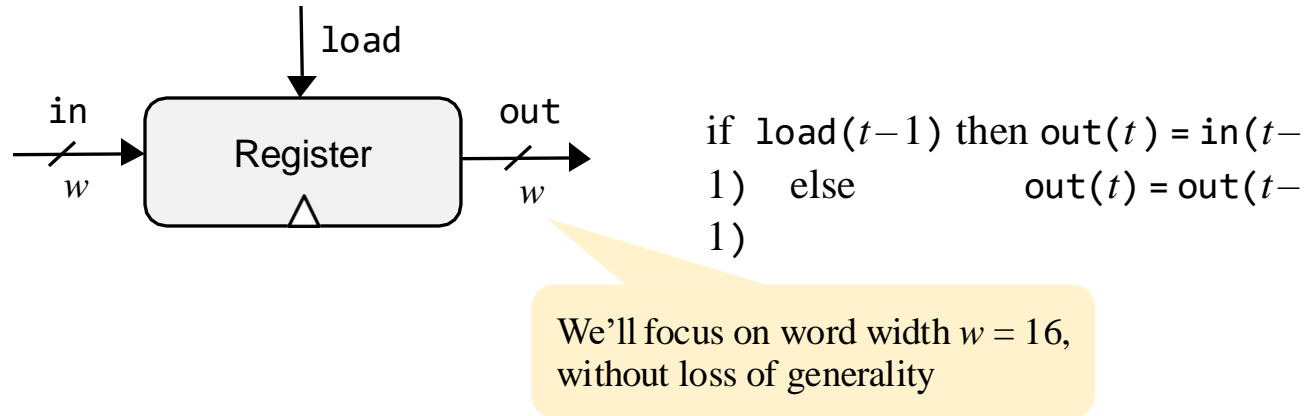
To read:

probe out (out always emits the register's state)

To write:

set $\text{in} = v$ Result: The register's state becomes v ;
set $\text{load} = 1$ From the next time-step onward, out will emit v

Multi-bit Register



Load / store behavior: Exactly the same as a 1-bit register

Read / write usage: Exactly the same as a 1-bit register



Chapter 3: Memory

Abstraction

✓ Representing time

✓ Clock

✓ Registers


➡ RAM

- Counters


Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

Computer architecture



RAM



Practice question:

Suppose that the RAM size $n = 8$ registers.
What should be the value of k ?


Answer:

$$k = \log_2 n$$

Abstraction: A sequence of n addressable, w -bit registers, with addresses 0 to $n-1$

Word width: Typically 16, 32, 64 bits (Hack computer: $w = 16$)

RAM



Behavior

If $load == 0$, the RAM maintains its state

If $load == 1$, $RAM[address]$ is set to the value of in

The loaded value will be emitted by out from the next time-step (cycle) onward, until the next load

(Only one RAM register is selected;
All the other registers are not affected)

Usage:

To read register i :

set $address = i$,

probe out (out always emits the value of $RAM[i]$)

To write v in register i :

set $address = i$,

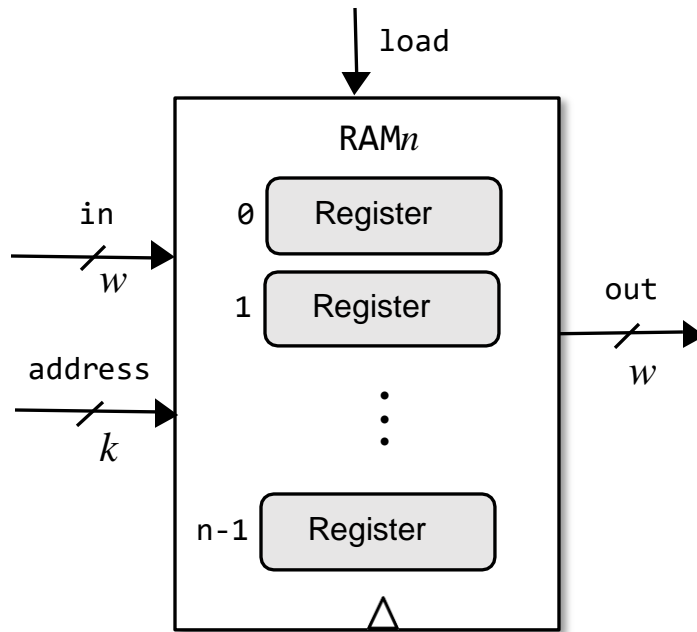
set $in = v$,

set $load = 1$

Result: $RAM[i] \leftarrow v$

From the next time-step onward, out will emit v

RAM



Why “Random Access Memory”?

Irrespective of the RAM size (n), every randomly selected register can be accessed “instantaneously”, at more or less the same speed.



Chapter 3: Memory

Abstraction

✓ Representing time

✓ Clock

✓ Registers


✓ RAM

➡ Counters

Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

Computer architecture



Counter

- Later in the course (chapter 5), we will see that the computer must keep track of which instruction should be fetched and executed next
- This task is regulated by a register typically called Program Counter
- We'll use the PC to store the address of the instruction that should be fetched and executed next
- The PC should support three abstractions:

Reset: fetch the first instruction

PC = 0


Next: fetch the next instruction

PC++

Goto: fetch instruction n

PC = n

Counter



| | |
|--------------------------|---------------------------------------|
| if $\text{reset}(t)$ | $\text{out}(t+1) = 0$ |
| else if $\text{load}(t)$ | $\text{out}(t+1) = \text{in}(t)$ |
| else if $\text{inc}(t)$ | $\text{out}(t+1) = \text{out}(t) + 1$ |
| else | $\text{out}(t+1) = \text{out}(t)$ |

Usage:

To read:

probe out

To set:

set in to v ,
assert load,
set the other control bits to 0

To reset:

assert reset,
set the other control bits to 0

To count:

assert inc,
set the other control bits to 0



Chapter 3: Memory

Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines



Chapter 3: Memory

Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

Implementation





Data Flip Flop

- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

DFF

Data Flip Flop (aka *latch*)



The most elementary sequential gate: Outputs the input in the previous time-step



From DFF to a 1-bit register


Data Flip Flop (aka latch)

The most elementary sequential gate: Outputs the input in the previous time-step



How can we “load” and then “maintain” a value (0 or 1) over time, without having to feed the value in every cycle?

From DFF to a 1-bit register



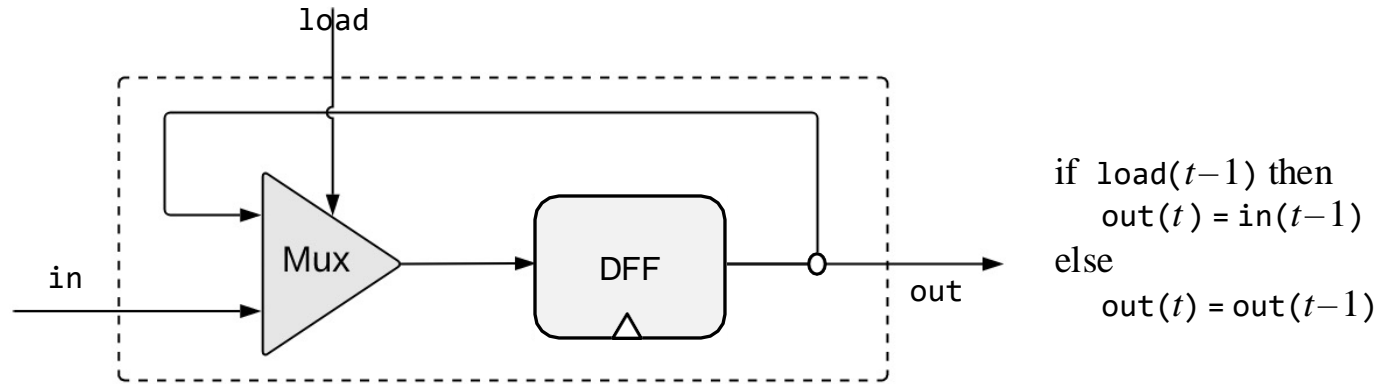
We have to realize a “loading” behavior and a ”storing” behavior, and be able to select between these two states

How can we “load” and then “maintain” a value (0 or 1) over time, without having to feed the value in every cycle?

From DFF to a 1-bit register

1-bit Register

Stores one bit
over time



We have to realize a “loading” behavior and a ”storing” behavior,
and be able to select between these two states


Behavior

if $\text{load} == 1$ the register’s value becomes in
else the register maintains its current value

Register

1-bit Register


Stores one bit
over time



Register

1-bit Register

Stores one bit
over time




zoom out...

```
if load( $t-1$ ) then
    out( $t$ ) = in( $t-1$ )
else
    out( $t$ ) = out( $t-1$ )
```

Register

1-bit Register


Stores one bit
over time



if $\text{load}(t-1)$ then
 $\text{out}(t) = \text{in}(t-1)$
else
 $\text{out}(t) = \text{out}(t-1)$

w-bit Register:

Stores w bits
over time






Partial diagram, showing
some of the chip-parts,
without connections

Chapter 3: Memory

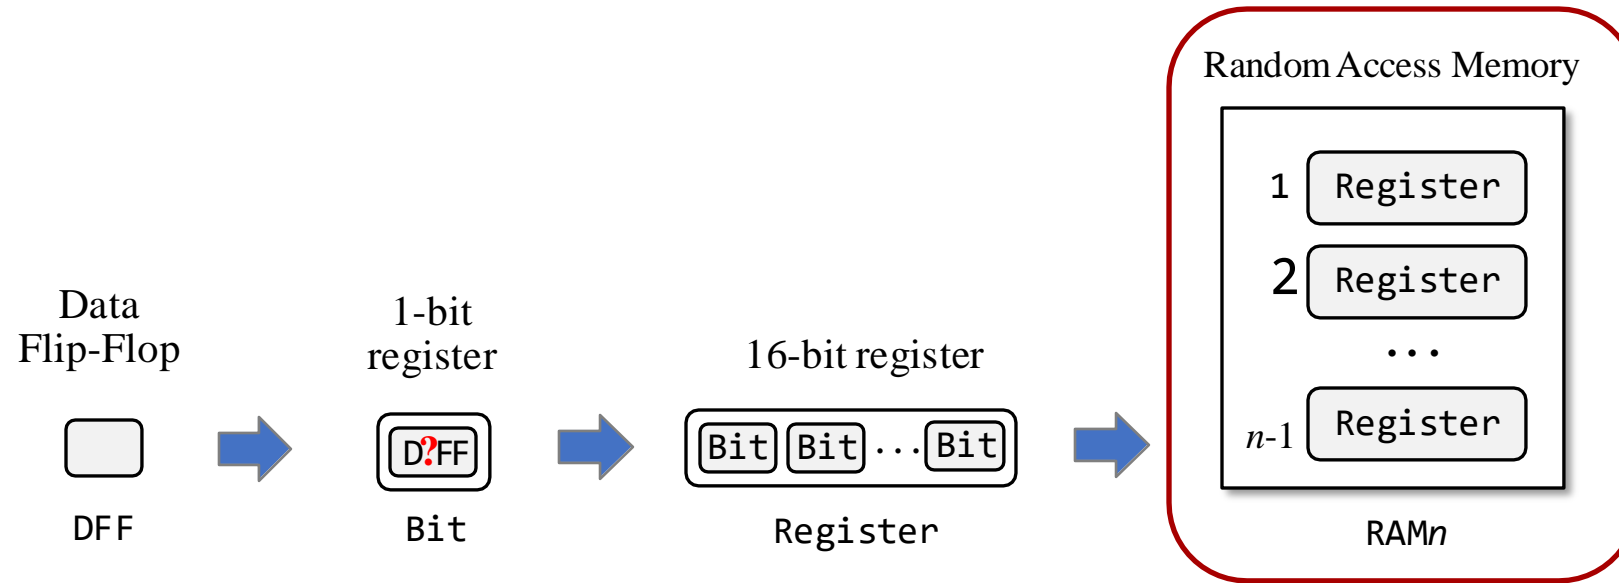
Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

Implementation

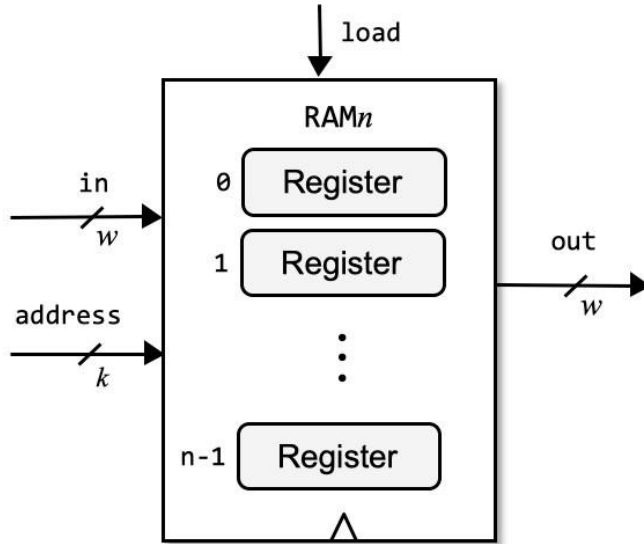
-  Data Flip Flop
-  Registers
-  RAM
 - Project 3: Chips
 - Project 3: Guidelines

Memory hierarchy



RAM: Abstraction

RAM of n registers:



Usage: **To read register i :**

set address = i ,

probe out (out always emits the state of $RAM[i]$)

To write v in register i :

set address = i ,

set in = v ,

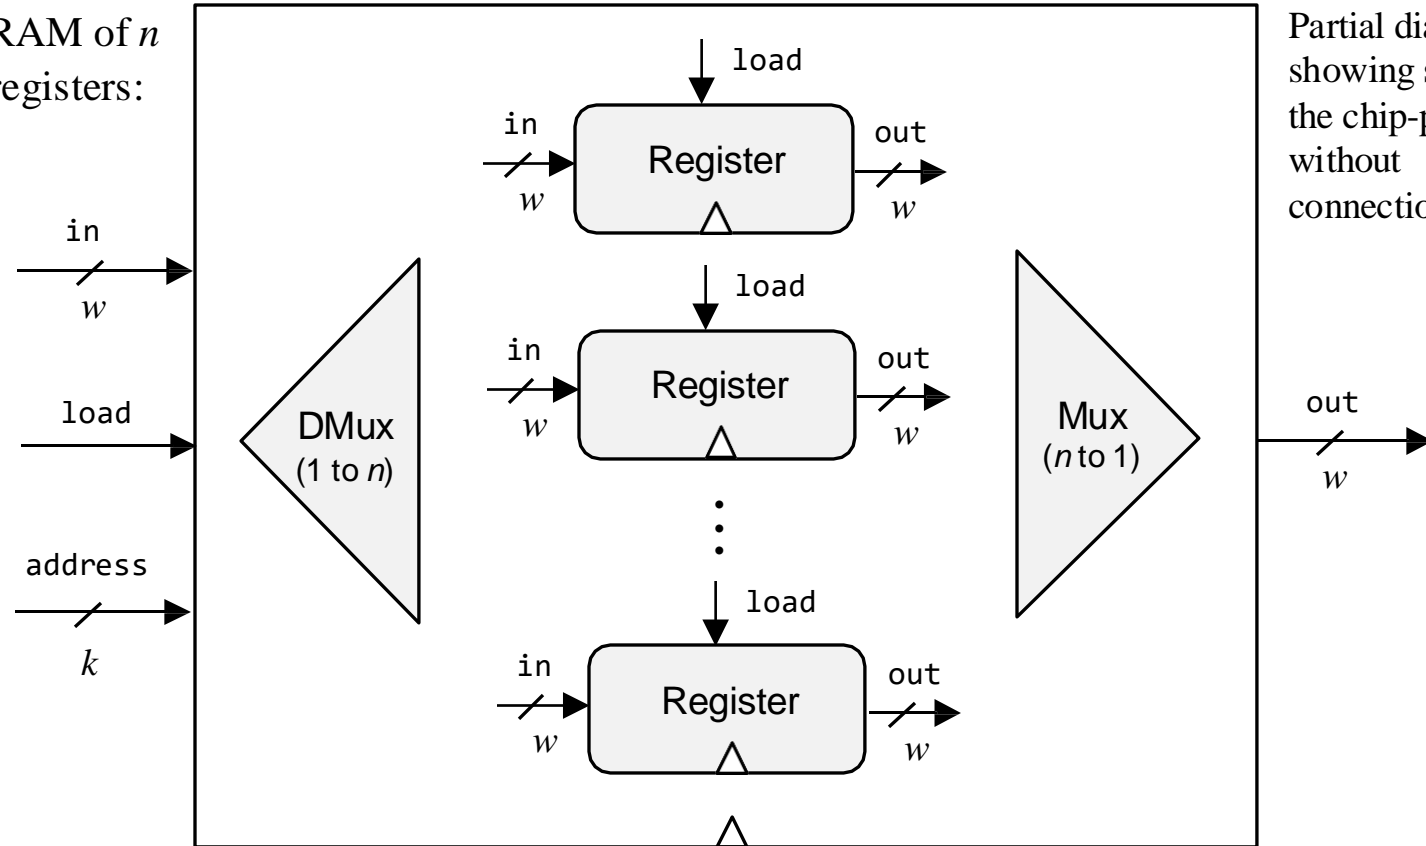
set load = 1

Result: $RAM[i] \leftarrow v$

From the next time-step onward, out emits v

RAM: Implementation

RAM of n registers:



Partial diagram,
showing some of
the chip-parts,
without
connections

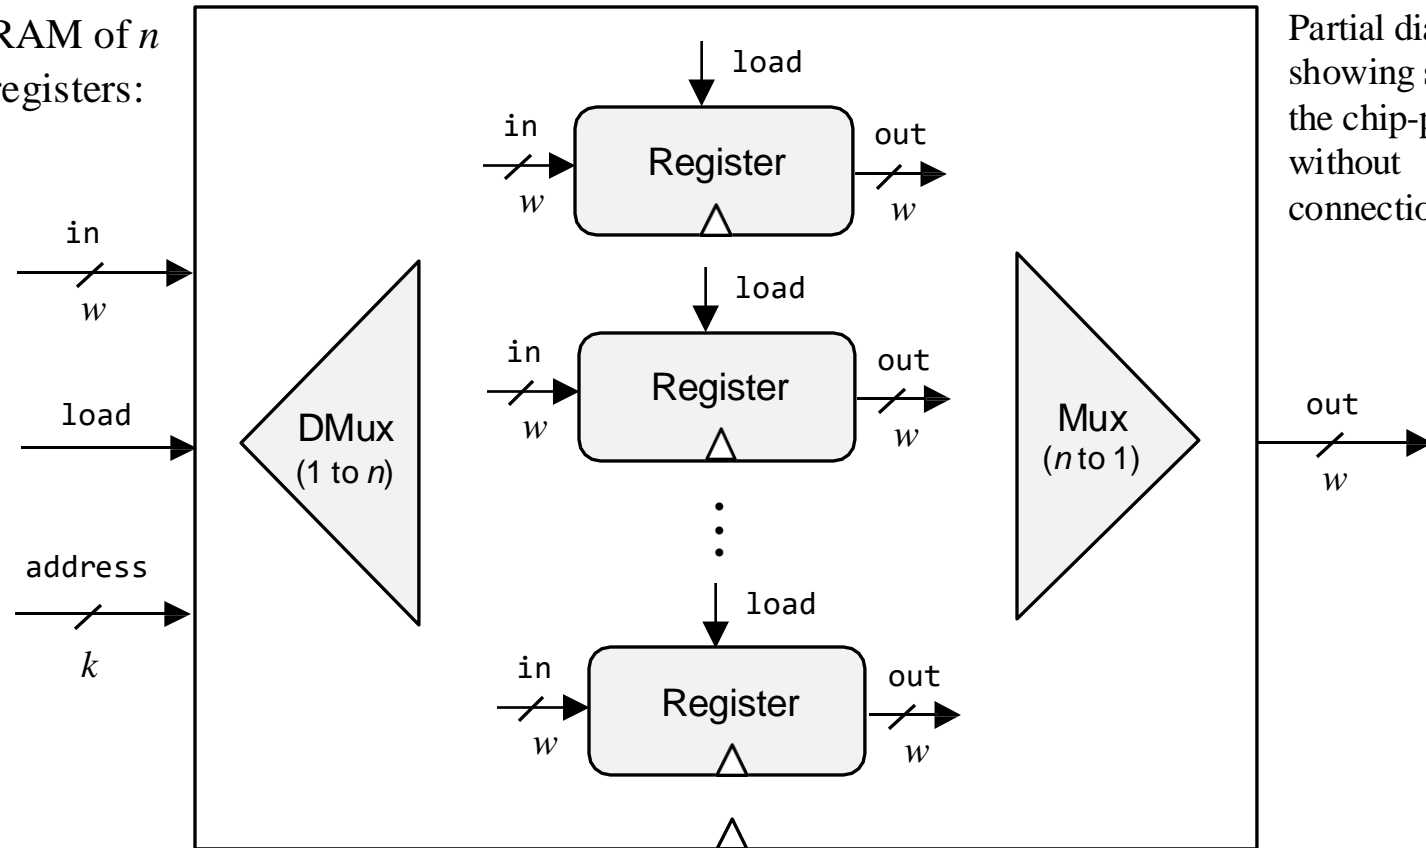
Reading: Can be realized using a Mux

Writing: Can be realized using a DMux

Connections?
You figure it out

RAM: Implementation

RAM of n registers:

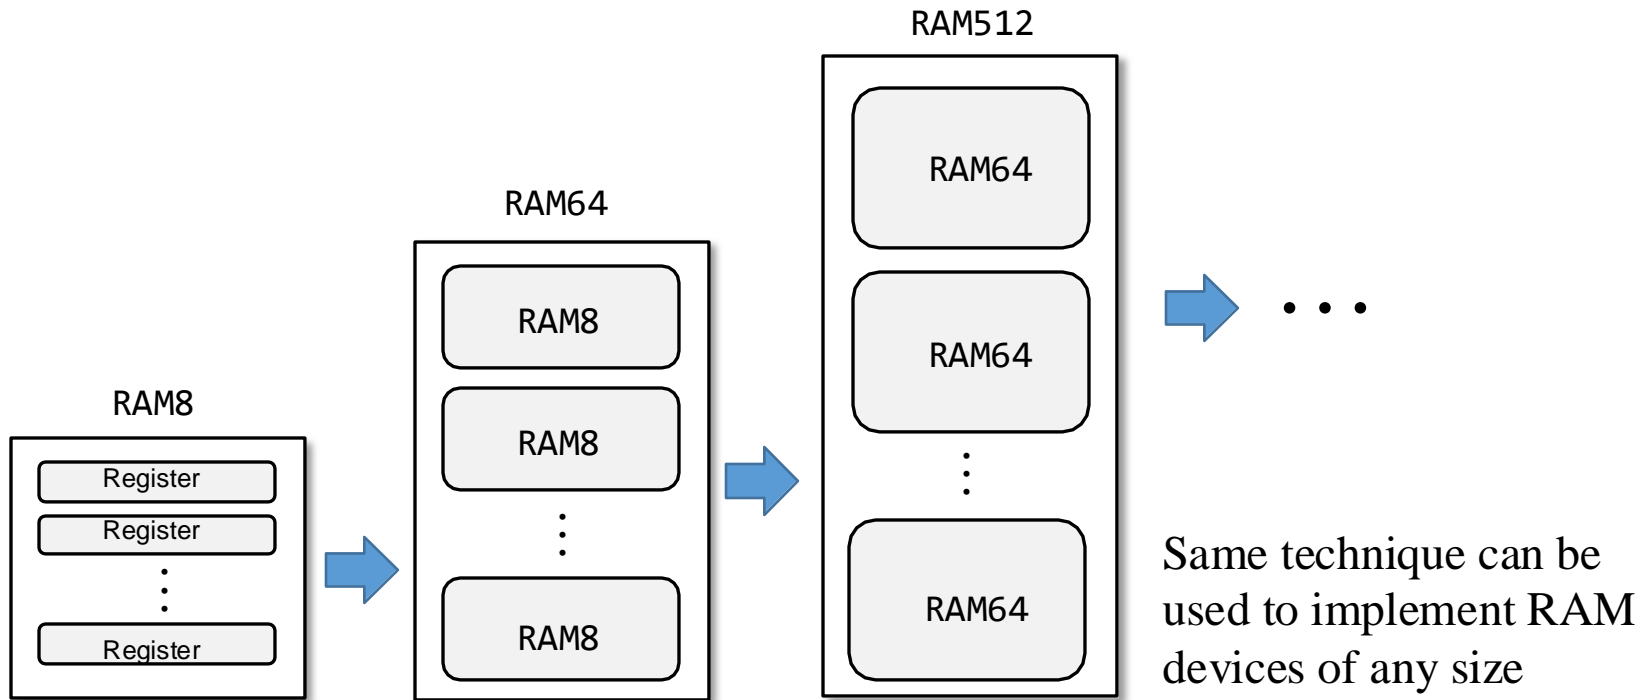


Partial diagram,
showing some of
the chip-parts,
without
connections

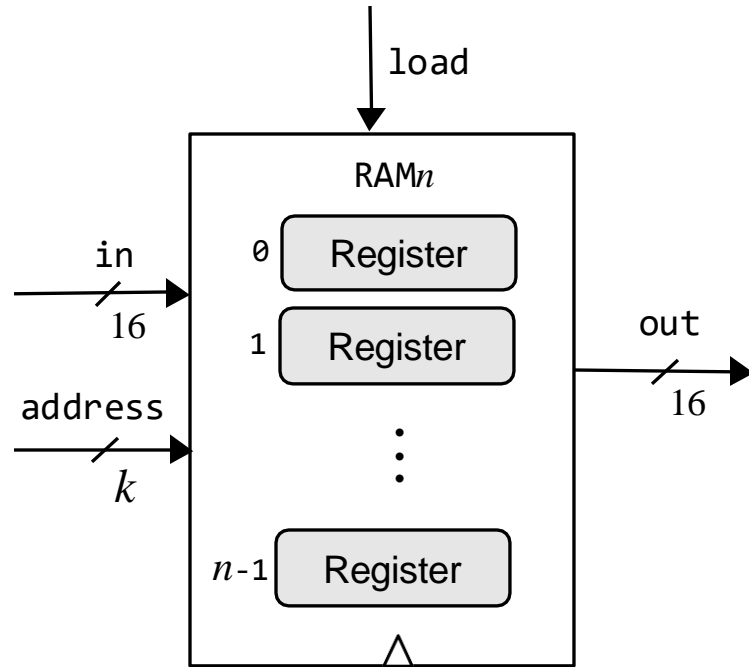
Observations

- The addressing/selection/reading logic is *combinational*
- The writing logic is (i) *sequential (clocked)*
(ii) embedded in the Register logic.

RAM: Implementation



Hack RAM



A family of 16-bit RAM chips:

| chip name | n | k |
|-----------|-------|-----|
| RAM8 | 8 | 3 |
| RAM64 | 64 | 6 |
| RAM512 | 512 | 9 |
| RAM4K | 4096 | 12 |
| RAM16K | 16384 | 14 |

Why these particular RAM chips?

Because that's what we need for building the Hack computer.

Chapter 3: Memory

Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

Implementation

- ✓ Data Flip Flop
- ✓ Registers
- ✓ RAM
- ➡ Project 3: Chips
 - Project 3: Guidelines

Project 3

Given:

- All the chips built in projects 1 and 2
- Data Flip-Flop (built-in DFF gate)


Build:



Bit

- Register
- PC
- RAM8
- RAM64
- RAM512
- RAM4K
- RAM16K

1-bit Register




Bit.hdl

```
/** 1-bit register:
    if load(t-1) then out(t) = in(t-1)
    else
        out(t) = out(t-1) */


CHIP Bit {
    IN in, load;
    OUT out;

    PARTS:
    // Put your code here:
}
```



Implementation tip:
Follow the chip diagram

16-bit Register

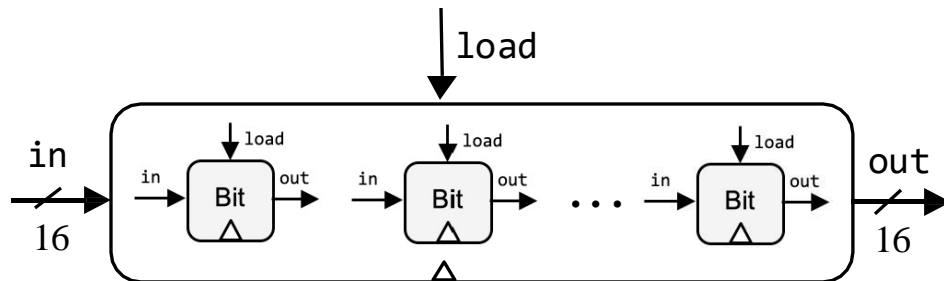


Register.hdl

```
/** 1-bit register:
    if load(t-1) then out(t) = in(t-1)
    else out(t) = out(t-1) */

CHIP Bit {
    IN in[16], load;
    OUT out[16];


    PARTS:
    // Put your code here:
}
```



Partial diagram, showing some of
the chip-parts, without connections

Implementation tip:
Follow the chip diagram

16-bit Counter



```
/**
  A 16-bit counter with control bits.
  if      reset(t - 1)  out(t) = 0           // resetting
  else if load(t - 1)   out(t) = in(t - 1)    // setting
  else if inc(t - 1)    out(t) = out(t - 1) + 1 // incrementing
  else                  out(t) = out(t - 1)    // maintaining
*/

CHIP PC {
  IN in[16], load, inc, reset;
  OUT out[16];

  PARTS:
    // Put your code here:
}
```

Implementation tip: Can be built from a Register, an Incrementer, and Mux's

Project 3

Given

- All the chips built in projects 1 and 2
- Data Flip-Flop (built-in DFF gate)

Build the following chips

✓ Bit

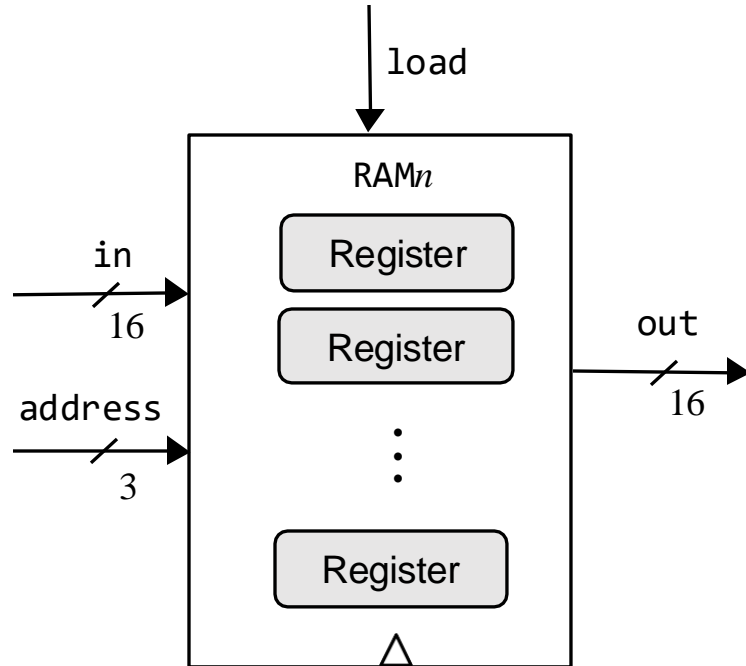
✓ Register

✓ PC

➡ RAM8

- RAM64
- RAM512
- RAM4K
- RAM16K

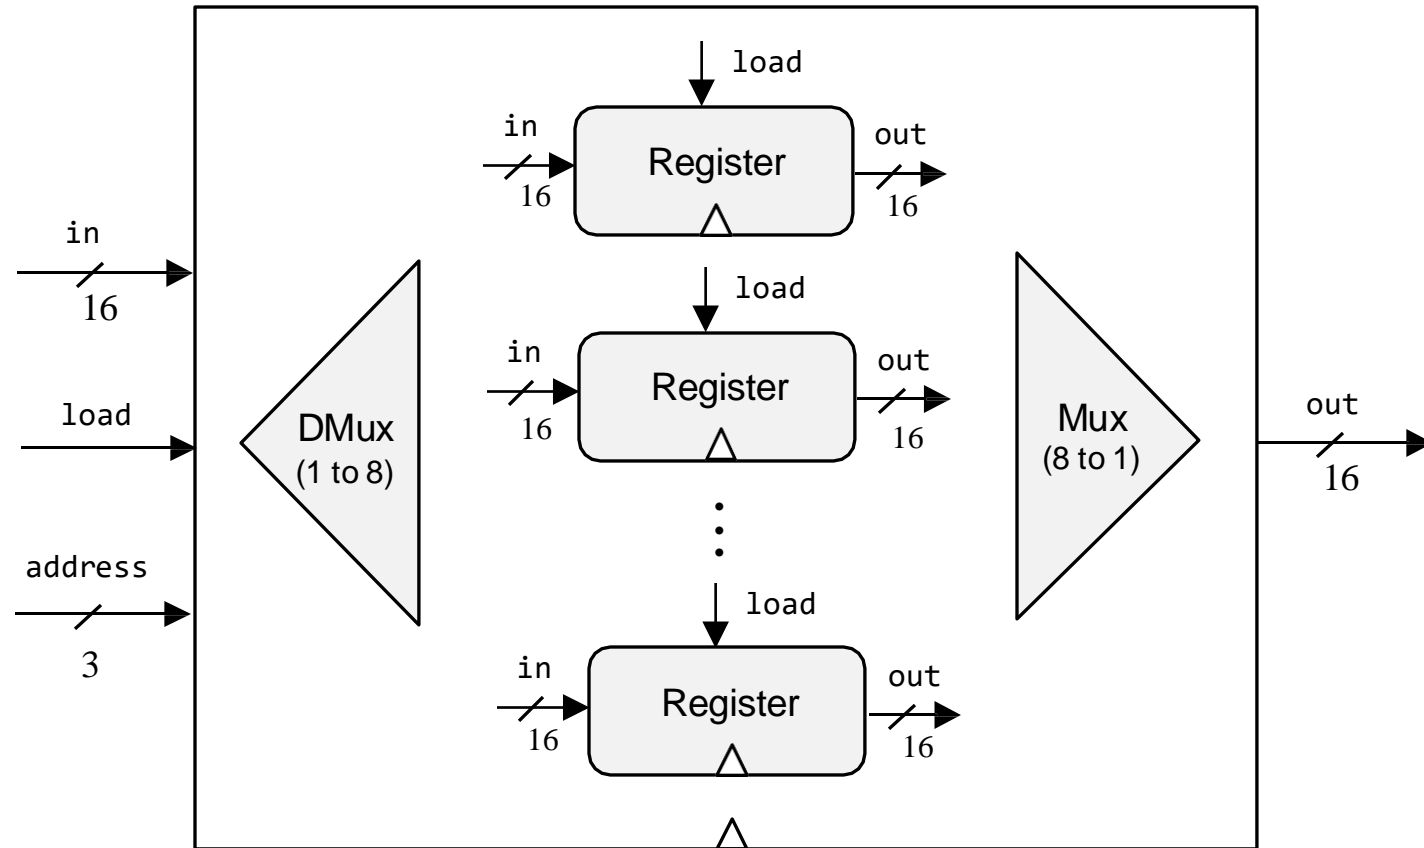
8-Register RAM: Abstraction



RAM8.hdl

```
/*  
  Let M stand for the state of the register  
  selected by address.  
  if load(t - 1) then {M = in(t), out(t) = M}  
  else  
    out(t) = M  
*/  
CHIP RAM8 {  
  IN in[16], load, address[3];  
  OUT out[16];  
  PARTS:  
    // Put your code here:  
}
```

8-Register RAM: Implementation



Partial diagram, showing some of the chip-parts, without connections

Implementation tip:

Follow the chip diagram

Project 3

Given

- All the chips built in projects 1 and 2
- Data Flip-Flop (built-in DFF gate)

Build the following chips

✓ Bit

✓ Register

✓ PC

✓ RAM8

• RAM64

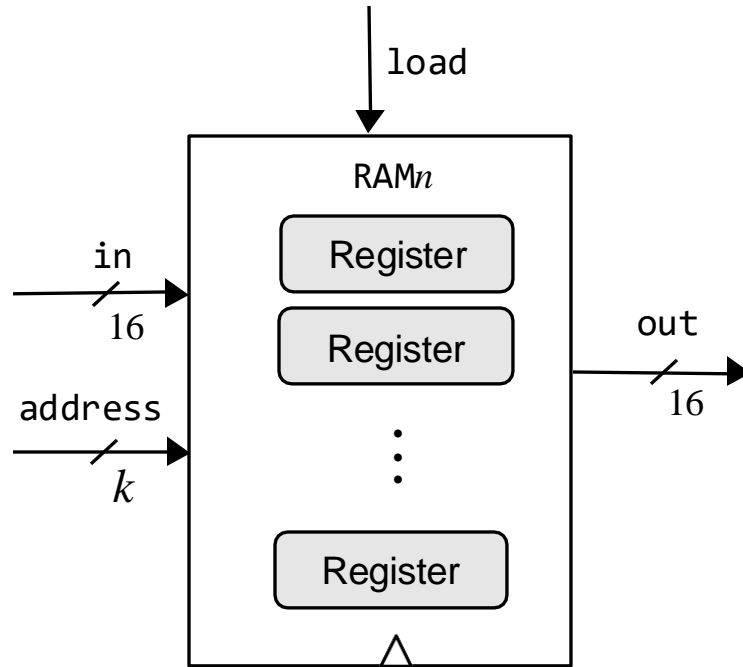
• RAM512

• RAM4K

• RAM16K

} A family of RAM chips

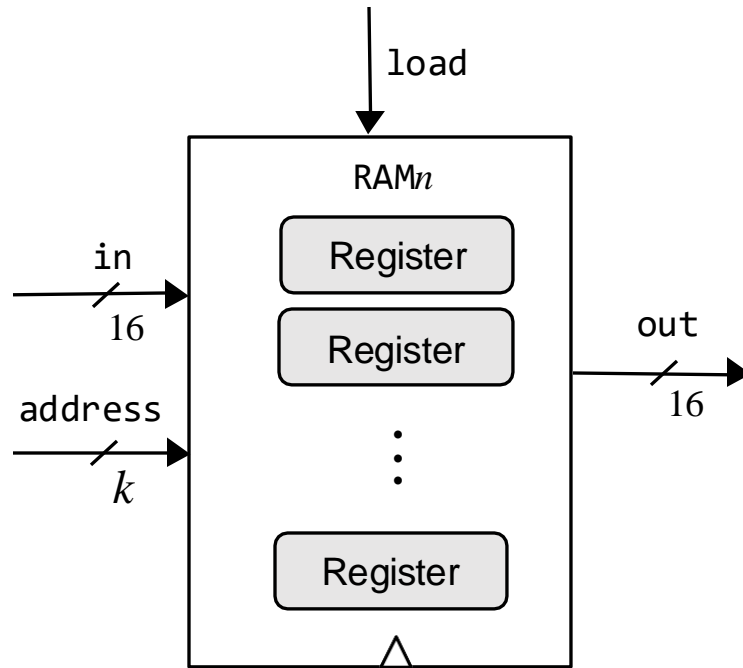
n -Register RAM



$RAM_n.hdl$

```
/*  
  Let M stand for the state of the register  
  selected by address.  
  if load(t - 1) then {M = in(t), out(t) = M}  
  else  
                                out(t) = M  
*/  
CHIP RAMn {  
  IN in[16], load, address[k];  
  OUT out[16];  
  PARTS:  
    // Put your code here:  
}
```

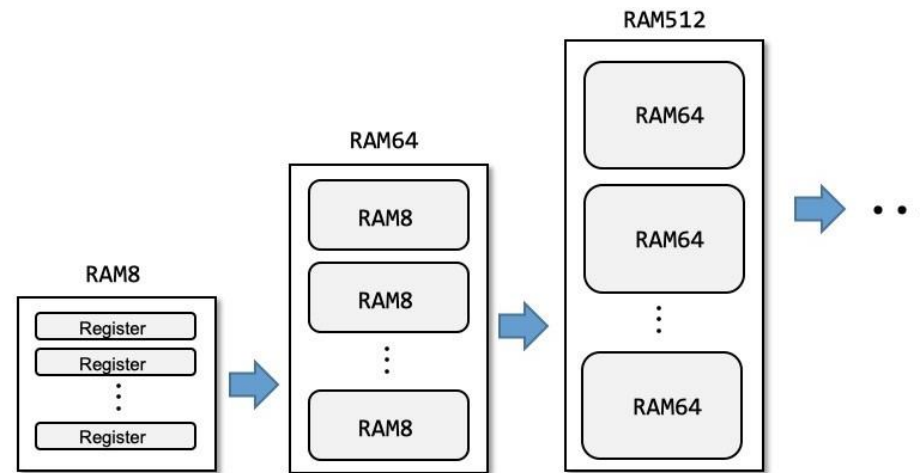

n -Register RAM



| chip name | n | k |
|-----------|-------|-----|
| RAM8 | 8 | 3 |
| RAM64 | 64 | 6 |
| RAM512 | 512 | 9 |
| RAM4K | 4096 | 12 |
| RAM16K | 16384 | 14 |

Implementation tips

- Think about the RAM's address input as consisting of two fields:
 - One field selects a RAM-part;
 - The other field selects a register within that RAM-part
- Use logic gates to effect this addressing scheme.








Chapter 3: Memory

Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

Implementation

-  Data Flip Flop
-  Registers
-  RAM
-  Project 3: Chips
-  Project 3: Guidelines

Project 3

From NAND to Tetris
Building a Modern Computer From First Principles
www.nand2tetris.org

Home
Prerequisites
Syllabus
Course
Book
Software
Terms
Papers
Talks
Cool Stuff
About
Team
Q&A

Project 3: Sequential Chips

Background

The computer's main memory, also called *Random Access Memory*, or *RAM*, is an addressable sequence of n -bit registers, each designed to hold an n -bit value. In this project you will gradually build a RAM unit. This involves two main issues: (i) how to use gate logic to store bits persistently, over time, and (ii) how to use gate logic to locate ("address") the memory register on which we wish to operate.

Objective

Build all the chips described in Chapter 3 (see list below), leading up to a *Random Access Memory* (RAM) unit. The only building blocks that you can use are primitive DFF gates, chips that you will build on top of them, and chips described in previous chapters.

Chips

| Chip (HDL) | Description | Test script | Compare file |
|------------|--------------------------------|--------------|--------------|
| DFF | Data Flip-Flop (primitive) | | |
| Bit | 1-bit register | Bit.tst | Bit.cmp |
| Register | 16-bit register | Register.tst | Register.cmp |
| RAM8 | 16-bit / 8-register memory | RAM8.tst | RAM8.cmp |
| RAM64 | 16-bit / 64-register memory | RAM64.tst | RAM64.cmp |
| RAM512 | 16-bit / 512-register memory | RAM512.tst | RAM512.cmp |
| RAM4K | 16-bit / 4096-register memory | RAM4K.tst | RAM4K.cmp |
| RAM16K | 16-bit / 16384-register memory | RAM16K.tst | RAM16K.cmp |
| PC | 16-bit program counter | PC.tst | PC.cmp |



All the necessary project 3 files are available in:
nand2tetris / projects / 03

Resources

Project 3 folder (.hdl, .tst, .cmp files): `nand2tetris/projects/03`

Tools

- Text editor (for completing the given .hdl stub-files)
- Hardware simulator: `nand2tetris/tools`

Guides

- [Hardware Simulator Tutorial](#)
- [HDL Guide](#)
- [Hack Chip Set API](#)

Best practice advice

- Implement the chips in the order in which they appear in the project guidelines
- If you don't implement some chips, you can still use their built-in implementations
- No need for “helper chips”: Implement / use only the chips we specified
- In each chip definition, strive to use as few chip-parts as possible
- You will have to use chips implemented in previous projects;
For efficiency and consistency's sake, use their built-in versions, rather than your own HDL implementations.

For technical reasons, the chips of project 3 are organized in two sub-folders named `projects/03/a` and `projects/03/b`

When writing and simulating the `.hdl` files, leave this folder structure as is.

That's It!
Go Do Project 3!

Chapter 3: Memory

Abstraction

- Representing time
- Clock
- Registers
- RAM
- Counters

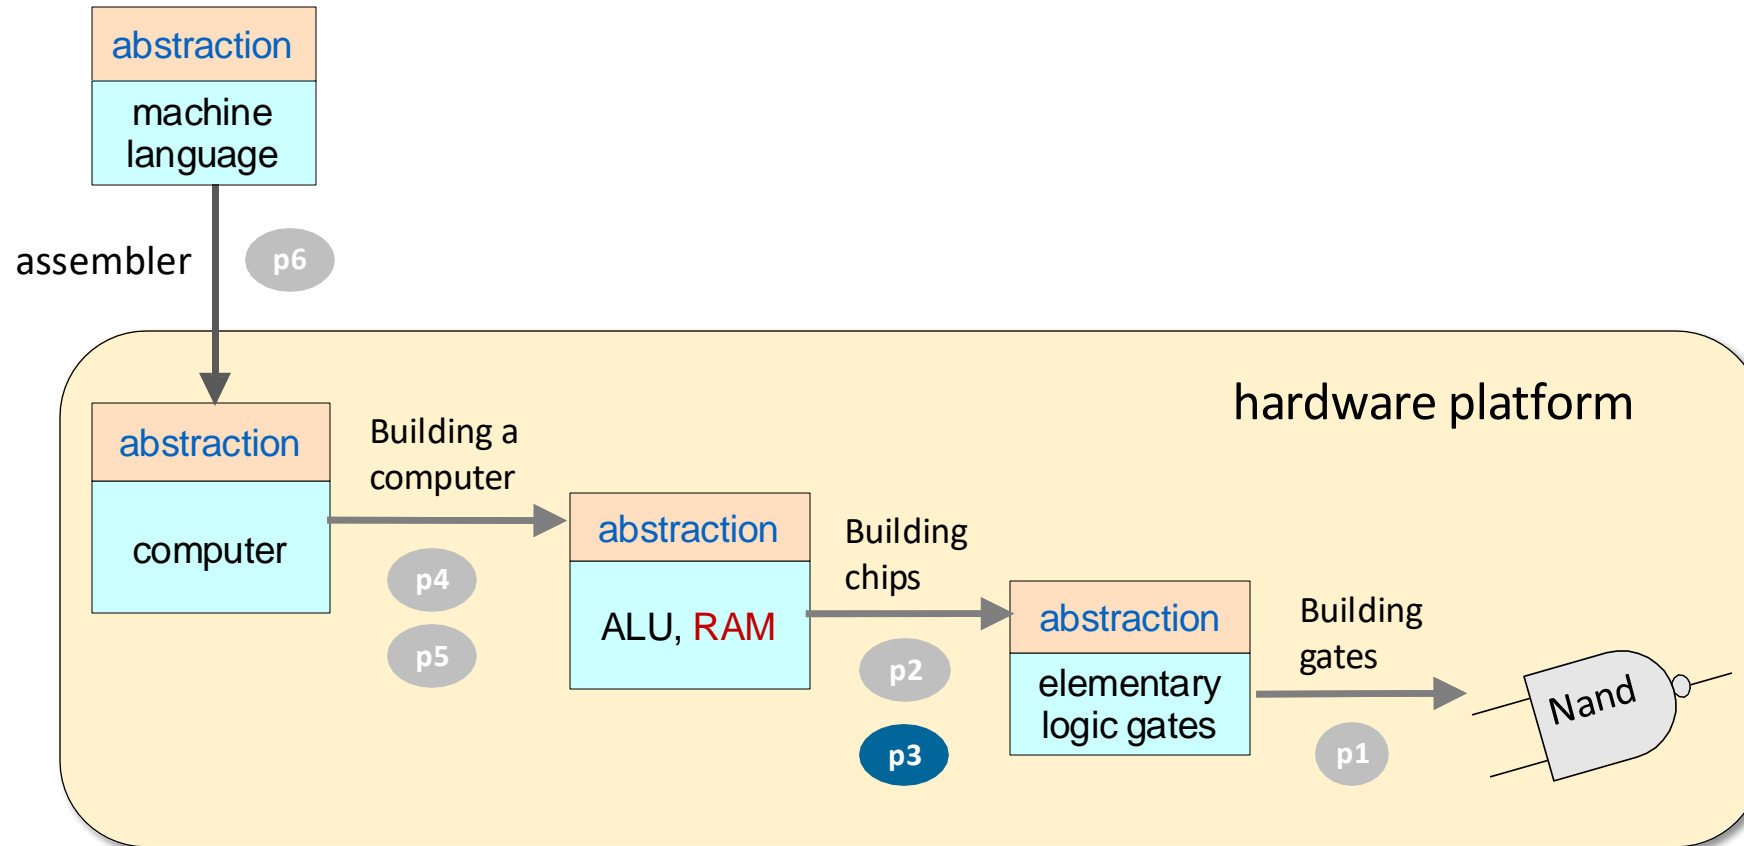


Implementation

- Data Flip Flop
- Registers
- RAM
- Project 3: Chips
- Project 3: Guidelines

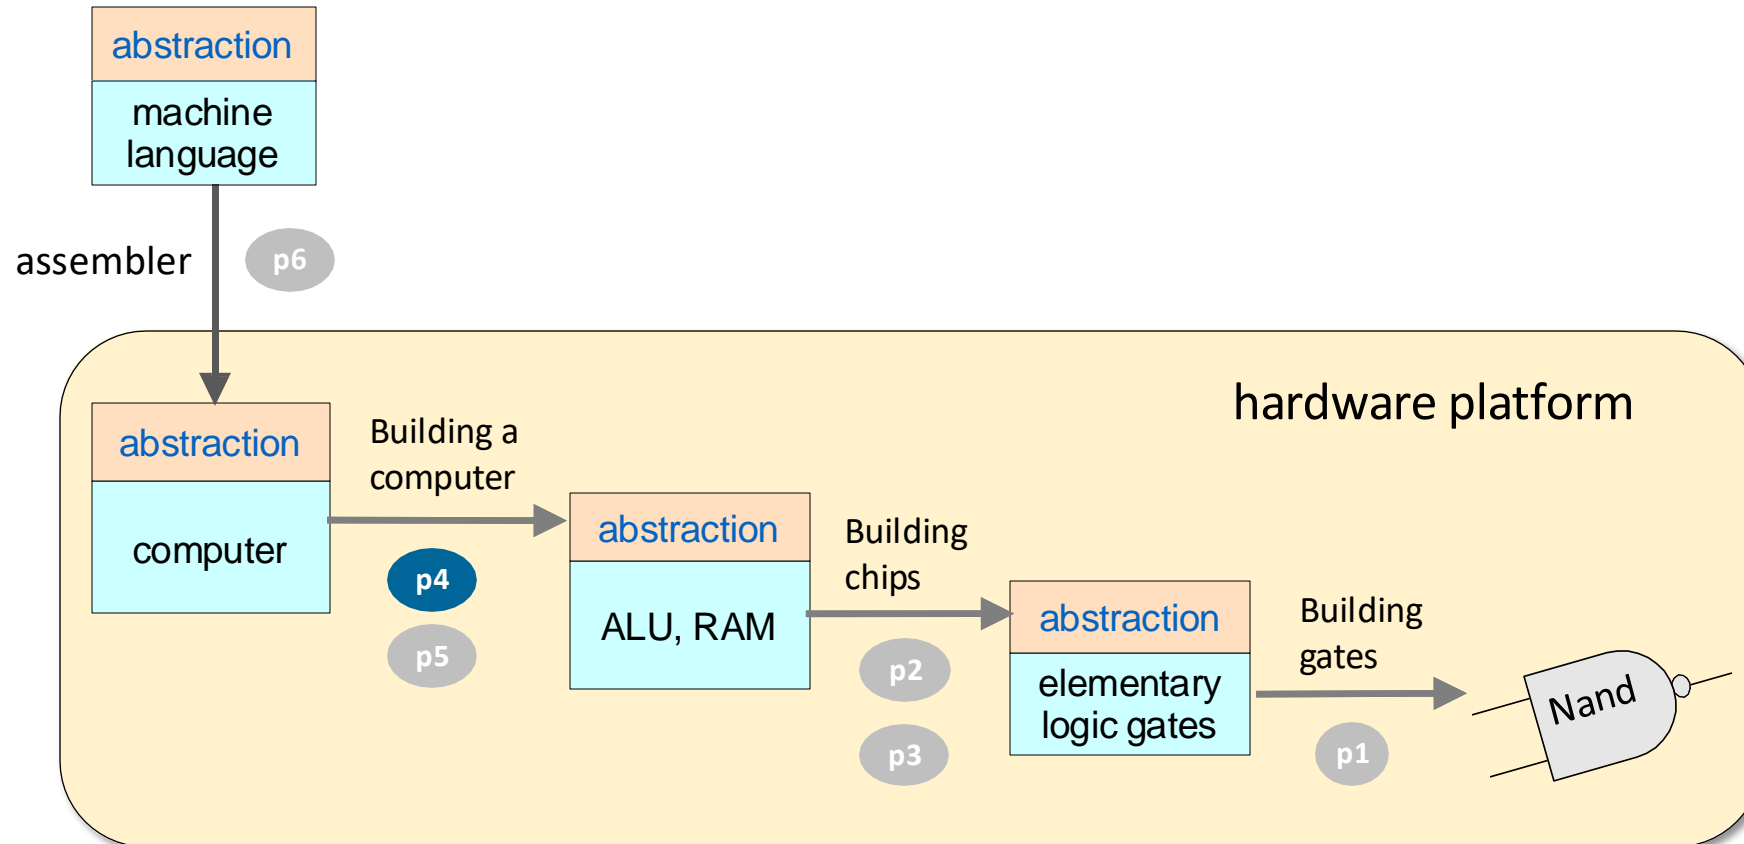


What's next?



This lecture / chapter / project:
Build the computer's RAM

What's next?



Next lecture / chapter / project:

- Get acquainted with the computer architecture
- Write machine language programs